

**EFFICIENT COMBINATORIAL SEARCH ALGORITHMS**

**A Thesis**

**Submitted to the Faculty**

**of**

**Purdue University**

**by**

**Chee Fen Yu**

**In Partial Fulfillment of the  
Requirements for the Degree**

**of**

**Doctor of Philosophy**

**December 1986**

## **EFFICIENT COMBINATORIAL SEARCH ALGORITHMS**

### **ACKNOWLEDGMENTS**

**A Thesis**

**Submitted to the Faculty**

**of**

**Purdue University**

**by**

**Chee Fen Yu**

**In Partial Fulfillment of the  
Requirements for the Degree**

**of**

**Doctor of Philosophy**

**December 1986**

I wish to express my thanks to my mentor, Professor Benjamin W. Wah. I also wish to thank Professor Philip H. Swain, Professor Jose A. Fortes, Professor John S. P. Wang, and the late Professor King-Sun Fu for their unstinting service on my thesis committee.

In addition, I gratefully acknowledge the support of the National Science Foundation through grants ECS81-05968 and DMC85-19649, and CIDMAC, a research unit of Purdue University, sponsored by Purdue University, Cincinnati Milicron Corporation, Control Data Corporation, Cummins Engine Company, Ransburg Corporation, and TRW.

<b>TABLE OF CONTENTS</b>	<b>Page</b>
<b>CHAPTER 1 - INTRODUCTION</b> .....	1
1.1 Introduction and Motivation .....	1
1.2 Terminology .....	2
1.3 Combinatorial Search Procedures .....	5
1.4 Improving the Efficiency of Search Procedures .....	9
1.5 Organization of this Thesis .....	11
1.6 Contributions of this Thesis .....	12
<b>CHAPTER 2 - B&amp;B ALGORITHMS IN A TWO-LEVEL MEMORY SYSTEM</b> .....	14
2.1 Introduction .....	14
2.2 A Comparison of the Best-First and Depth-First B&B Algorithms .....	22
2.3 B&B Algorithms in a Paged Virtual-Memory Environment .....	29
2.4 A Virtual Memory System for the Best-First B&B Algorithm .....	33
2.5 The Modified B&B Algorithm .....	35
2.5.1 Modified B&B Algorithm for a File System .....	35
2.5.1.1 Expected Overhead for Real-Valued Problems .....	39
2.5.1.2 Expected Overhead for Discrete-Valued Problems .....	40
2.5.1.3 Choosing Stack Boundaries .....	49
2.5.2 Modified B&B Algorithm for a Virtual-Memory System .....	49
2.5.3 Summary .....	51
<b>CHAPTER 3 - LEARNING DOMINANCE RELATIONS</b> .....	55
3.1 Introduction .....	55
3.2 Methodology .....	55
3.3 Problem Definition .....	58
3.4 Representation .....	59
3.5 Constraints in the Problem Domain .....	64
3.6 Applicable Learning Mechanisms .....	68
3.6.1 Learning Dominance Relations by Deduction .....	78
3.6.2 Learning Dominance Relations by Analogy .....	83
3.6.3 Learning Dominance Relations by Experimentation .....	87
3.6.4 Combining the Different Learning Mechanisms .....	88
3.7 Summary .....	90
<b>CHAPTER 4 - LEARNING DOMINANCE RELATIONS BY EXPERIMENTATION</b> .....	91
4.1 Introduction .....	91
4.2 MUUT-1 .....	93
4.3 Scheduler .....	93
4.4 Description Manager .....	95
4.5 Forester .....	97
4.6 Representation of Candidate Relations .....	97
4.6.1 Rule Space .....	102
4.6.2 Value Inheritance .....	104
4.6.3 Parallelism .....	105
4.7 Searcher .....	106
4.7.1 Cartographer .....	111
4.7.2 Assayer .....	117
4.7.3 Derivator .....	119
4.7.4 Tester .....	122
4.7.5 Scrutinizer .....	125
4.7.6 Generator .....	127
4.8 Verification .....	127
4.9 Summary .....	128
<b>CHAPTER 5 - CONCLUSIONS AND FUTURE WORK</b> .....	129
<b>LIST OF REFERENCES</b> .....	132
<b>LIST OF FIGURES</b> .....	v1
<b>ABSTRACT</b> .....	x
<b>CHAPTER 1 - INTRODUCTION</b> .....	viii

## APPENDICES

Appendix A : A Prototype of the Searcher .....	142
Appendix B : A Trace for the 0/1-Knapsack Problem .....	188
VITA.....	206

## LIST OF TABLES

Table	Page
1.1 Comparison of several polynomial and exponential time complexity functions.....	3
1.3 Effect of improved technology on several polynomial and exponential time algorithms.....	4
2.1 Number of nodes expanded by the best-first (BFS) and depth-first (DFS) B&B algorithms for randomly generated 30-object 0/1-knapsack, 80-node vertex cover, and 20-by-20 integer programming problems.....	30
2.2 Normalized overheads for the analytical (AA) and equal-width (EWA) boundary assignments for five randomly generated, 80-node vertex-cover problems and five randomly generated, 20-by-20 integer programming problems.....	50
2.3 Summary of B&B algorithms to be used.....	54
3.1 Similarities between the problems descriptions of the 0/1-knapsack, scheduling to minimize weighted completion time, and inventory problems.....	62
3.2 Examples of the different classes of dominance relations.....	67
3.3 Upper bounds on the implementation overheads of different classes of dominance relations (LEV(A) is the level of node A).....	69
3.4 Identity descriptions for scheduling to minimize cumulative completion time and scheduling to minimize cumulative lateness.....	85

Table	Page
4.1 Search tree $r_1$ .....	99
4.2 $\delta$ -values of parallel rules.....	107
4.3 Results of evaluating $n_c$ for several values of $m$ and $n$ .....	109
4.4 Possible outcomes of the matching process.....	124

Page	
4.1 Search tree $r_1$ .....	99
4.2 $\delta$ -values of parallel rules.....	107
4.3 Results of evaluating $n_c$ for several values of $m$ and $n$ .....	109
4.4 Possible outcomes of the matching process.....	124

#### LIST OF FIGURES

Figure	Page
--------	------

1.1 Search procedures.....	7
2.1 (a) A 0/1-knapsack problem instance and (b) its best-first B&B tree.....	18
2.2 A $s$ -way state-space tree.....	23
2.3 $E[N_B(0,0)]$ for a state-space tree with $k=2$ , $b=15$ , $s=10$ , $b=20$ ,.....	27
2.4 Differences between $E[N_B(0,0)]$ and $E[N_B(0,0)]$ for a state-space tree with $k=2$ , $b=15$ , $s=10$ , $b=20$ ,.....	28
2.5 A virtual-memory system for the best-first B&B algorithm.....	34
2.6 Data structure for the modified B&B algorithm.....	35
2.7 The modified B&B algorithm.....	38
2.8 Wall model of the best-first B&B algorithm.....	40
2.9 Simulation results for B&B algorithms on two 20-by-20 integer programming problems. (Main memory size is 128 Kbytes.).....	52
2.10 Simulation results for B&B algorithms on two 80-node vertex cover problems. (Main memory size is 128 Kbytes. The compilation time for the modified B&B scheme is the same for both 16 Kbytes/track and 64 Kbytes/track in the second problem.).....	53

Figure	Page
3.1 Description frames for the 0/1-knapsack problem.....	65

3.2 A simplified model of a machine learning system.....	70
--	----

3.3 Resolution tree for $=\{L(x), L(y)\} \& >\{R(x), R(y)\}$ & $<\{W(x), W(y)\} \rightarrow >\{B(x), B(y)\}$	81
3.4 Resolution tree for $=\{L(x), L(y)\} \& >\{R(x), R(y)\}$ & $>\{W(x), W(y)\} \rightarrow >\{B(x), B(y)\}$	82

3.5 Interactions between the learning mechanisms.....	80
4.1 Organization of MUTT-1.....	83
4.2 Frame describing value-description of 0/1-knapsack problem.....	95

4.3 (a) The p-frame for the 0/1-knapsack problem, and (b) the p-frame instance for the 0/1-knapsack instance, $\Pi_1$ .....	98
4.4 A fragment of the rule graph for the 0/1-knapsack problem.....	103
4.5 Organization of the searcher.....	110
4.6 $G^1$ .....	113
4.7 $G^2$ .....	114
4.8 $G^3$ .....	115
4.9 $G^4$ .....	116

Appendix	Page
A1 Listing of file knap.....	189
A2 Listing of file des.....	190
A3 Listing of file t0.....	101
A4 Listing of file t1.....	102

## ABSTRACT

Yu, Chee Pen. Ph.D., Purdue University. December 1986. Efficient Combinatorial Search Algorithms. Major Professor: Benjamin W. Wah.

The search for acceptable solutions in a combinatorially large problem space is an important problem in the fields of artificial intelligence, operations research, and computer science. Ways of improving the efficiency of this search process includes reducing the number of nodes searched by the process and reducing the average time to search each node.

One way to reduce the number of nodes searched is to use an appropriate selection scheme. Best-first branch-and-bound (B&B) algorithms generally expand less nodes than depth-first B&B algorithms. However, a best-first B&B algorithm requires exponential memory space. A depth-first B&B algorithm requires linear memory space. In a two-level memory system, this difference in space requirements can cause the average expansion time of each node to differ for the two B&B strategies. This thesis presents results on the efficiency of the B&B strategies in a two-level memory. The best B&B strategy depends on the characteristics of the problem domain as well as the characteristics of the two-level memory. A best-first B&B algorithm should be used when it expands much less nodes than the corresponding depth-first B&B algorithm, and when the secondary memory is very fast. A depth-first B&B algorithm should be used when it expands approximately the same number of nodes as the corresponding best-first B&B algorithm, and when the secondary memory is very slow. The choice for intermediate situations is not so clear. One result of

our research is the modified B&B algorithm, which can be used in these intermediate situations. The modified B&B algorithm has been designed to match the characteristics of a two-level memory system.

#### Dominance relations prune unnecessary nodes in search graphs and so

reduce the number of nodes searched by the search process. There is no systematic procedure to derive dominance relations because they are problem-dependent. A possible approach to this problem is the use of knowledge-based techniques. This thesis also contains a study of the machine learning of dominance relations. Among the results are a representation for optimization problems that exposes useful domain information, a classification of dominance relations that identifies the useful forms of dominance relations, and the applicable learning mechanism. This research also includes a study of a system that learns dominance relations by experimentation. A prototype of this system has been able to learn dominance relations for several scheduling and mathematical programming problems.

## CHAPTER I INTRODUCTION

### 1.1 Introduction and Motivation

Many problems in engineering and science involve finding one or more correct solutions among a set of possible solutions. The set of possible solutions constitutes the *problem space* for the problem.

A function  $f(n)$  is of  $O(g(n))$  when there exists a constant  $c$  such that  $|f(n)| \leq c|g(n)|$  for all value of  $n \geq 0$ . A complexity function has polynomial complexity if the complexity function is  $O(p(n))$  for some polynomial function  $p$ , where  $n$  denotes the problem size. Any complexity function that cannot be so bounded has exponential complexity. This definition of exponential complexity includes certain non-polynomial functions like  $n \log n$  which are not normally regarded as polynomial functions. A problem space is said to be combinatorially large if the number of possible solutions to the problem is of exponential complexity.

If the knowledge of the problem is complete in terms of the objects, properties, and relations relevant to the solution of the problem, then we can construct a rigorous computation model of the problem and solve it by evaluation. If the knowledge of the problem is incomplete, then it will be necessary to try some or all of the possible solutions in the problem space until the desired solution or set of solutions is found. In other words, a search of the problem space is required. Thus, a *combinatorial search problem* is a problem whose solution involves a search in a combinatorially large problem space and a *combinatorial search algorithm* is a search algorithm used to solve combinatorial search problems. Examples of combinatorial search problems include theorem proving, game playing, evaluating a logic program, solving a database query, designing a computer system, assigning registers for a compiler, finding the shortest path in a graph, evaluating the maximum flow in a network, solving a mathematical programming problem, and searching for a permutation order to sort a set of numbers. These problems exist in a large spectrum of engineering and science applications, including artificial intelligence and operations research.

Efficient search methods are very important for combinatorial search problems. An efficient search method requires significantly less time to solve any non-trivial instance of a combinatorial search problem than an inefficient search method. Often this difference in time requirements determines whether or not a problem instance may be solved. Table 1.1 is a dramatic illustration of the importance of efficient search algorithms for combinatorial search problems. It compares the growth rates among several complexity functions, where the functions express execution time in microseconds. Notice the much more explosive growth rates of the two exponential complexity functions.

Even more revealing is an examination of the effects of improved computer technology on algorithms having these time complexity functions. Table 1.2 shows how the largest problem instance solvable in one hour would change if we had a computer 100 or 1000 times faster than the present machine. Observe that with the  $2^n$  algorithm a thousand-fold in computing speed only adds 10 to the size of the largest problem instance we can solve in an hour, whereas with the  $N^6$  algorithm this size almost quadruples.

Before we consider the problem of improving the efficiency of combinatorial search algorithms, we will introduce some of the terminology and basic concepts of combinatorial search algorithms.

### 1.2 Terminology

Combinatorial search problems are divided into decision problems and optimization problems. A *decision problem* seeks any solution that satisfies a given set of constraints while an *optimization problem* seeks the solution that satisfies the set of constraints as well as optimizes a given objective function. A *feasible solution* is any solution that satisfies the given set of constraints and an *optimal solution* is any feasible solution that optimizes the objective function.

A search problem can be represented as a *search graph*. By convention, the places in a search graph are called *nodes* and the connections between nodes are called *links*. Each search graph has a *root node* and one or more *terminal nodes*. No link enters the root node while no link leaves a terminal node. A special case of the search graph is the *search tree* in which each node except the root has exactly one link entering it. A search graph can always be transformed into a search tree by duplicating common nodes. By convention, the links in a search tree are called *branches*. It is common to talk about trees

Table 1.1 Comparison of several polynomial and exponential time complexity functions.

Time Complexity Function	Size n					
	10	20	30	40	50	60
$n$	0.0001	0.0002	0.0003	0.0004	0.0005	0.0006
$n^2$	0.0001	0.0004	0.0009	0.0016	0.0025	0.0036
$n^3$	0.001	0.008	0.027	0.064	0.125	0.216
$n^5$	0.1	3.2	24.3	1.7	5.2	13.0
$2^n$	0.001	1.0	17.9	127	35.7	366
$3^n$	0.059	58	6.5	$3855$	$2 \times 10^8$	$1.3 \times 10^{13}$

Table 1.2 Effect of improved technology on several polynomial and exponential time algorithms.

Time Complexity	Size of Largest Problem Instance Solvable in 1 Hour	
	With Present Computer	With Computer 100 Times Faster
$n$	$N_1$	$100 N_1$
$n^2$	$N_2$	$10 N_2$
$n^3$	$N_3$	$4.64 N_3$
$n^5$	$N_4$	$2.5 N_4$
$2^n$	$N_5$	$N_5 + 0.64$
$3^n$	$N_6$	$N_6 + 4.19$
		$N_6 + 6.29$

Each branch in a search tree represents the assignment of a value to an unassigned parameter. For example, the unassigned parameters in 0/1-knapsack problem are the set of unassigned items. Expansion of a node involves the selection of an unassigned item and the creation of two alternatives: (a) include the item in the knapsack, thereby decreasing the capacity of the knapsack by the size of the item and increasing its profit by the value of the item, and (b) exclude the item from the knapsack.

Each node in a search tree, except the terminal nodes, can be either an AND node or an OR node. An AND node is a problem reduction from the problem represented by this node to multiple smaller subproblems represented by its children. To solve the problem represented by an AND node, all the subproblems represented by its children must be solved. An example of an AND node is one that adds the solutions from all the subtrees expanded from this node. In contrast, an OR node represents a state-space approach that transforms the state represented by this node, by operators represented by the arcs, to states represented by its children. An OR node is solved if any of its children can be solved. The need to search for the desired solution is due to the presence of OR nodes.

### 1.3 Combinatorial Search Procedures

There are two types of combinatorial search procedures: problem-specific procedures and general procedures. Problem-specific procedures are tailored to the problem characteristics. Examples of problem-specific procedures include polynomial-time approximation algorithms with guaranteed bounds [GaJ70, SaJ71], polynomial-time algorithms that optimize the average performance [Smib6] and algorithms that are guaranteed to yield optimal or near-optimal solutions with a given probability [Kar76]. However, experience has shown that efficient algorithms for slightly different problems can be very different and that it is very difficult to design good problem-specific algorithms for most

problems. We do not study problem-specific procedures in this thesis.

General procedures are applicable to a wide range of problems [KS74, Iba76b, Sel81], but may also incorporate some problem-specific knowledge. Figure 1.1 shows some of these general procedures.

The first group of general procedures are appropriate for problems involving an adversary: *minimax search*, *alpha-beta pruning*, *progressive deepening*, *heuristic pruning*, and *heuristic continuation*. They are common in programs that play board games and use techniques that depend on the presence of an adversary. We will not study these procedures in this thesis.

The second group of procedures consists of simple, basic procedures which can only be used for decision problems. This group of procedures includes *depth-first search*, *hill climbing*, *breadth-first search*, *beam search*, and *best-first search*. Each of these procedures has the following basic form:

```

insert root node into active list;
while active list is not empty or goal has not been reached [
    remove first element from the active list;
    expand the first element;
    add the first element's children, if any, to the active list; ]
    if goal node has been found then announce success else announce failure
The different procedures differ mainly in the way they add the first element's children to the active list. A depth-first search always adds them to the front of the active list but a breadth-first search always adds them to the end of the active list. Hill climbing procedures use a heuristic measure of the distance from the node to its nearest solution node (i.e., goal distance). A hill climbing procedure sorts the best element's children by their heuristic values before adding these children to the front of the active list. Beam search is like breadth-first search but only expands the best w nodes (by the estimated remaining distance) at each level. A best-first search orders the entire active list by the estimated goal distance.
```

The relative efficiencies of these methods depends on the nature of the problem space. A depth-first search is good when blind alleys do not get too deep. A breadth-first search is good when the number of alternatives at the choice points are not too large. Hill climbing is good when there is a natural measure of goal distance and a good choice is likely to be among the promising choices at each choice point. Beam search is good when there is a natural measure of goal distance and a good path is likely to be among the promising

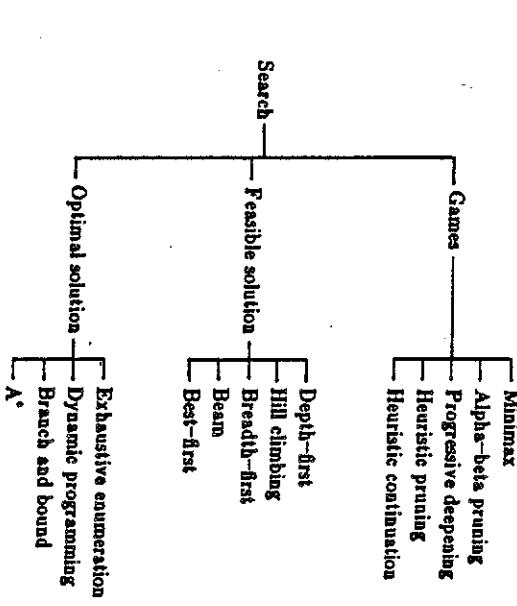


Figure 1.1 Search procedures.

partial paths at all levels. Best-first search is good when there is a natural measure of goal distance and a good path may look bad at shallow levels.

The final group of general procedures finds the best solution. This group includes exhaustive enumeration, discrete dynamic programming [BeD62, DiL77, Kal67, Iba73, Whi68], branch and bound (B&B) [LaW68, Mit70], and A\* [Ni71].

Exhaustive enumeration tries all possible solutions. The other procedures discards paths that are known not to lead to the optimal solution. Thus, these procedures have the following basic form:

```

insert root node into active list;
while active list is not empty or goal has not been reached [
    if first element in active list is not the goal node then]
        remove first element from the active list;
        expand the first element;
        prune unnecessary nodes, if any;
        add the first element's children, if any, to the active list; ]
]

if goal node has been found then announce success else announce failure
```

The dynamic programming principle [BeD62] states that the path from a node, S, to a non-terminal node, I, does not affect the path from node I to a solution node, G. Consequently when looking for the best path from S to G, all paths through the intermediate node, I, other than the best path from S to I, can be ignored. Dynamic programming procedures may use depth-first, or breadth-first, or best-first organization of the active list. The key concept in branch-and-bound algorithms is the *lower bound value*, which is a heuristic underestimate of the value of the best solution that can be derived from a node. No further work needs to be done on a node whose lower-bound value exceeds some known solution because the value of the best solution derived from this node cannot be smaller than its underestimate. This principle is known as the *lower-bound test* [LiW68]. The method used to order the active list leads to the identification of three different B&B algorithms: depth-first B&B, breadth-first B&B, and best-first B&B. The A\* algorithm combines the best-first B&B algorithm with the dynamic programming principle. The B&B procedures and dynamic programming algorithms have been extended to include general pruning rules known as *dominance relations* [BeD62, AlH74, Iba74].

Exhaustive enumeration is only useful for small problem spaces. Dynamic programming is good when many different paths reach common non-terminal nodes in the search tree. Branch-and-bound search is good when there is a

good lower-bound. The A\* procedure is good when both the basic B&B procedure and dynamic programming are good. Both A\* and dynamic programming may be viewed as a special case of the B&B search with dominance relations [KuK83].

#### 1.4 Improving the Efficiency of Search Procedures

A comparison of the general procedures described in Section 1.3 shows that the more efficient procedure is generally one that searches a smaller portion of the problem space, i.e., the procedure with the smaller search tree (or graph). Ways of reducing the size of the search tree (or graph) include the use of appropriate representation, the use of an appropriate selection scheme, and the use of appropriate pruning schemes. Another way to reduce the time required to solve an instance of the problem is to reduce the time required to expand each node. As suggested by Table 1.2 in Section 1.1, this method is only useful for procedures that have reasonable time complexities. It is not very useful if used to extend the solvable problem size.

A problem can often be represented in multiple forms. For example, the 0/1-knapsack problem may be represented by either an OR-tree and solved with a B&B procedure [IlsS78] or an AND/OR-tree and solved by dynamic programming. It is possible for different representations to lead to procedures with drastically different time complexities. A dramatic example of this disparity in time complexities is the checkerboard and dominoes problem. The upper left and lower right squares are removed from a  $n \times n$  checkerboard. Each domino is one square by two squares in size. The problem is to decide whether or not the checkerboard can be covered completely by dominos. An OR-tree description leads to search procedures that try different arrangements of the dominos on the checkerboard. Such procedures have exponential complexity. This problem can as be represented as the repeated removal of two adjacent squares from a board consisting of  $n^2 - 2$  red squares and  $n^2$  black squares. This representation leads to a trivial solution of the problem.

Although efficient representations have been established for many search problems, the general question of deriving the optimal representation is an open question, especially the kinds of innovative problem descriptions which can lead to dramatic improvements in the effort required to solve a problem. We do not explicitly study the problem of the optimal representation for a particular problem. However we will investigate some representation issues in conjunction with the other ways of reducing the size of the state tree.

The selection scheme orders the elements of the active list and thus chooses nodes for expansion (which is always the first element in the active list). The three basic selection schemes are depth-first selection, breadth-first selection, and best-first selection. For problems with a suitable lower-bound function, a best-first scheme always has the smallest search tree (or graph) if no node has a lower-bound value equal to the optimal solution value [LaW86, LiW84, PeS84]. B&B procedures commonly use relaxation-guided lower-bound functions. A relaxation of a discrete optimization problem is a problem with the same set of inputs and a set of feasible solutions that is a superset of the set of feasible solutions for the discrete optimization problem, and has a cost function such that common feasible solutions have the same cost in both problems. A relaxation-guided lower-bound function is a fast algorithm for solving the relaxed problem. Such lower-bound functions are problem dependent and are very difficult to design because they require extensive knowledge about the different types of problems.

A pruning scheme identifies and eliminates ineffective nodes known not to lead to a solution better than another existing node. One pruning scheme is the dynamic programming principle—if two nodes have identical states, then the search procedure needs only expand one of the two nodes.

Another pruning scheme involves dominance relations [BaD82, AHU77, HoS77, HoS88]. Let subtree  $T_i$  be the subtree rooted at node  $N_i$ . If  $N_i$  dominates  $N_j$ , then  $T_i$  contains a solution node with a value that is no worse than the value of the best solution in  $T_j$ . The converse is not true. A solution in  $T_i$  no worse than the best solution in  $T_j$  does not imply the presence of a dominance relation between  $N_i$  and  $N_j$ . Good dominance relations can result in a search tree of polynomial complexity.

As an example, consider the problem of finding the shortest path between city  $S$  and city  $T$ . Let  $P_1$  and  $P_2$  be paths from  $S$  to an intermediate city  $N$  and  $P_3$  be a path from  $N$  to  $T$ .  $P_1$  dominates  $P_2$  if  $P_1$  is shorter than  $P_2$ . For each  $P_3$ , the path with length  $P_1 + P_3$  is shorter than the path with length  $P_2 + P_3$ . Without the above dominance relations, the shortest-path problem requires exponential time.

The lower-bound test used in B&B algorithms is a dominance relation. The pruning power of the lower-bound tests depends on the quality of the lower-bound function [WaY85]. In the lower-bound test, the best solution found so far dominates all nodes with lower bounds not greater than its solution value. Other more general dominance relations have also been used [HoS78], but there is no systematic procedure to derive dominance relations

because they are problem-dependent. Identifying and encoding problem-specific dominance relations is a painstaking and difficult process, particularly for complex problems. One possible approach to this problem is to use artificial intelligence techniques, specifically, machine learning techniques. LiS85] and will not be studied in this thesis.

For many serial computers, the desire to have memories of large capacity and moderate cost have resulted in the use of multi-level memories. In such memories, the cost of accessing an item of information depends on its location. The effects of this architectural feature has not been considered in previous work in the design of search procedures. It can affect the choice of the appropriate selection scheme as well as the implementation of the pruning scheme. The B&B algorithms provide a convenient context for studying this issue since the set of B&B algorithms includes all three selection schemes.

In short, the objective of this thesis is to improve the efficiency of serial combinatorial search algorithms. This objective is achieved by studying effective B&B algorithms in a two-level memory environment and by studying machine learning of dominance relations.

### 1.5 Organization of this Thesis

The remainder of this thesis is organized into four chapters. The next chapter deals with the question of efficient B&B algorithms in a two-level memory. This chapter is followed by two chapters dealing with machine learning of dominance relations. The final chapter contains conclusions and suggestions for future work.

Chapter 2 attempts to find the best B&B algorithm on a two-level memory system. We first establish, in two steps, that the best-first B&B algorithm is not always the best B&B algorithm on a two-level memory system in two steps. We will show that the difference between the number of nodes expanded by the best-first B&B algorithm and the number of nodes expanded by the depth-first B&B algorithm for a given problem instance can take on a wide range of values depending on the accuracy of the lower-bound function. We will then show that for some values of this difference, the depth-first B&B algorithm is better than the best-first B&B algorithm. Having shown that the best-first B&B

algorithm is not always the best choice, we then seek to develop guidelines for choosing the appropriate B&B algorithm. This includes the development of a new B&B algorithm, the modified B&B algorithm, which can be used when neither the best-first B&B algorithm nor the depth-first B&B algorithm is particularly suitable.

Chapter 3 considers task of learning dominance relations. The methodology used follows the caveats of a good methodology in artificial intelligence as outlined by Winston [Win84]. We have focused this study on the number optimization problem. By definition, some objects in a number optimization problem must be numbers, and the magnitude of the largest number cannot be bounded by a product of a polynomial and the number of symbols used to describe the problem. One result of applying this methodology is a problem representation which exposes semantic information about the problem. Another result is a classification of dominance relations that uncovers constraints in the domain of dominance relations. These constraints permit some classes of dominance relations to be eliminated from consideration. We have also identified the applicable learning mechanisms, and characterized their strengths, areas of application, and limitations. The primary learning mechanism is experimentation while the secondary learning mechanisms are analogy and deduction. The study of each mechanism is a major research problem in itself and we have concentrated on the primary mechanism.

Chapter 4 discusses a prototype program which learns dominance relations by experimentation. This program uses a generate-and-test paradigm based upon heuristic constraints derived from manipulations of the standard descriptions as well as constraints in the space of relations to reduce the number of candidate relations proposed and tested. In addition, the program is able to automatically define and manipulate new variable symbols. This chapter concludes with a summary of the problems to which this program has been applied and the results that we have obtained.

useful in situations where both the depth-first and best-first B&B algorithms do not perform well. This new B&B algorithm demonstrates the effectiveness of including architectural considerations in the design of an algorithm; in particular, it shows the effectiveness of the no-wait principle in the design of non-deterministic search algorithms.

As indicated in Section 1.3, the problem of deriving dominance relations is an important, open research problem. We do not claim to have solved this problem. Nonetheless, we believe that the work in this thesis has taken several important steps towards an eventual solution. Perhaps, the most important step is the demonstration that machine learning of dominance relations is possible. We have also laid the ground work for future efforts with our work on the fundamental issues in this problem: a problem representation which exposes relevant semantic information, a classification of dominance relations which identifies the useful classes of dominance relations, and a characterization of the applicable learning mechanisms.

The work on learning dominance relations by experimentation also contributes to the field of learning by experimentation, in particular, learning problem solving heuristics by experimentation. Neither number optimization problems nor pruning heuristics between two arbitrary nodes on the same level have been previously studied in experimentation research. Specific contributions in this area include the mechanisms (knowledge) needed to drive such a learning system, and the work on automatic derivation of new semantically sensible symbols.

#### 1.6 Contributions of this Thesis

The work in this thesis has contributed to a better understanding about B&B algorithms in a two-level memory. It shows a need to revise the belief that the best-first B&B algorithm is always the best B&B algorithm. This belief is true only in a single-level memory system. In a two-level memory system, the best B&B algorithm depends on the characteristics of the problem and the memory system. We have also developed a new B&B algorithm that is

## CHAPTER 2 B&B ALGORITHMS IN A TWO-LEVEL MEMORY SYSTEM

The branch and bound (B&B) algorithm [Law68, Mirro] has emerged as the principal method for solving constrained optimization problems for which 'direct' methods of solution either do not exist or are inefficient. Formally, a constrained optimization problem is a problem having the form:

$$\begin{aligned} \text{Minimize} \quad & C_0(x) \\ \text{subject to} \quad & g_i(x) \leq 0 \quad i = 1, 2, \dots, m \end{aligned} \quad (2.1)$$

$X$  represents the domain of optimization defined by the  $m$  constraints, normally an Euclidean  $n$ -space, and  $x$  denotes a vector  $(x_1, x_2, \dots, x_n)$ . A *feasible solution* is a solution vector that lies in  $X$ . An *optimal solution* is a feasible solution for which  $C_0(x)$  is minimal. Among the constrained optimization problems to which the B&B algorithm has been applied are scheduling problems [Len70], puzzles and other cognitive tasks [Nil71], optimal decision trees [Mam78], knapsack problems [Grif70, HoS74, InK77], traveling salesman problems [BeN68, BeM71, Gar73], flow shop and job shop sequencing [GSS51], facility allocation [EIR66], and integer programming [GaN72, GaM72]. It has also been suggested by Kumar and Kanal [KuK83] that the B&B algorithm is a generalization of many heuristic search algorithms including A\* [Nil71], AO\* [Nil71], SSS\* [StR83], B\* [Ber70], alpha-beta [Knu75] and dynamic programming [BeD62, DrL77, Kal67, Iba73, Whi69].

The search process generated by the B&B algorithm is often represented by a search tree. Nodes in this tree represent subinstances. We will use the terms subinstance and node interchangeably in the rest of this chapter. Branches in the search tree represent the decomposition of subinstance  $P_i$  into additional subinstances.

The B&B algorithm has four major components. The *branching rule* determines if and how a given node is to be decomposed into additional nodes. The branching rule creates multiple nodes by heuristically selecting some

unassigned parameter(s) in the node and assigning alternative values to the selected parameter(s). For example, the unassigned parameters in the traveling-salesman problem are the set of untraversed edges. The B&B process expands a node by selecting an untraversed edge  $(i,j)$  and creating two alternatives: (a) the edge is traversed and the salesman goes directly from City  $i$  to City  $j$ , and (b) the edge is not traversed and the salesman stays at City  $i$ .

The second component of the B&B algorithm is a *lower-bound function*. Let  $g(P_i)$  be the value of the best solution that can be derived from node,  $P_i$ . The lower-bound function,  $g$ , must satisfy the following conditions:

(a)  $g(P_i)$  never exceeds  $|f(P_i)|$ ;

(b)  $g(P_i)$  has the same value as  $|f(P_i)|$  whenever  $P_i$  is feasible solution;

(c) every child of  $P_i$  has a lower-bound value that exceeds  $g(P_i)$ .

The lower-bound function is problem-dependent. A linear program with relaxed integer constraints is often used in the integer-programming problem [LaD80]. An assignment algorithm or a spanning-tree algorithm can be used for the traveling-salesman problem [Faas58].

The third component of the B&B algorithm, the *elimination rule*, uses the lower-bound function to eliminate those nodes which are known not to contain an optimal solution to the problem instance. The *incumbent*  $z$  is the best solution value found so far in the B&B process. In minimization problems, a node  $P_i$  is pruned if its lower-bound exceeds the incumbent as it cannot lead to a solution value better than the incumbent. In other words,  $P_i$  is eliminated if:

$$g(P_i) \geq z \quad (2.2)$$

The B&B process continues until all nodes have been either expanded or eliminated.

A relaxed form of this elimination rule will find a suboptimal solution with a guaranteed accuracy [LaW68]. Suppose that a deviation of 10% from the optimum is tolerable. A feasible solution of 150 permits the elimination of all nodes with lower-bounds not less than  $136.4$  (or  $150/(1+0.1)$ ) as these nodes cannot lead to solutions which are better than 150 by more than 10%. This technique significantly reduces the amount of intermediate storage and time required by the B&B process. We can define the *allowance function*  $\epsilon(z): \mathbf{R} \rightarrow \mathbf{R}$  (set of reals) such that  $P_i$  is terminated if:

$$g(P_i) \geq z - \epsilon(z). \quad (2.3)$$

The final incumbent  $z_F$  will be related to the optimal solution value  $z_0$  by the inequalities [Iba76a]:

$$z_F - \epsilon(z_F) \leq z_0 \leq z_F$$

Two common allowance functions are:

$$\epsilon(z) = \epsilon \quad \epsilon \geq 0 \text{ (absolute error deviation) and} \quad (2.5)$$

$$\epsilon(z) = \frac{\epsilon^2}{1+\epsilon} \quad \epsilon \geq 0, z \geq 0 \text{ (relative error deviation)} \quad (2.6)$$

The final component of the B&B algorithm is the selection rule that chooses one of the currently active nodes for expansion. The B&B process maintains the set of active nodes in an ordered list and always selects the first element of this list. A breadth-first B&B algorithm maintains the active list in a first-in/first-out order; a depth-first B&B algorithm maintains the active list in a last-in/first-out order; a best-first B&B algorithm maintains the active list in increasing order of lower-bounds.

The B&B algorithm for minimization problems can be summarized as follows:

```

incumbent z = infinity;
active list = {P0};
while active list ≠ ∅ do [
    apply selection rule to active list;
    expand selected node by the branching rule;
    for P ∈ {children generated} do [
        if P is a feasible solution then [
            update incumbent z;
            apply elimination rule];
        else if g(P) < (z - ε(z)) then
            insert P into active list
    ]
]
```

We will illustrate the B&B process with the integer-programming problem [Mit70]. The integer-programming problem is the constrained optimization:

$$\begin{aligned}
&\text{Minimize} && CX \\
&\text{subject to} && AX \geq B \\
&&& x_i: \text{non-negative integer}, i=1, 2, \dots, n; \\
&&& A \text{ is a } 1\text{-by-}n \text{ constant matrix;} \\
&&& B \text{ is a } m\text{-by-}1 \text{ constant matrix;} \\
&&& C \text{ is a } n\text{-by-}m \text{ constant matrix;}
\end{aligned}$$

These problems differ from ordinary linear-programming problems in that the variables can only have non-negative integer values.

One formulation of the B&B algorithms for the integer-programming problem uses the dual simplex method as the lower-bound function. The dual simplex method ignores the integrality constraints on the values of the  $x_i$ 's and solves the problem instance as a linear program. If the simplex solution is integral, then a feasible solution has been generated; otherwise the B&B process applies the branching rule. The branching rule selects a variable with a non-integral value (say  $x_1 = 4.4$ ) and uses this variable to create two new nodes.

One node restricts the value of  $x_1$  to not more than the floor of its value in the simplex solution ( $x_1 \leq \lfloor 4.4 \rfloor$  or  $x_1 \leq 4$ ). The other node restricts the value of  $x_1$  to not less than the ceiling of its value in the simplex solution ( $x_1 \geq \lceil 4.4 \rceil$  or  $x_1 \geq 5$ ). The branching rule always chooses the variable with the greatest (up or down) penalty. The up penalty for a variable  $x_i$  with value of  $s_i$  is the estimated increase in the simplex solution value due to the introduction of the integral constraint  $x_i \geq s_i$ . The down penalty is similarly associated with the constraint  $x_i \leq s_i$ . The lower-bound of each new node is the sum of the value of the simplex solution of its parent and the associated penalty.

Figure 2.1(b) shows the best-first B&B tree for the problem in Figure 2.1(a). Initially, there is only the problem instance so the best-first B&B process selects this problem instance for expansion. Application of the dual simplex method yields a solution of 14.2. This is not a feasible solution as the variables have non-integral values and the B&B process calculates the up and down penalties.  $x_1$  has the greatest penalty ( $U = 1.8$ ). This results in the creation of two new nodes, one with  $x_1 = 0$ , and the other with  $x_1 \geq 1$ .

The best-first B&B process always selects the active node with the smallest lower-bound;  $N_2$  has a smaller lower-bound than  $N_1$ . Application of the dual simplex method to node  $N_2$  yields a simplex solution in which all variables have integral values. This feasible solution has a value of 15.  $N_3$  has a lower-

$$\begin{aligned} \min x_0 &= 7x_1 + 3x_2 + 4x_3 \\ x_1 + 2x_2 + 3x_3 &\geq 8 \\ 3x_1 + x_2 + x_3 &\geq 5 \\ x_1, x_2, x_3 &\geq 0, \text{ integer} \end{aligned}$$

(a)

Optimal dual simplex solution

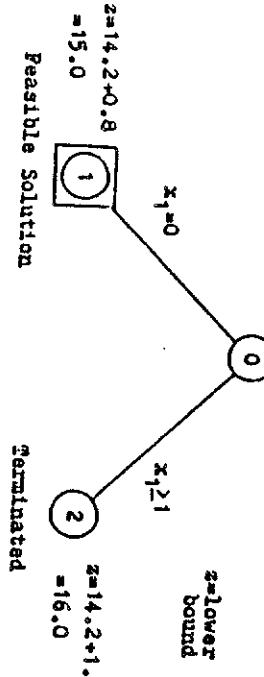
$$x_0 = 14.2$$

$$x_1 = 0.4$$

$$x_2 = 3.8$$

$$x_3 = 0$$

Variable	Down Penalty	Up Penalty
$x_1$	0.8	1.8
$x_2$	0.3	0.13



bound of 15.0 and the elimination rule eliminates  $N_4$ . There are no more active nodes and the best-first B&B process terminates with an optimal solution of 15.

The best-first B&B algorithm expands the smallest number of nodes if no node other than the optimal solution node has a lower-bound value equal to the value of the optimal solution [LiW86, LiW84]. Unfortunately, the best-first B&B algorithm has a maximum memory-space requirement whose expectation is an exponential with a sublinear exponent [WaY82, WaY85]. In contrast, the number of nodes in last-in/first-out stack of the depth-first B&B algorithm never exceeds the height of the search tree. For many search problems, the height of the search tree is equal to the number of variables in the problem and is reasonably small. The space requirement of the breadth-first B&B algorithm is similar to that of the depth-first B&B algorithm. Thus, for any non-trivial problem instance, the best-first B&B algorithm will require much more memory than the other two B&B algorithms.

Computer memories of very large overall capacity are often organized into a hierarchy of at least two levels, main memory and secondary memory. The smaller main memory usually consists of faster but more expensive semiconductor memory, while the larger secondary memory usually consists of slower disk memory. A program can only reference its information (i.e. instruction code and data) when it resides in main memory; information having immediate likelihood of being referenced should be kept in main memory. This storage allocation problem, i.e. determining how information is to be distributed among the levels of memory at each moment of time, must be solved before a memory hierarchy can be used effectively.

An elegant and universally accepted solution to the storage allocation problem is the idea of virtual memory [Den70]. At the heart of this idea is the notion that the memory address used by a program to reference information is a concept distinct from the physical memory location in which the information item is held. The computer hardware and software directs program-generated addresses to the memory locations containing the information addressed and moves information between main memory and secondary memory. The term *virtual address* denotes an address used by the programmer, while the term *memory address* denotes an address used by the memory. There is no a priori correspondence between virtual addresses and memory addresses, so an *address map* associates virtual addresses with memory addresses.

Figure 2.1 (a) A 0/1-knapsack problem instance and (b) its best-first B&B tree.

missing item is retrieved from secondary memory and placed in main memory. The mapping device then redefines the address map and completes the reference. Whenever main memory is full, some item has to be removed to make way for the item entering, the particular item being chosen at the discretion of the *replacement rule*. There is also a *fetch rule* which decides when to load an item. A *demand rule* loads an item into main memory only when a fault for it occurs. *Prefetching* is said to occur if action is taken to load an item before it is referenced.

An address map of reasonable size requires the grouping of information into sets of contiguous addresses called blocks. The entries in the mapping table will now refer to blocks, which are far less numerous than individual addresses. One method-segmentation-organizes address spaces into blocks (segments) of arbitrary size. Another-paging-organizes memory space into blocks (pages) of fixed size. Paging is the more popular of the two methods. The equal sized blocks of main memory are known as *page frames*. Page frames serve as sites of residence for pages which are blocks of virtual addresses of the same size.

Virtual memory gives the programmer the illusion of a very large main memory even though the computer has a relatively small main memory. However, virtual memory is not without problems. Space in a virtual memory may be an illusion; unnecessarily large and carelessly organized programs may generate excessive overhead in the automatic storage allocation mechanism. Virtual memory management mechanisms perform most efficiently when programs tend to exhibit good *locality*, i.e. they tend to concentrate their references in small regions of address space. Programs with poor locality will continually be interrupted while missing items are retrieved from secondary memory.

Studies have been carried out on non-search algorithms in a paged virtual-memory environments. These include rules for enhancing locality [Mor86], automatic restructuring of programs [HaG71, Fer74, Fer75, Fer76, BaS76], matrix manipulation [McG60, FP70], sorting [BGH70, AEH84], and database operations [SmI78].

In this chapter, we investigate B&B algorithms in a single-user computer system with a two-level memory. The time-efficiency of an algorithm in single-user systems may be characterized by its expected completion time.

We will start by establishing the need to revise the belief that the best-first B&B algorithm is always the best B&B algorithm. Previous theoretical studies of B&B algorithms [Bar73, Bar77, Bar78, K857, and

NiI80] have concentrated on the complexity of the algorithm, either its worst-case behavior or its average complexity. These complexity functions are not precise enough to guide a consideration of the trade-offs between the different B&B algorithms. The effects of the accuracy of the lower-bound function on the average performance of the A\*-algorithm (best-first search) and informed backtracking (depth-first search) have been studied for decision problems [Per81]. Other models of the cutoff mechanism which do not explicitly consider the accuracy of the lower-bounds have also been used to study the average performance of best-first B&B algorithms [SmI84] and depth-first B&B algorithms [SmI84, Sto84]. Section 2.2 describes an analysis of the numbers of nodes expanded by best-first and depth-first B&B algorithms using the accuracy of the lower-bound function and the distribution of the solution values. This will show that the possible differences between the number of nodes expanded by the best-first B&B and depth-first B&B algorithms span a wide range of values, including values where the best-first B&B algorithm may not be the best B&B algorithm.

Section 2.3 investigates the relationship between the number of nodes expanded and the effectiveness of the best-first B&B and depth-first B&B algorithms. We then consider alternatives that are effective in situations when neither the best-first B&B algorithm nor the depth-first B&B algorithm is effective. Section 2.4 describes a virtual memory management system tailored to the characteristics of the best-first B&B algorithm. Section 2.5 considers the modified B&B algorithm that is a compromise between the depth-first and best-first B&B algorithms.

Both of these alternatives are of limited value in multi-user systems. A specialized virtual-memory management system adversely affects other users who are not running B&B algorithms. Moreover, there is a lack of a suitable metric for the time efficiency that can be used as an optimization measure. The completion time of a program in a multi-user system depends on both the algorithm and the system load. Other possible metrics such as the processing time and the sum of the processing and secondary-storage times do not adequately describe either the trade-off or the overlap between computations and secondary-memory accesses.

### 2.3 A Comparison of the Best-First and Depth-First B&B Algorithms

Any analytical model of the depth-first B&B algorithm capable of predicting the number of node expansions must consider the accuracy of the lower-bound function. To compare the performances of the B&B algorithms for all possible distributions of the accuracy is a non-trivial, if not impossible, problem. This section presents a mathematically tractable solution that assumes a constant distribution for the accuracy.

We assume a state-space which is a complete  $s$ -way tree of height  $h$  (Figure 2.2). The root of the tree is at Level 0, and the leaves (solution nodes) are at Level  $h$ . Let  $\mu(i) = s^i$  be the number of nodes in Level  $i$ , and let  $P_{ij}$ ,  $0 \leq i \leq h$ ,  $0 \leq j \leq \mu(i)-1$ , be the  $j$ 'th node in the  $i$ 'th level of the tree. In addition, let  $S_j$  be the value of solution node  $P_{hj}$ ,  $0 \leq j \leq \mu(h)-1$ . We assume that the solution values are independently and identically distributed (iid) random variables with distribution function  $F_S(s)$  and density function  $f_S(s)$ , where

$$f_S(x) = \begin{cases} \frac{\beta e^{\beta x}}{(e^{\beta b} - e^{\beta a})} & a \leq x \leq b \\ 0 & \text{otherwise} \end{cases} \quad (2.8)$$

This is a truncated exponential distribution between  $a$  and  $b$  and it represents many interesting situations as  $\beta$  is varied. When  $\beta = -\infty$ , it is a delta function at  $a$  and represents the situation in which all solutions are optimal. As  $\beta$  is increased, the standard deviation of the distribution first increases and then decreases. At  $\beta = \infty$ , the function is a delta function at  $b$ . A large standard deviation may be interpreted as a low probability of finding solution values close to the optimal solution value.

The subtree rooted at  $P_{ij}$  contains  $\mu(h-i)$  solution nodes, so its best solution value,  $z_{ij}$ , is

$$z_{ij} = \min \left[ S_{j+\mu(h-i)}, \dots, S_{j+\mu(h-i)-1} \right] \quad (2.9)$$

The optimal solution,  $z_0$  will be given by  $z_0, z_{1,0}, \dots, z_{1,\mu(1)-1}$  are iid random variables because the  $S_j$ 's are iid random variables and their distribution is

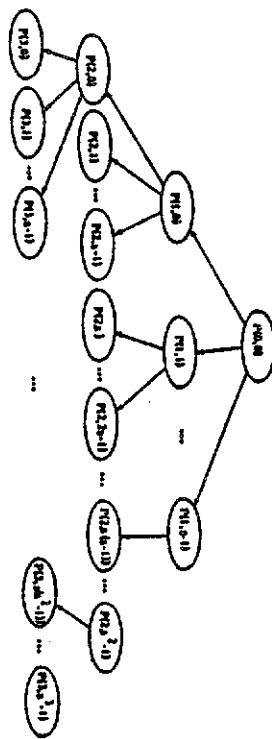


Figure 2.2 A  $s$ -way state-space tree.

$$F_{s,i}(x) = \begin{cases} 0 & x < a \\ 1 - [1 - F_S(x)]^{\mu(b-i)} = 1 - \left[ \frac{e^{\beta_b} - e^{\beta_a}}{e^{\beta_b} - e^{\beta_a}} \right]^{\mu(b-i)} & a \leq x \leq b \\ 1 & x > b \end{cases} \quad (2.10)$$

Let  $\xi_{ij}$  be the lower-bound of node  $P_{ij}$ , and let  $\delta_{ij}$  be the accuracy of the lower-bound function at node  $P_{ij}$ .  $\xi_{ij} = \delta_{ij} z_{ij}$ . By definition,  $\delta_{ij} = 1$  as  $P_{kj}$  is a solution node. We have only studied the case of  $\delta_{ij} = \delta$  in this thesis, where  $i < b$  and  $\delta$  is a constant in the range  $[a/b, 1]$ . The elimination rule is ineffective for  $\delta < a/b$ , and the entire state-space will have be searched.  $\xi_{i,0}, \dots, \xi_{i,\mu(i)-1}$  are iid random variables as  $z_{i,0}, \dots, z_{i,\mu(i)-1}$  are iid random variables, and their distribution function for  $i < b$  is

$$F_{s,i}(x) = \begin{cases} 0 & x < \delta a \\ F_{s,i}(x/\delta) & \delta a \leq x \leq \delta b \\ 1 & \delta b < x \end{cases} \quad (2.11)$$

We will use  $N_{B(i,j)}$  to denote the number of nodes expanded by the best-first B&B algorithm in the subtree rooted at  $P_{ij}$ .  $P_{ij}$  will be expanded if and only if  $\xi_{ij} \leq z_0$ , so  $N_{B(i,j)}, 0 \leq j \leq \mu(i)-1$ , can be expressed as a following recurrence equation.

$$N_B(i,j) = \begin{cases} 1 + \sum_{i_1=0}^{j-1} \left[ \Pr\{\xi_{i+i_1,j+i_1} \leq z_0\} \cdot N_B(i+1, j+i_1) \right] & i < b-1 \\ 1 & i = b-1 \end{cases} \quad (2.12)$$

$\Pr\{\xi_{ij} \leq z_0\}$  and  $N_B(i,j)$  are independent. This permits  $E[N_B(i,j)]$ , the expected value of  $N_B(i,j)$ , to be expressed as

$$E[N_B(i,j)] = 1 + \sum_{i_1=0}^{j-1} \left[ E[\Pr\{\xi_{i+i_1,j+i_1} \leq z_0\}] \cdot E[N_B(i+1, j+i_1)] \right] \quad i < b-1 \quad (2.13)$$

$E[N_B(0,0)]$  is the expected number of nodes expanded by the best-first B&B algorithm as the root is always expanded.  $\xi_{i,j}, j = 0, \dots, \mu(i)-1$ , are iid random variables so  $E[N_B(i,j)] = \dots = E[N_B(j,\mu(j)-1)], 0 \leq i \leq b$ , and

$$E[N_B(i,0)] = 1 + s \cdot E[\Pr\{\xi_{i+1,0} \leq z_0\}] \cdot E[N_B(i+1,0)] \quad i < b-1 \quad (2.14)$$

where

$$E[\Pr\{\xi_{ij} \leq z_0\}] = \int_a^b F_{s,i}(x) f_{s,i}(x) dx$$

$$= \int_a^b F_{s,i}(x/\delta) f_{s,i}(x) dx + \int_b^b f_{s,i}(x) dx \quad (2.15)$$

From Eq. (2.9),  $z_0$  will have the following density function.

$$f_{z_0}(x) = \begin{cases} \mu(b) f_S(x) [1 - F_S(x)]^{\mu(b)-1} & a \leq x \leq b \\ 0 & \text{otherwise} \end{cases}$$

$$= \begin{cases} \frac{\mu(b) \beta e^{\beta a} (e^{\beta b} - e^{\beta x})^{\mu(b)-1}}{(e^{\beta b} - e^{\beta a})^{\mu(b)}} & a \leq x \leq b \\ 0 & \text{otherwise} \end{cases} \quad (2.16)$$

Substituting Eq's (2.10) and (2.16) into Eq. (2.15) yields

$$E[\Pr\{\xi_{ij} \leq z_0\}] = 1 - \int_a^b \frac{\frac{\mu(b) \beta e^{\beta a} (e^{\beta b} - e^{\beta x})^{\mu(b)-1} (e^{\beta b} - e^{\beta x+i})^{\mu(b)-1}}{(e^{\beta b} - e^{\beta a})^{\mu(b)}}}{dx} \quad (2.17)$$

Similarly, let  $N_B(i,j)$  be the number of nodes expanded in the subtree rooted at  $P_{ij}$  by the depth-first B&B algorithm. Given that node  $P_{ij}$  is expanded and  $t_{ij}$  is the value of the incumbent when  $P_{ij}$  is the first element in the active list. The depth-first B&B algorithm expands  $P_{ij}$  if and only if  $\xi_{ij} \leq t_{ij}$ . Thus,  $N_B(i,j)$  can be described by the recurrence equation:

$$N_B(i,j) = \begin{cases} 1 + \sum_{i_1=0}^{j-1} \left[ \Pr\{\xi_{i+i_1,j+i_1} \leq t_{i+i_1,j+i_1}\} \cdot N_B(i+1, j+i_1) \right] & i < b-1 \\ 1 & i = b-1 \end{cases} \quad (2.18)$$

Taking the expectations of Eq. (2.18) yields

$$E[N_B(i,j)] = 1 + \sum_{i_1=0}^{j-1} \left[ E[\Pr\{\xi_{i+i_1,j+i_1} \leq t_{i+i_1,j+i_1}\}] \cdot E[N_B(i+1, j+i_1)] \right].$$

$$\boxed{E[N_B(i,0)] = 1 + s \cdot E[\Pr\{\xi_{i+1,0} \leq z_0\}] \cdot E[N_B(i+1,0)] \quad i < b-1} \quad (2.19)$$

$$\begin{aligned}
 E[Pr\{\bar{s}_{i,j} \leq t_{i,j}\}] &= \int_a^b F_{a,i}(x) f_{t_{i,j}}(x) dx \\
 &= \int_a^b F_{a,i}(x/\delta) f_{t_{i,j}}(x) dx + \int_a^b f_{t_{i,j}}(x) dx
 \end{aligned} \tag{2.20}$$

The expected number of nodes expanded by a depth-first B&B will be  $E[D(0,0)]$ . If  $\mu(b-i)$  solution nodes have already been found when  $P_{i,j}$  becomes the first element in the active list. Thus,

$$t_{i,j} = \min \left\{ S_0, \dots, S_{j,\mu(b-i)-1} \right\} \tag{2.21}$$

and  $t_{i,j}$  has the following density function.

$$\begin{aligned}
 f_{t_{i,j}}(x) &= \begin{cases} j\mu(b-i)S(x)[1-F_S(x)]^{j\mu(b-i)-1} & a \leq x \leq b \\ 0 & \text{otherwise} \end{cases} \\
 &= \begin{cases} j\mu(b-i)\beta e^{\beta x}[e^{\beta b} - e^{\beta x}]^{j\mu(b-i)-1} & a \leq x \leq b \\ (e^{\beta b} - e^{\beta a})^{j\mu(b-i)} & \text{otherwise} \end{cases} \tag{2.22}
 \end{aligned}$$

Substituting Eq's (2.10) and (2.22) into Eq. (2.20) yields

$$E[Pr\{\bar{s}_{i,j} \leq t_{i,j}\}] = 1 - \int_a^b j\mu(b-i)\beta e^{\beta x}[e^{\beta b} - e^{\beta x}]^{j\mu(b-i)-1} [e^{\beta b} - e^{\beta x}]^{j\mu(b-i)-1} dx \tag{2.23}$$

Eq's (2.17) and (2.23) have no closed form solutions but can be transformed into recurrences. However,  $O(hs)$  recursions are required, and this number is very large for non-trivial  $h$  and  $s$ . Instead, Eq's (2.17) and (2.23) were evaluated by numerical integration.

Figure 2.3 shows the results of evaluating  $E[N_B(0,0)]$  for a tree with  $s=2$ ,  $h=15$ , and various values of  $\delta$  ( $a=10$ ,  $b=20$ ). Figure 2.4 shows the difference between  $E[N_B(0,0)]$  and  $E[N_E(0,0)]$  for the same tree. The best-first B&B algorithm always expands less nodes than the depth-first B&B algorithm. The difference between  $E[N_B(0,0)]$  and  $E[N_E(0,0)]$  first increases and then decreases as  $\delta$  increases from 0.55. A small  $\delta$  corresponds to a very inaccurate lower-bound function. Such a lower-bound function renders and the elimination rule ineffectual, so the best-first B&B process has to expand most of the problem space. This will be very similar to the portion of the problem space explored by the depth-first B&B process. A  $\delta$  close to 1 corresponds to a

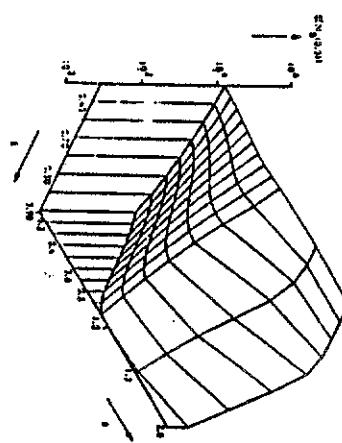


Figure 2.3  $E[N_B(0,0)]$  for a state-space tree with  $k=2$ ,  $h=15$ ,  $a=10$ ,  $b=20$ .

very accurate lower-bound function which is able to eliminate most nodes in the search space. For any particular problem, both the best-first and depth-first B&B algorithms use the same elimination rule. Hence, for a problem with a large  $\delta$ , the number of nodes expanded by the best-first B&B algorithm is close to the number of nodes expanded by the depth-first B&B algorithm. Between these two extremes, there are substantial disparities between the number of nodes expanded by the depth-first and best-first B&B algorithms. The regions for which the above behavior takes place depend on the value of  $\beta$ .

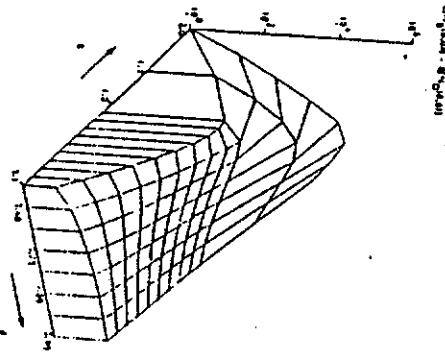
The above behavior can be illustrated by the integer-programming, vertex-cover, and 0/1-knapsack problems.

The vertex-cover problem involves a graph  $G = \{V, E\}$ , where  $V$  is the set of vertices and  $E$  is the set of edges. A vertex-cover for the graph  $G$  is a subset  $V' \subseteq V$  such that for each edge  $\{u, v\} \in E$ , at least one of  $u$  and  $v$  belongs to  $V'$ . The vertex-cover problem seeks the vertex-cover with the smallest number of vertices.

The 0/1-knapsack problem involves a set  $U$  of  $n$  items. The knapsack has a capacity  $K$ . Each item  $u \in U$  has a weight  $w(u)$  and a value  $v(u)$ . The 0/1-knapsack problem seeks the subset  $U' \subseteq U$  which maximizes  $\sum_{u \in U'} v(u)$  and satisfies  $\sum_{u \in U'} w(u) \leq K$ .

The lower-bound functions for the 0/1 knapsack and vertex-cover problems are very accurate greedy algorithms. The lower-bound function for the integer-programming problem is a linear program with the integrality constraints removed for all unassigned variables and is less accurate than the lower-bound functions of the vertex-cover and 0/1-knapsack problems. Table 2.1 shows some results from our simulations with these problems. For the knapsack and vertex-cover problems, the number of nodes expanded by the depth-first B&B algorithm was almost identical to that expanded by the best-first B&B algorithm. For integer-programming problems, the ratio of the number of nodes expanded in a depth-first B&B algorithm to that in a best-first B&B algorithm varies between 1.4 and 17.5.

Figure 2.4 Difference between  $E[N_D(0,0)]$  and  $E[N_B(0,0)]$  for a state-space tree with  $k=2$ ,  $b=15$ ,  $a=10$ ,  $b=20$ .



### 2.3 B&B Algorithms in a Paged Virtual-Memory Environment

In this section, we examine the overheads of the best-first B&B and depth-first B&B algorithms in a paged virtual memory environment. One implementation of the best-first B&B algorithm uses an unordered list of nodes and a priority queue of pointers to nodes. In each iteration, the best-first B&B

\* The priority queue can be implemented by a heap, which is a complete binary tree such that the value of each node is at least as small as that of its descendants.

Table 2.1

Number of nodes expanded by the best-first (BFS) and depth-first (DFS) B&B algorithms for randomly generated 30-object 0/1-knapsack, 80-node vertex-cover, and 20-by-20 integer programming problems.

Problem	Number of Nodes Expanded					
	0/1 Knapsack		Vertex-Cover		Integer-Program	
	BFS	DFS	BFS	DFS	BFS	DFS
1	1701	1731	4351	6205	1541	2169
2	5475	5519	5839	7290	2098	9425
3	6394	6442	8045	9750	2237	5358
4	1405	1474	5349	7134	1242	6752
5	1103	1145	4301	5186	594	1178
6	15309	15403	8253	8277	1398	3051
7	808	860	12140	12463	355	2887
8	7137	7202	7663	7647	1458	14485
9	18265	18396	12847	21892	240	3852
10	14257	14337	14104	15524	3475	7301

process gets a pointer to the node with the smallest lower-bound from the priority queue and deletes this node from the node list. After expanding the selected node, it inserts the newly-generated nodes into the node list and updates the priority queue by inserting pointers to the new nodes. A pointer to the node is much smaller in size than a node, so the pointer heap can be assumed to be resident in main memory and incurs no paging overhead.

For most problems, the best-first B&B algorithm generates few feasible solutions before finding the optimal solution. These few invocations of the elimination rule incurs little paging overhead. Thus, we can assume that the overhead due to the elimination rule is small compared to the overhead due to the expansion of nodes and may be ignored.

Let  $n_s$  be the number of nodes that can be expanded without exceeding the size of the main memory. The initial free memory space is  $(n_m n_p - 1)$  nodes as only the root of the B&B tree is present in main memory. Each node expansion removes one node from main memory but generates  $s$  new nodes; thus, each node expansion increases the number of nodes in main memory by  $s-1$ .

$$n_s = \left\lfloor \frac{n_m n_p - 1}{s-1} \right\rfloor \quad (2.24)$$

The node with the smallest lower-bound is equally likely to be in any of the pages occupied by the node list as the node list is not ordered. Consider the  $i$ -th node expansion. The  $\left( (i-1)+1 \right) / n_p$  active nodes occupy  $\left( (i-1)+1 \right) / n_p$  pages.  $n_m$  of these pages are resident in main memory; hence, the retrieval of the node with the smallest lower bound causes a page fault with probability  $(1 - \lfloor n_m / ((i-1)+1) / n_p \rfloor)$ . Hence, the expected number of page reads in a best-first B&B process is given by:

$$r_{\text{bf}} = \sum_{i=n_1+1}^{n_m} \left[ 1 - \frac{i(s-1) + 1}{n_p} \right] \quad n_{\text{bf}} > n_t \quad (2.25)$$

Each node expansion generates  $s$  new nodes. One new node can be inserted into the space vacated by the expanded node. Thus, the expected number of page writes is

$$w_{\text{bf}} = \left\lceil \frac{n_{\text{bf}}(s-1) + 1}{n_p} \right\rceil - n_m \quad n_{\text{bf}} > n_t \quad (2.26)$$

If there are sufficient secondary memory buffers, the best-first B&B process can resume as soon as the contents of the page selected for replacement has been copied into a secondary memory buffer. The actual page write is overlapped with processing. In contrast, the best-first B&B process has to be suspended until the page read has been completed.

Let  $h_c$  be the CPU time required to expand a node, and let  $h_d$  be the average page-read time. The overhead incurred by a best-first search,  $D_{\text{bf}}$ , can be approximated by the summation of the total page-read times and the maximum of the CPU times and the page write times.

$$D_{\text{bf}} = r_{\text{bf}} h_c + \max(w_{\text{bf}}, h_d, n_{\text{bf}} h_c) \quad (2.27)$$

From Eq. (2.25),  $r_{\text{bf}} \approx w_{\text{bf}} \approx n_{\text{bf}}$  when  $n_{\text{bf}} \gg n_t$ . In this case, the best-first B&B algorithm will be more efficient than the depth-first B&B algorithm if  $\Xi$  exceeds  $\Xi_e$ , where

$$\Xi_e = \frac{\max\{h_c, h_d/n_p\} + h_d/n_p}{h_c} \quad (2.28)$$

Similarly, the depth-first B&B algorithm will be more time efficient than the best-first B&B algorithm when  $\Xi$  is less than  $\Xi_e$ . The vertex-cover and 0/1-knapsack problems are two problems for which this condition is likely to be true. Thus, it is possible for the best-first B&B algorithm to be less efficient than the depth-first B&B algorithm in a two-level memory system.

The minimum time required by any B&B algorithm to solve a problem instance is  $b_c \cdot n_{\text{bf}}$ . The depth-first B&B algorithm approaches this optimal performance when  $n_{\text{bf}} \ll n_{\text{bf}}$ . The best-first B&B algorithm should be chosen only when  $\Xi \gg \Xi_e$ . Similarly, a depth-first search should be used when  $\Xi < \Xi_e$ . For problems characterized by intermediate values of  $\Xi$ , the choice between the depth-first and best-first searches is less clear. Either a depth-first

or a best-first search may be used, although neither is particularly suitable. In the next two sections, we study alternatives which may be used for problems with intermediate values of  $\Xi$ .

#### 2.4 A Virtual Memory System for the Best-First B&B Algorithm

In work prior to the work of this thesis, we have designed a virtual memory system which is tailored to the characteristics of the best-first B&B algorithm [YuW83, Yu83]. Figure 2.5 depicts the virtual-memory system. Nodes in secondary memory are organized as a B<sup>+</sup>-tree.<sup>\*\*</sup> Each leaf in a B<sup>+</sup>-tree is a page, and the nodes reside in the leaves of the B<sup>+</sup>-tree. Non-terminal nodes in the B<sup>+</sup>-tree are pointer nodes that are much smaller than the leaf nodes, so the memory management system keeps most of the pointer tree in main memory. This significantly reduces the number of secondary-memory accesses required to access a node in secondary storage.

The main memory contains a partial list of nodes (m-list) and a heap of pointers to nodes in the m-list. The best-first B&B process inserts newly generated nodes into the m-list. Whenever the m-list is full, the memory management system moves part of the m-list to pages in the B<sup>+</sup>-tree. The actual transfer of nodes will be overlapped with computation if there are sufficient secondary-memory buffers.

The best-first B&B process always expands the active node with the smallest lower-bound. This node may either be in the m-list or the first page of the B<sup>+</sup>-tree. The overhead of accessing the active node with the smallest lower-bound will be reduced by keeping part or all of the first B<sup>+</sup>-tree page in main memory (s-buffer).

<sup>\*\*</sup> A B-tree of order  $b$  is a search tree that is either empty or satisfies the following properties [Com79]: (i) the root node has at least two children; (ii) each node contains at least  $b/2$  keys and  $b+1$  pointers; and (iii) each node contains at most  $2b$  keys and  $2b+1$  pointers. A B<sup>+</sup>-tree is a variant of the B-tree in which all records reside in the leaves. The upper levels are organized as a B-tree and serve as an index to locate a record. The leaf nodes are linked from left to right for easy sequential processing.



empty block in the bit map and writing the block when the disk head is properly positioned.

A suitable replacement algorithm for this virtual-memory system is the *Buck-Replacement Policy*,  $S(\phi)$ , which replaces the  $\phi$  nodes with the  $\phi$ -largest lower bounds. A value of  $\phi$  between 0.7  $n_m$  and 0.9  $n_m$  has been found to be suitable for integer-programming and vertex-cover problems. The page size should be between 0.85  $n_m$  and 0.85  $n_m$ .

### 2.5 The Modified B&B Algorithm

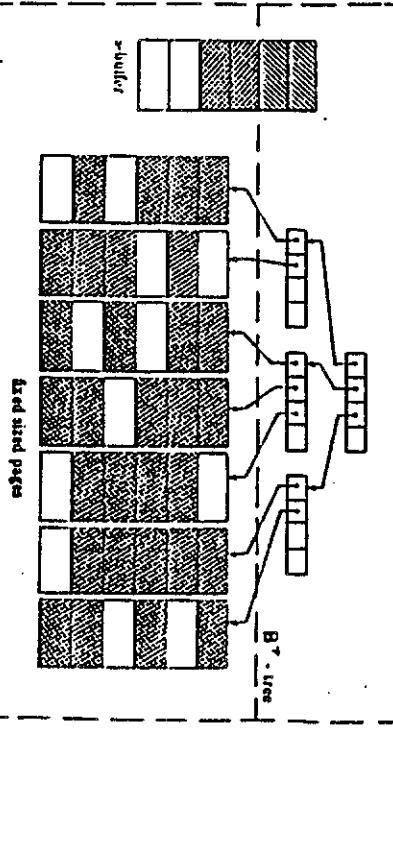
The modified B&B algorithm partitions the range of possible lower-bounds into  $n$  disjoint regions,  $\{[u_0, u_1], [u_1, u_2], \dots, [u_{n-1}, u_n]\}$  (Figure 2.6). Normally  $u_0 = g_0$  and  $u_n = g_{opt}$ , where  $g_0$  is the lower-bound of the initial problem and  $g_{opt}$  is an estimate of the optimal solution value. The set of active nodes in each region is usually large, so some of the nodes may have to be stored in secondary memory. Nodes in each region  $[u_i, u_{i+1}], i = 0, \dots, n-1$ , are not ordered by lower-bounds but are kept in a LIFO stack  $S_i$  because of the high overhead of accessing nodes in secondary storage.

The implementation of the modified B&B algorithm depends on whether it uses the file system or the virtual-memory system. We will describe the version for a file system and note the needed changes for a virtual-memory system.

#### 2.5.1 Modified B&B Algorithm for a File System

Each stack  $S_i$  consists of a secondary-memory file,  $F_i$ , and a main-memory buffer,  $B_i$ . Buffer  $B_i$  is the top portion of  $S_i$ . The main-memory space allocated to the storage of nodes is organized as a pool of fixed-sized frames. Each buffer consists of a number of frames. Nodes falling within the range of  $S_i$  are inserted into  $B_i$ . If the frames allocated to  $B_i$  are all full, the modified B&B process will request a new empty frame from the frame pool manager. The frame manager will always try to satisfy this request from its frame pool but if the frame pool is empty, it has to reclaim the oldest full frame from the buffers. Assuming sufficient secondary-memory buffers, the overhead of reclaiming a buffer is the time to copy the contents of the reclaimed frame into a secondary memory buffer. The actual transfer of nodes from the secondary-memory buffer to secondary memory is completely overlapped with computation.

Figure 2.5 A virtual-memory system for the best-first B&B algorithm.



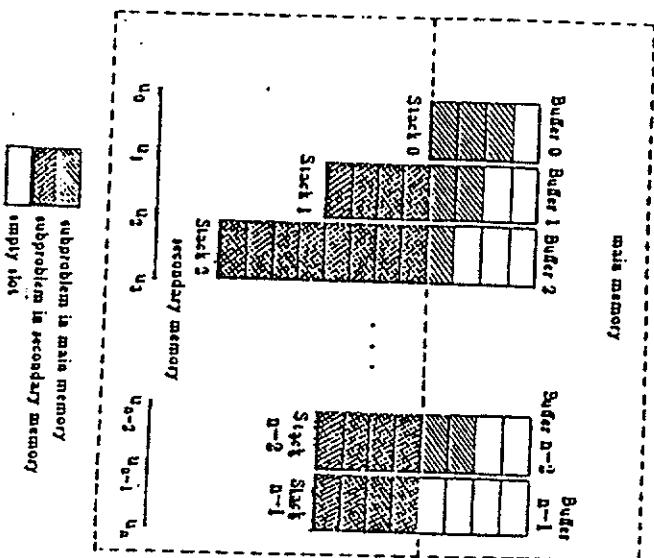


Figure 2.6 Data structure for the modified B&B algorithm.

The modified B&B algorithm with one stack is equivalent to the depth-first B&B; the modified B&B algorithm with infinitely many stacks is equivalent to a best-first B&B algorithm. As the number of stacks increases, the number of expanded nodes generally decreases, while the secondary-memory overhead generally increases. The stack boundaries and the number of stacks should be chosen to maximize the overlap between computation and secondary-memory accesses.

Let the current stack be the stack that contains the active node with the smallest lower-bound, and let the optimal stack be the stack containing the optimal solution. A stack is selected when it becomes the current stack. Similar labels apply to the corresponding buffers and files. Stack  $S_j$  is defined to be smaller than stack  $S_k$  if  $j < k$ , and larger if  $j > k$ . Similar relationships hold for the buffers  $B_j$  and  $B_k$ .

The new data structure entail modifications to the selection and elimination rules. The modified selection rule selects the top element of the smallest non-empty buffer. This may or may not be current buffer. An empty current buffer triggers a transfer of nodes from the current file into the current buffer. The modified selection rule enables nodes to be expanded while this transfer is being carried out and avoids the long wasteful wait incurred by strict adherence to the failible lower-bound heuristic.

Let the incumbent be  $z$ , and let the allowance function be  $\epsilon(z)$ . Suppose that  $(z - \epsilon(z))$  lies in the range of stack  $S_i$ . The modified lower-bound elimination rule eliminates all active stacks  $S_i$  with  $i > k$  that have not been previously eliminated. Nodes in  $F_k$  with lower-bounds greater than  $(z - \epsilon(z))$  are not eliminated immediately because of the high overhead of accessing them in secondary memory. Furthermore, it is possible that a future incumbent will eliminate stack  $S_k$ . Thus, the culling of nodes in  $F_k$  is postponed until they have been read into main memory during the expansion of  $S_k$ . This necessitates an additional lower-bound test on each selected node before it is expanded. Figure 2.7 summarizes the modified B&B algorithm.

The modified B&B algorithm expands at least as many nodes as the corresponding best-first B&B algorithm. The depth-first expansion of the solution stack,  $S_r$ , results in the expansion of some nodes in  $S_r$  with lower-bounds greater than the optimal solution value. In addition, the modified selection rule results in the expansion of some nodes in stacks  $S_{r+1}, \dots, S_{n-1}$ . The best-first B&B algorithm never expands these nodes.

```

procedure modified B&B algorithm;
  /* Range of possible lower-bounds is divided into regions, each of which is occupied by a stack  $S_i$ ,  $0 \leq i \leq n-1$ .  $S_i$  has a main-memory buffer  $B_i$  and a secondary-memory file  $F_i$ . */
  |
  |   incumbent =  $\infty$ ;  $B_0 = \{P_0\}$ ;  $S_0 = \emptyset$ ,  $i = 1, \dots, n-1$ ;
  |   while there are non-empty stacks do [
  |     let  $S_k$  be the current stack;
  |     if  $B_k$  is empty and no read from  $F_k$  to  $B_k$  is outstanding then
  |       initiate read from  $F_k$  to  $B_k$ ;
  |       if all buffers are empty then
  |         wait until read from any buffer is completed;
  |         apply modified selection rule;
  |       if (lower-bound of selected node) < ( $\tau - \epsilon(z)$ ) then [
  |         expand selected node;
  |         for  $P \in \{\text{expanded nodes}\}$  do [
  |           if  $P$  is a feasible solution then [
  |             update incumbent;
  |             apply modified elimination rule];
  |           else if ( $\epsilon(P)$ ) < ( $\tau - \epsilon(z)$ ) then [
  |             if  $B_i$  is full then [
  |               if frame pool is empty then
  |                 initiate write from  $B_i$  to  $F_i$ ;
  |                 get frame from frame pool];
  |                 push  $P$  into the corresponding  $B_i$ ];
  |             ];
  |           ];
  |         ];
  |       ];
  |     ];
  |   ];
  |

```

Figure 2.7 The modified B&B algorithm.

Before we present some strategies for choosing the stack boundaries, we will develop analytical expressions for the expected overhead for the modified B&B process. These analytical expressions form the basis of one scheme for choosing stack boundaries.

### 2.5.1.1 Expected Overhead for Real-Valued Problems

This analysis uses some results from the wall model of the best-first B&B algorithm. The wall model was developed prior to the work of this thesis and has been reported elsewhere [Yu83, WaY82, WaY85]. We will only include the relevant aspects of the model in this thesis; specifically, the expected numbers of nodes which have already been expanded and generated when the best-first B&B process expands a node with a specified lower-bound. As shown in Figure 2.8, this model consists of two moving walls. The left wall represents the lower-bound of the node being expanded. The back wall represents the incumbent. These values are normalized by the lower-bound of the initial problem,  $\epsilon(P_0) = g_0$ . For any node  $P_j$ , let  $\Delta(P_j) = \epsilon(P_j) - g_0$ . Initially, the left wall is at  $\Delta(P_0) = 0$ , and the back wall is at infinity. The left wall always moves to the right as the best-first B&B process always expands the active node with the smallest lower-bound which always increases the smallest lower-bound value. Moreover, the back wall always moves to the left as the incumbent can only decrease in value. Thus, successive node expansions cause the two walls to approach each other. The best-first B&B process terminates when the two walls meet.

The difference between the lower-bound of a node and the lower-bound of its parent is assumed to be an exponential random variable with density function,  $f_G(y) = \lambda e^{-\lambda y}$ . In addition, each node expansion is assumed to generate  $s$  new nodes.

Let  $N(x)$  be the number of expanded nodes when the front wall is at  $x$ . The best-first B&B process must have already expanded all nodes with lower-bounds less than  $x$  before it expands a node with lower bound  $x$ . Thus,  $N(x)$  may also be interpreted as the number of nodes in the search tree with lower-bounds not exceeding  $x$ . Let  $P_j$  be a child of  $P_i$ . Suppose that the increase in lower-bounds during this node expansion is known to be  $y$ . Then,  $\Delta(P_j) \leq x$  if and only if  $\Delta(P_i) \leq x-y$ . This together with the assumption that each node expansion generates  $s$  new nodes leads to a renewal equation for the expectation of  $N(x)$ .

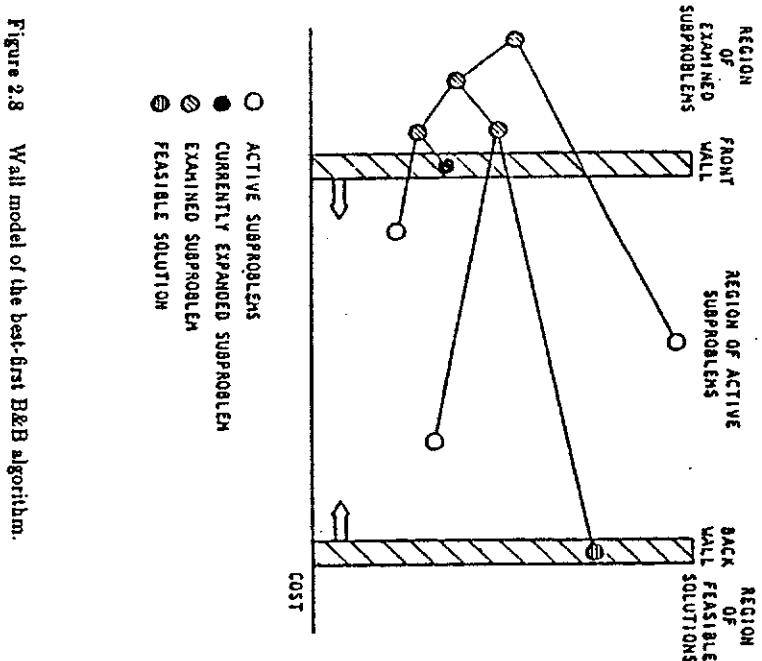


Figure 2.8 Wall model of the best-first B&B algorithm.

$$E[N(x)] = 1 + \int_{y=0}^x s E[N(x-y)] f_G(y) dy \quad (2.29)$$

Assuming a reasonably large  $x$ ,  $f_G(y) \approx 0$  for any  $y > x$  as  $f_G(y) \rightarrow 0$  as  $y \rightarrow \infty$ . This assumption simplifies Eq. (2.29) to

$$E[N(x)] = 1 + s \int_{y=0}^{\infty} E[N(x-y)] f_G(y) dy \quad (2.30)$$

Eq. (2.30) has a generic solution with the form  $E[N(x)] = \left[ k e^{mx} - \frac{1}{s-1} \right]$

Substituting the generic solution for  $E[N(x)]$  in Eq. (2.30) results in the identity:

$$k e^{mx} - \frac{1}{s-1} = 1 + s \int_{y=0}^{\infty} \left[ k e^{m(s-y)} - \frac{1}{s-1} \right] f_G(y) dy$$

$$\text{or} \quad 1 = s \int_{y=0}^{\infty} e^{-my} f_G(y) dy \quad (2.31)$$

$$\text{or} \quad m = \lambda(s-1)$$

Solution of the constant  $k$  requires the use of the boundary condition  $E[N(x_0)] = N_0$ , where  $x_0 + \epsilon_0$  is the lower-bound of the  $N_0$ -th node to be expanded. The value of  $N_0$  should be chosen to eliminate any initial transients in the density function  $f_G(y)$ . Setting  $x$  to 0 in the generic solution yields  $k = \frac{N_0(s-1)}{s-1} + \frac{1}{s-1} e^{-ms_0}$ . Therefore,

$$E[N(x)] = \frac{N_0(s-1) + 1}{s-1} e^{\lambda(s-1)m - m s_0} - \frac{1}{s-1} \quad (2.33)$$

Eq. (2.33) is an approximation when  $s$  is not an integer. Let  $M(x)$  be the number of generated nodes when the front wall is at  $x$ . Each node expansion generates  $s$  new nodes; so the expectation of  $M(x)$  is

$$E[M(x)] = s E[N(x)] + 1 \quad (2.34)$$

Recall that  $N(x)$  may be interpreted as the number of nodes in the search space with lower-bounds not exceeding  $x$ . The expressions for  $E[N(x)]$  and  $E[M(x)]$  will be used in the analysis of the modified B&B algorithm in this context. Thus,  $E[N(u)]$  represents the number of nodes expanded during the expansion of stacks  $S_1, \dots, S_{t-1}$ , and  $E[M(u)]$  is the expectation of the corresponding number of generated nodes.

Let  $Z_1, \dots, Z_{N(x)}$  be iid random variables with distribution  $F_s(z) | x, x \geq z$ . The assumption that the  $N(x)$  expanded nodes are independent of each other implies the stochastic equivalence of  $Z_i$  and  $\Delta(P_i)$ . Let  $z_i$  be the value of  $Z_i$  and assume, without any loss of generality, that  $z_1 \leq \dots \leq z_{N(x)}$ . For any  $z < x$ ,  $E[N(z)]$  is the expected number of expanded nodes with lower-bounds not exceeding  $z$ . Stochastically equivalent random variables have identical means.

$$N(x)F_s(z) | x, x \geq z = E[N(z)]. \quad (2.35)$$

$E[N(z)]$  in Eq. (2.33) may be replaced by its expression in Eq. (2.35). This replacement leads to

$$F_s(z | x, x \geq z) = \frac{1}{(s-1)N(x)} \left[ s e^{\lambda(u_1 - z)} - 1 \right] \quad (2.36)$$

The density function,  $f_s(z | x, x \geq z)$ , is the differential of  $F_s(z | x, x \geq z)$  with respect to  $z$ .

$$f_s(z | x, x \geq z) = \frac{1}{N(x)} s \lambda e^{\lambda(u_1 - z)} \quad (2.37)$$

Let  $\Pi_j, j = 1, \dots, N(x)$  be the  $N(x)$  expanded nodes and  $\Psi_k, k = 1, \dots, M(x)$  the  $M(x)$  generated nodes when the left wall is at  $x$ . Suppose that  $\Pi_j$  is  $\Psi_k$ 's parent. If  $l_k$  is the difference between the lower-bounds of  $\Psi_k$  and  $\Pi_j$  then

$$\Delta(\Psi_k) = \Delta(\Pi_j) + l_k = z_j + l_k \quad (2.38)$$

Eq. (2.38) implies that  $\Delta(\Psi_1), \dots, \Delta(\Psi_{M(x)})$  are iid random variables as both  $Z_1, \dots, Z_{N(x)}$  and  $I_1, \dots, I_{M(x)}$  are sets of iid random variables. Let  $\gamma_i(x)$  be the probability that  $\Delta(\Psi_k), k = 1, \dots, M(x)$  lies within the range of Stack  $i$ , (that is,  $(u_i - \varepsilon_0) < \Delta(\Psi_k) \leq (u_{i+1} - \varepsilon_0)$ ) when the front wall is at  $x$ . Suppose that  $Z_j = z$ .  $\Delta(\Psi_k)$  lies within  $[u_i - \varepsilon_0, u_{i+1} - \varepsilon_0]$  if  $I_k \in [u_i - \varepsilon_0 - z, u_{i+1} - \varepsilon_0 - z]$ . Thus, the value of  $\gamma_i(x)$  given  $Z_j$  will be

$$\begin{aligned} (\gamma_i(x) | Z_j = z) &= \int_{y=u_i-\varepsilon_0-z}^{u_{i+1}-\varepsilon_0-z} \lambda e^{-\lambda y} dy \\ &= \left[ e^{-\lambda(u_i-\varepsilon_0)} - e^{-\lambda(u_{i+1}-\varepsilon_0)} \right] e^{\lambda z} \end{aligned} \quad (2.39)$$

Eq. (2.39) may be unconditioned by using Eq. (2.37). Hence

$$\gamma_i(x) = \int_{z=0}^x \left[ e^{-\lambda(u_i-\varepsilon_0)} - e^{-\lambda(u_{i+1}-\varepsilon_0)} \right] \frac{s \lambda e^{\lambda z}}{N(x)} dz \quad (2.40)$$

There are  $M(x)$  generated nodes. Each generated node has probability  $\gamma_i(x)$  of falling within the range of Stack  $S_i$ . The stochastic equivalence of  $Z_i$  and  $\Delta(P_i)$  implies that the expected number of nodes in the region  $[u_i - \varepsilon_0, u_{i+1} - \varepsilon_0]$  is

$$E[W_i(x)] = E[M(x)] \gamma_i(x) \approx s E[N'(x)] \gamma_i(x)$$

$$\approx \left[ e^{-\lambda(u_i-\varepsilon_0)} - e^{-\lambda(u_{i+1}-\varepsilon_0)} \right] s e^{\lambda x} \quad (2.41)$$

Let  $g_{opt}$  be the optimal solution value and let  $S_r$  be the solution stack (i.e.,  $u_r < g_{opt} \leq u_{r+1}$ ). In addition, let  $E[q | r]$  be the expected number of nodes expanded before the modified B&B process finds the optimal solution in  $S_r$ . It is extremely difficult to model the LIFO expansion of  $S_r$  as it requires the lower-bounds of the individual nodes in  $S_r$ . Moreover,  $g_{opt}$  is not known until the problem is solved. Instead, we use the simplification that all nodes in  $S_r$  are expanded. This implies

$$E[q | r] = E[N(u_{r+1} - \varepsilon_0)] \approx \frac{s}{s-1} e^{-\lambda(u_{r+1} - \varepsilon_0)} \quad (2.42)$$

The optimal solution value is assumed to be uniformly distributed over the range  $[g_0, g_{opt}]$ . The actual distribution depends on the particular problem instance, and the assumed distribution is selected for reasons of mathematical tractability. With this assumption, the probability that  $S_r$  is the solution stack,  $\pi_r$ , for a problem with real-valued solutions will be given by

$$\pi_r = \frac{u_{r+1} - u_r}{g_{opt} - g_0} \quad (2.43)$$

For a problem with integer-valued solutions,  $\pi_r$  is

$$\pi_r = \frac{\left| \begin{array}{|c|c|} \hline u_{r+1} & - u_r \\ \hline \end{array} \right|}{\left| \begin{array}{|c|c|} \hline g_{opt} & - g_0 \\ \hline \end{array} \right|} \quad (2.44)$$

The modified B&B process expands  $E[g | r]$  nodes when the optimal solution is in  $S_r$ . The probability of this event is  $\pi_r$ . Let  $b_r$  be the mean computational

time required to expand a node. The expected computational overhead,  $E[d_i]$ , is

$$E[d_i] = \sum_{r=0}^{n-1} E[d_i|r] b_i \pi_r \\ = \sum_{r=0}^{n-1} \left[ \frac{s}{s-1} e^{\lambda(r-1)(u_{i+r}-\omega)} \right] b_i \pi_r \quad (2.45)$$

Only the secondary-memory overhead due to the transfer of nodes will be considered in this analysis. The main memory allocated to node storage is divided into  $n_m$  frames. Each frame holds  $n_p$  nodes.  $n_0$  frames are reserved for the depth-first expansion of the current stack.

Let  $m_i$  be the number of nodes in  $S_i$  when it is either selected or eliminated, and let  $E[m_i|r]$  be its expectation when  $S_i$  is the solution stack. If  $1 \leq i \leq r$  then  $m_i$  is the number of nodes in  $S_i$  when it becomes the current stack,  $E[m_i|r] \approx E[W_i|u_i-\omega]$  if few nodes in  $S_i$  have been expanded before  $S_i$  becomes the current stack. If  $(r+1) \leq i \leq (n-1)$  then  $m_i$  is the number of nodes in  $S_i$  when  $S_i$  is eliminated,  $E[m_i|r] \approx E[W_i|u_{i+r}-\omega]$  if no stacks are eliminated prior to the generation of the optimal solution. This is a good approximation when the modified B&B process generates few feasible solutions before it finds the optimal solution, as in the integer-programming and vertex-cover problems. In short,

$$E[m_i|r] \approx \begin{cases} E[W_i|u_i-\omega] & 1 \leq i \leq r \\ E[W_i|u_{i+r}-\omega] & (r+1) \leq i \leq (n-1) \end{cases} \\ \approx \begin{cases} \left[ e^{-\lambda(u_i-\omega)} - e^{-\lambda(u_{i+r}-\omega)} \right] s e^{\lambda(u_i-\omega)} & 1 \leq i \leq r \\ \left[ e^{-\lambda(u_{i+r}-\omega)} - e^{-\lambda(u_{i+r}-\omega)} \right] s e^{\lambda(u_{i+r}-\omega)} & (r+1) \leq i \leq (n-1) \end{cases} \quad (2.46)$$

$n_0$  frames are always reserved for the depth-first expansion of the current stack. A correct choice of the value of  $n_0$  ensures that there is no overflow of nodes from the current buffer to the current file. Furthermore, the lower-bound test will remove all nodes falling within the range of an eliminated stack so no nodes will ever be written to an eliminated file. Hence, no more than  $n_i$  nodes from Stack  $S_i$  will be written to the secondary memory. Let  $b_i$  frames be required to hold  $n_i$  nodes. The expectation of  $b_i$  given the solution stack  $S_i$  is

$$E[b_i|r] = \left\lceil \frac{E[m_i|r]}{n_p} \right\rceil \\ \approx \left\lceil \frac{\left[ \left( e^{-\lambda(u_i-\omega)} - e^{-\lambda(u_{i+r}-\omega)} \right) s e^{\lambda(u_i-\omega)} \right]}{n_p} \right\rceil \quad 1 \leq i \leq r \\ \left\lceil \frac{\left[ \left( e^{-\lambda(u_{i+r}-\omega)} - e^{-\lambda(u_{i+r}-\omega)} \right) s e^{\lambda(u_{i+r}-\omega)} \right]}{n_p} \right\rceil \quad (r+1) \leq i \leq (n-1) \quad (2.47)$$

Let  $c_i$  be the number of frames allocated to  $B_i$  when  $S_i$  either becomes the current stack or is eliminated. If no stack has been previously eliminated, the number of buffers holding frames is  $(n-i-1)$  when  $1 \leq i \leq r$  and  $(n-r-1)$  when  $(r+1) \leq i \leq (n-1)$ . The frame-pool manager always reclaims the oldest full frames so the number of frames allocated to each buffer is proportional to the total number of nodes in each stack. Thus, the expected value of  $c_i$  when  $S_i$  is the solution stack is

$$E[c_i|r] = \begin{cases} \left[ \left( e^{-\lambda(u_i-\omega)} - e^{-\lambda(u_{i+r}-\omega)} \right) (n_m - n_0) \right] & 1 \leq i \leq r \\ \left[ \left( e^{-\lambda(u_{i+r}-\omega)} - e^{-\lambda(u_{i+r}-\omega)} \right) (n_m - n_0) \right] \sum_{j=\min(i,r+1)}^{n-1} E[b_j|r] & (r+1) \leq i \leq (n-1) \end{cases} \quad (2.48)$$

The modified B&B process writes at most  $(b_i - c_i)$  frames to file F. These frames will be read back into the main memory if  $S_i$  is selected but will be discarded if  $S_i$  is eliminated. Let  $b_d$  be the overhead of transferring one frame of nodes between main memory and secondary memory. The expected secondary-memory overhead for  $S_i$ , given the solution stack  $S_i$ , will be the product of  $b_d$  and the number of frames transferred between main memory and secondary memory.

$$E[d_i|r] = \begin{cases} 0 & i=0 \\ 2(E[b_i|r] - E[c_i|r]) b_d & 1 \leq i \leq r \\ (E[b_i|r] - E[c_i|r]) b_d & (r+1) \leq i \leq (n-1) \end{cases} \quad (2.49)$$

Unconditioning for the solution stack  $S_i$  yields

$$\mathbb{E}[t_d] = \sum_{i=0}^{n-1} \sum_{r=0}^{s-1} \mathbb{E}[d_i|r] \pi_r \quad (2.50)$$

The expected overhead of an algorithm in our computer model is the larger of the expected computation time and the secondary memory overhead. The expected overhead for the boundary assignment  $\{u_0, \dots, u_n\}$  is

$$\mathbb{E}[D(u)] = \max \{\mathbb{E}[t_c], \mathbb{E}[t_d]\} \quad (2.51)$$

### 2.5.1.2 Expected Overhead for Discrete-Valued Problems

The analysis for discrete-valued problems is similar to that for real-valued problems. We will only present the highlights of the analysis in this section. The increases in lower-bounds during a node expansion are now governed by a discrete distribution, the negative-binomial distribution,

$$f(y|r,p) = \begin{cases} r+y-1 \\ y \end{cases} p^y (1-p)^{r-y} \quad y = 0, 1, 2, \dots \quad (2.52)$$

$r$  and  $p$  are parameters to be determined from the problem under consideration. This distribution has been verified experimentally for the vertex-cover problem.

Using the arguments and notations of Section 2.5.1.1, we can derive an approximate renewal equation for the expected number of expanded nodes when the front wall is at  $x$ .

$$\mathbb{E}[N(x)] = 1 + s \sum_{y=0}^{\infty} \mathbb{E}[N(x-y)] \times f(y|x,p) \quad x = 0, 1, \dots \quad (2.53)$$

The solution to Eq. (2.53) is:

$$\mathbb{E}[N(x)] = \frac{N(s-1)+1}{s-1} \left[ \frac{q}{1-ps/l_r} \right]^{x-r+1} - \frac{1}{s-1} \quad x = 0, 1, \dots \quad (2.54)$$

As before, the expectation of the corresponding number of generated nodes is

$$\mathbb{E}[M(x)] \approx s \mathbb{E}[N(x)] + 1 \quad (2.55)$$

Eq. (2.54) also enables us to derive an expression for the expected computational overhead.

$$\mathbb{E}[t_d] = \sum_{r=0}^{n-1} \frac{N(s-1)+1}{s-1} \left[ \frac{q}{1-ps/l_r} \right]^{u_{r+1}-x_0} b_c \pi_r \quad (2.56)$$

where  $b_c$  is the mean computational time to expand a node, and  $l_r$  is the

probability that stack  $S_r$  is the optimal stack.

As mentioned in Section 2.5.1.1,  $N(x)$  also represents the number of nodes in the search space with lower-bounds not exceeding  $x$ . The difference between  $N(y)$  and  $N(y-1)$  is the number of nodes with lower-bound  $y$ . This results in the following approximate expression for  $\mathbb{E}[L(y)]$ , the expected number of nodes with lower-bound  $y$ .

$$\mathbb{E}[L(y)] \approx \begin{cases} \mathbb{E}[N'(y)] - \mathbb{E}[N'(y-1)] & y=1, \dots, x \\ \mathbb{E}[N'(0)] & y=0 \end{cases} \quad (2.57)$$

A node generated by the expansion of a node with lower-bound  $z$  lies within Stack  $S_i$  if the increase in lower-bound lies within  $[u_i^{-z}, u_{i+1}^{-z}]$ . The increase in lower-bound during a node expansion is an iid negative binomial random variable, so the probability of this event is  $\sum_{v=u_i-z}^{u_{i+1}-z} f(v|r,p)$ . Thus, the

expected number of generated nodes in the region  $[u_i^{-z}, u_{i+1}^{-z}]$  when the front wall is at  $x$  can be expressed as

$$\mathbb{E}[W_i(x)] = \sum_{y=0}^x L(y) \left[ \sum_{z=u_i^{-y}}^{u_{i+1}-y-1} f(z|r,p) \right] \quad (2.58)$$

As in Section 2.5.1.1,  $m_i$  denotes the number of nodes in stack  $S_i$  when it is either eliminated or selected. Using Eq. (2.45), we can show that the expectation of  $m_i$  given the optimal stack,  $S_o$ , is

$$\mathbb{E}[m_i|r] \approx \begin{cases} \mathbb{E}[W_i(u_i^{-z_0})] & 1 \leq i \leq r \\ \mathbb{E}[W_i(u_{r+1}^{-z_0})] & (r+1) \leq i \leq (n-1) \end{cases}$$

$$\begin{aligned} & \approx \begin{cases} \sum_{y=0}^{u_i^{-z_0}-1} L(y) \left[ \sum_{z=u_i^{-y}}^{u_{i+1}-y-1} f(z|r,p) \right] & 1 \leq i \leq r \\ \sum_{y=0}^{u_{r+1}^{-z_0}-1} L(y) \left[ \sum_{z=u_i^{-y}}^{u_{i+1}-y-1} f(z|r,p) \right] & (r+1) \leq i \leq (n-1) \end{cases} \end{aligned} \quad (2.59)$$

The number of frames required to hold these  $\mathbb{E}[m_i|r]$  nodes will be given by

$$\mathbb{E}[b_i|r] = \left\lceil \frac{\mathbb{E}[m_i|r]}{n_p} \right\rceil$$

$$\begin{aligned}
 & \approx \left[ \frac{\sum_{y=0}^{u_i - b_p^{-1}} L(y) \left[ \sum_{s=u_i y}^{u_{i+1} - b_p^{-1}} l(z; r, p) \right]}{n_p} \right] \\
 & \quad \text{if } i \leq i \leq r \\
 & \quad (r+1) \leq i \leq (n-1) \\
 & \left[ \frac{\sum_{y=0}^{u_{i+1} - b_p^{-1}} L(y) \left[ \sum_{s=u_i y}^{u_{i+1} - b_p^{-1}} l(z; r, p) \right]}{n_p} \right]
 \end{aligned} \tag{2.60}$$

Using the arguments, assumptions, and notation of Section 2.5.1.1, we can then determine the number of frames allocated by each stack when it is either selected or eliminated.

$$\begin{aligned}
 E[b_i | r] = & \left[ \frac{\sum_{y=0}^{u_{i+1} - b_p^{-1}} L(y) \left[ \sum_{s=u_i y}^{u_{i+1} - b_p^{-1}} l(z; r, p) \right]}{\sum_{y=0}^{u_{i+1} - b_p^{-1}} L(y) \left[ \sum_{s=u_i y}^{u_{i+1} - b_p^{-1}} l(z; r, p) \right]} \right] \sum_{i=\min(i, r+1)}^{n-1} E[b_i | r] \geq n_m - n_o
 \end{aligned} \tag{2.61}$$

$$\sum_{i=\min(i, r+1)}^{n-1} E[b_i | r] < n_m - n_o$$

As before, the expected secondary memory overhead for Stack  $S_i$  due to the transfer of nodes when  $S_i$  is the solution stack will be given by

$$E[d_i | r] = \begin{cases} 0 & i = 0 \\ 2(E[b_i | r] - E[c_i | r])b_d & 1 \leq i \leq r \\ (E[b_i | r] - E[c_i | r])b_d & (r+1) \leq i \leq (n-1) \end{cases} \tag{2.62}$$

where  $b_d$  is the overhead for each page read. Only the secondary memory overhead due to the transfer of nodes is considered in this analysis, so the modified B&B process incurs a secondary memory overhead given by

$$E[d] = \sum_{i=0}^{n-1} \sum_{r=0}^{n-1} E[d_i | r] \pi_r \tag{2.63}$$

The expected overhead for  $n$  stacks and a boundary assignment  $\{u_0, \dots, u_n\}$  is

$$E[D(n)] = \max(E[t_d], E[t_d]) \tag{2.64}$$

### 2.5.1.3 Choosing Stack Boundaries

The analytical expressions for  $E[D(n)]$  form the basis of one scheme for choosing stack boundaries. This analytical strategy finds the assignment that minimizes  $E[D(n)]$  for a given  $n$  by enumerating  $E[D(n)]$  over feasible assignments. analytical scheme is the equal-width strategy which is a heuristic method that divides the range of possible lower-bounds into equal-width regions.

We have carried out extensive simulations of the analytical (AA) and equal-width (EWA) strategies using integer-programming and vertex-cover problems. Table 2.2 shows some of these results; the values for the overhead have been normalized with  $b_d$ , the mean computational time required to expand a node. These simulations suggest that the analytical and equal-width strategies have generally similar performances.

These extensive simulations also suggest the use of either three or four stacks for integer-programming and vertex-cover problems to minimize the overhead. Table 2.2 clearly shows this.

### 2.5.2 Modified B&B Algorithm for a Virtual-Memory System

The stacks reside in virtual space when the modified B&B algorithm is implemented with the virtual-memory management system. The virtual-memory management system automatically handles all node transfers between main memory and secondary memory so there is no need for the frame pool and its associated routines. However, reads from secondary memory to main memory are no longer overlapped with computations and the overhead analysis must be altered accordingly. Otherwise, the analysis is similar to that for the file system and will not be shown here.

### 2.6 Comparison of the Different B&B Algorithms

We have carried out extensive simulations to compare the different B&B algorithms. The simulation programs were written in the C language. These simulations used over one hundred hours of CPU time on the VAX 11/780. These simulators simulate a computer with a main memory of 128 Kbytes (10% of the virtual space).

Table 2.2

Normalized overheads for the analytical (AA) and equal-width (EWA) boundary assignments for five randomly generated, 80-node vertex-cover problems and five randomly generated, 20-by-20 integer programming problems.

Number of Stacks	Normalized Overhead						
	Problem 1	Problem 2	Problem 3	Problem 4	Problem 5		
EWA	AA	EWA	AA	EWA	AA	EWA	AA
Vertex-cover problems with $b_d = 5b_e$ and $b_d = 10b_e$ .							
2	8426	6196	8690	7274	11752	9714	11891
3	7047	6139	7157	6363	6231	9548	10423
4	5887	5942	8671	8257	8213	9120	11873
5	8271	8402	7218	6975	9888	10216	9894
Integer-programming problems with $b_d = 5b_e$							
2	4598	2958	2580	3561	581	846	2906
3	2555	1703	3412	3334	1419	612	2953
4	2040	3793	2200	2102	976	581	2237
5	1541	1585	2051	3610	680	2468	2287
Integer-programming problems with $b_d = 10b_e$							
2	4058	2058	5160	3561	581	846	2906
3	4370	5270	2880	2976	1300	605	4480
4	3860	3580	4400	4970	976	1710	3400
5	1880	2090	4700	3900	1360	713	3410
							3750
							3360

Figure 2.9 shows some of our simulation results for 20-variable, 20-constraint integer-programming problems. Each node requires 1874 bytes of memory space. The modified B&B scheme uses three stacks and a logical file block size of 8 Kbytes. The direct implementation uses a page size 8 Kbytes.

The modified virtual-memory scheme uses a page size of 16 Kbytes. We assumed a disk with a mean seek time of 28 ms., a mean rotational delay of 8.33 ms., and track capacities of 16 and 64 Kbytes (corresponding to transfer rates of approximately 1 and 4 Mbytes/sec.). The overhead is the completion time of the algorithm normalized by  $b_e$ . For the range of  $b_d/b_e$  simulated, the modified B&B scheme performs significantly better than the other schemes. Other simulations exhibited similar results.

Figure 2.10 shows some of our simulation results for the 80-node vertex-cover problems. Each node requires 330 bytes of memory. Here, the depth-first B&B algorithm performs as well as the modified B&B algorithm. This agrees with our expectation as the vertex-cover problem has a very accurate lower-bound function.

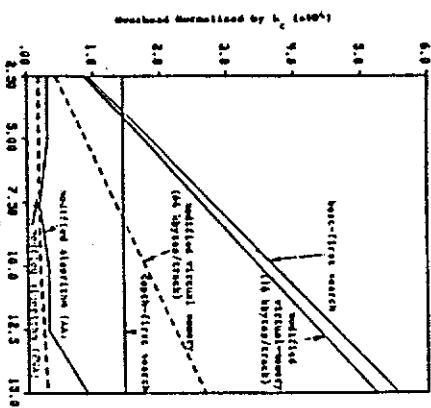
### 2.7 Summary

In a two-level memory system, the most time-efficient B&B algorithm depends on the numbers of nodes expanded by the B&B algorithms and the relative speeds of the main and secondary memories (Table 2.3).

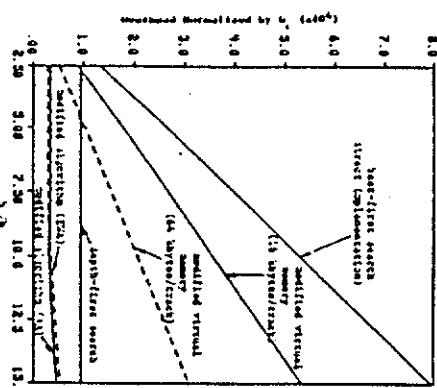
The best-first B&B algorithm expands the smallest number of nodes but requires a very large memory space. In contrast, the depth-first B&B algorithm is space efficient but expands more nodes than the best-first B&B algorithm. The modified B&B algorithm represents a compromise between the best-first and depth-first B&B algorithms.

Thus, the best-first B&B algorithm should only be used when the extra overhead due to its memory requirement is small. This requirement is met when the average time to transfer a node between main memory and secondary memory is close to the average time to expand a node. For slower secondary memories, the best-first B&B algorithm should be chosen only when it expands much less nodes than the other B&B algorithms.

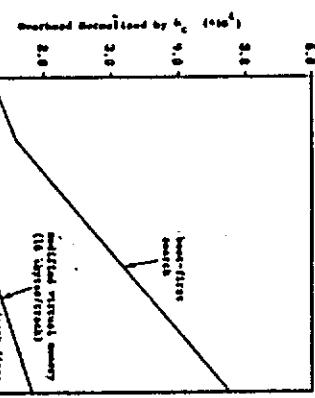
The depth-first B&B algorithm should be used to solve the problem for which there is only a small difference between the number of nodes expanded by the depth-first B&B algorithm and the number of nodes expanded by the best-first B&B algorithm. A greater difference is acceptable if the secondary



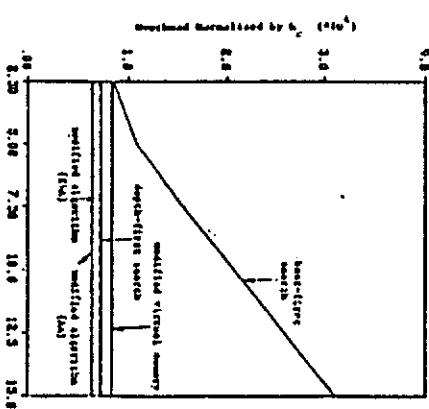
(a) Problem 1



(b) Problem 2



(a) Problem 1



(b) Problem 2

Figure 2.10 Simulation results for B&amp;B algorithms on two 80-node vertex cover problems. [Main memory size is 128 Kbytes. The completion time for the modified B&amp;B scheme is the same for both 16 Kbytes/truck and 64 Kbytes/truck in the second problem.]

Figure 2.9 Simulation results for B&amp;B algorithms on two 20-by-20 integer programming problems. [Main memory size is 128 Kbytes.]

memory is much slower than the main memory.

There will be problems for which both the best-first B&B algorithm and the depth-first B&B algorithms perform poorly: the depth-first B&B algorithm expands an excessive number of nodes, while the best-first B&B algorithm uses an excessive amount of memory. The modified B&B algorithm may be used for these problems.

The design of the modified B&B algorithm illustrates one important characteristic of many non-deterministic search problems. The non-determinism arises because the search is guided a fallible heuristic function that does not always lead in the right direction and results in backtracking when a wrong decision is made. Strict adherence to the heuristic function may not always be rewarding when it incurs a high overhead. In this case, adherence to the heuristic function should be relaxed. In virtual memory systems, this relaxation may be accomplished by processing other nodes while waiting for the node indicated by the heuristic function to be transferred from the secondary memory to the main memory. Item indicated by the

Table 2.3

Summary of B&B algorithms to be used.

Relative Speed of Secondary Memory			
$n_{dh} / n_{bfs}$	fast	moderate	slow
small	dls	dls	dls
moderate	bfs/mbb	mbb	dls/mbb
large	bfs	bfs	bfs/mbb

Keys:      bfs: depth-first search;

        bfs: best-first search;

        mbb: modified branch-and-bound algorithm;

        n<sub>bfs</sub>: number of nodes evaluated in bfs;

        n<sub>dh</sub>: number of nodes evaluated in dls.

## CHAPTER 3 LEARNING DOMINANCE RELATIONS

### 3.1 Introduction

Dominance relations are powerful, problem-dependent pruning rules used in combinatorial search algorithms to eliminate nodes that will not lead to the optimal solution. The problem-dependent nature of dominance relations has inhibited the development of a general procedure for the derivation of dominance relations. However, it is possible to overcome this obstacle by knowledge-based artificial intelligence techniques, specifically machine learning techniques. This chapter studies machine learning of dominance relations. Section 3.2 describes the five step methodology used in this study. Sections 3.3 to 3.6 describes the results of applying the first three steps and part of the fourth step in this methodology to the problem of learning dominance relations. Chapter 4 presents the rest of the results.

### 3.2 Methodology

Winston [Win84] has suggested the following commandments of a good AI methodology:

- First, identify the problem.
- Then select or devise an appropriate representation.
- Next expose constraints or regularities.
- Only now create particular procedures.
- Finally, verify via experiments.

Our methodology closely follows these commandments. We start by identifying the problem that we want to solve. Successful AI systems often employ domain knowledge. Combinatorial search problems come from many different problem domains. Moreover, little is known about the process of deriving dominance relations, so it is prudent to identify a particular subset of the combinatorial search problems to study instead of tackling the general problem at this stage. Section 3.3 discusses this topic in greater detail.

This is followed by a study of appropriate representations. A representation is a set of conventions about how to describe a class of things. A description makes use of these conventions to describe a member of this class of things. In theory, all computer-based representations are identical. In practice, one scheme may be better than another because it offers more convenience to the user. There are two issues to address in this area: the representation for the dominance relations and the representation for the search problems. Section 3.4 discusses these issues.

Next, we look for constraints in the problem domain which can simplify the task of finding dominance relations. This involves a two-step process. First, we seek a characterization of the structure of the domain occupied by dominance relations. Next, we seek the constraints, if any, on the permissible forms of the dominance relations exposed by their characterization. Possible constraints are trivial relations and relations with implementation overheads that clearly exceed the savings resulting from their use. Section 3.5 covers this topic.

We then set out to devise the learning mechanisms which carry out the actual work of learning the dominance relations. We employ a two-phase procedure. The initial phase seeks to identify and describe the applicable learning mechanisms. These descriptions include the aspects of the problem to which each learning mechanism may be applied, the important issues in their implementations, and the interactions, if any, among the different learning mechanisms. Section 3.6 discusses the initial phase of the process of devising learning mechanisms for the dominance relations.

The refinement phase develops the actual procedures (production rules) that implement the identified learning mechanisms. Each learning mechanism is developed using the *design-prototype-test* method commonly used in the development of knowledge intensive systems. Here, an initial design is created and a prototype implementing the design is tested. The test results often suggest modifications to the design. Once the appropriate modifications have been made to the prototype, further testing of the prototype is carried out. The new test results, in turn, may suggest additional modifications to the design. This process continues until the design goals have been met. At times, the test results may require such extensive modifications that a redesign of the entire system may be needed.

When there are more than one learning mechanism, it is prudent to interlace the refinements of the different learning mechanisms. Work on a learning mechanism may bring new insights to the design of another learning

mechanism. This allows work to proceed more smoothly than if attention were paid solely to one mechanism.

Sometimes, it may also be necessary to modify the structures resulting from steps one, two, and three of the methodology so that they are consistent with the results generated by the continuing refinement of the different learning mechanisms.

Chapter 4 describes the design of a prototype system that implements one of the learning mechanisms for dominance relations. We have experimented and verified the capabilities of the prototype system. Chapter 4 also summarizes the results of applying the prototype system to several test problems.

### 3.3 Problem Definition

As indicated in Section 1.2, combinatorial search problems may be divided into decision problems and optimization problems. Two functions may be associated with the optimization problem, I, [Ga79]: Length[] is the number of symbols used to describe I under some reasonable encoding scheme, and Max[] is the magnitude of the largest integer occurring in problem I. The optimization problem, I, is a number optimization problem if there is no polynomial, P, such that  $\text{Max}[I] \leq p[\text{Length}[I]]$ . A similar definition holds for the number decision problem. The solution of a combinatorial search problem requires pseudo-polynomial time if its time complexity function is bounded from above by a polynomial function of the two variables Length[] and Max[]. Most known dominance relations are associated with problems solved by dynamic programming, and these problems usually have pseudo-polynomial time solutions.

It has often been said among AI researchers that learning something new requires knowing a lot about it already. Thus, it is prudent to study a problem domain known to be rich in dominance relations. Many problems that are known to have dominance relations are number optimization problems; thus, the problems studied in this thesis are the number optimization problems. Moreover, dominance relations are generally more effective in problems with pseudo-polynomial time solutions than in problems with exponential time solutions. Dominance relations do not change the time complexity of the search process. As a result, dominance relations generally produce a smaller percentage improvement in problems with exponential time solutions.

### 3.4 Representation

The first facet of the representation problem is the choice of a representation for the dominance relations. A dominance relation is essentially a statement that if two nodes satisfy the conditions set forth in the dominance relation, then one node will contain a solution no worse than the other node. This statement can be naturally expressed by an IF-THEN rule (production rule). The left-hand side establishes the conditions that two nodes must satisfy in order for one node to dominate the other node. The right-hand side establishes the relationship between the best solutions contained in the two nodes. This production rule can easily be transformed into an implication in first-order predicate logic, so first-order predicate logic is an equally good representation for dominance relations. This easy transformation permits the switching between production rules and predicate logic as the situation dictates.

The other facet of the representation problem is the representation for the number optimization problems. A good representation exposes aspects of the problem involved in learning dominance relations. This requirement implies the need for relevant domain knowledge about the problem. A good representation also permits techniques developed for one problem to be used for other related problems. An important aspect of this requirement is a systematic method of incorporating domain knowledge in the representation so that relationships may be drawn between the domain knowledge of different problems.

One way of expressing a number optimization problem is in the form of a constrained optimization:

$$\begin{aligned} &\text{Minimize} && C_0(x) \\ &\text{subject to} && g_i(x) \leq 0 \quad i = 1, 2, \dots, m \end{aligned}$$

This description is inadequate for learning dominance relations as it does not include any domain information. The relevant domain information is related to the subject matter of the problem. One classification of combinatorial search problems [Ga79] lists twelve major subject areas: graph theory, network design, sets and partitions, storage and retrieval, sequencing and scheduling, mathematical programming, algebra and number theory, games and puzzles, logic, automata and languages, program optimization, and miscellaneous. Each major subject area is divided into several minor subject areas. The number optimization problems can occur in any of these subject areas except logic.

Studying all of these subject areas is non-trivial and time-consuming. As this study is an initial study of the problem of learning dominance relations, we choose a more detailed study of some of the subject areas instead of a less detailed study of more subject areas. The work of this thesis uses problems in mathematical programming and problems in sequencing and scheduling. We have substantial experience with problems in these two areas, and have abundant heuristic knowledge about them. Another reason for this choice is that the problems in these two subject areas are somewhat related. Mathematical programming problems studied in this thesis are the 0/1-knapsack problem, the inventory problem, and machine reliability by duplication. Problems in sequencing and scheduling studied in this thesis are sequencing to minimize weighted completion time, sequencing to minimize latest completion time, sequencing to minimize weighted lateness, sequencing to minimize task weights, sequencing to minimize tardy task weights, and the two-machine flowshop problem.

We will use the 0/1-knapsack problem, the inventory problem, and sequencing to minimize weighted completion time to illustrate the design of the problem representation.

The 0/1-knapsack problem consists of a set  $U$  of  $n$  items. The knapsack has a capacity  $K$ . Each item  $u \in U$  has a weight  $w(u)$  and a value  $v(u)$ . The desired solution is the subset  $U' \subseteq U$  which maximizes  $\sum_{u \in U'} v(u)$  and satisfies  $\sum_{u \in U'} w(u) \leq K$ .

In sequencing to minimize weighted completion time, there is a set  $T$  of  $n$  tasks. Each task  $t \in T$  has a length  $\ell(t)$ , a weight  $w(t)$ , and a deadline  $d(t)$ . The desired solution is the one-processor schedule  $\sigma$  for  $T$  which minimizes  $\sum_{t \in T} (\sigma(t) + \ell(t))w(t)$  and satisfies  $(\sigma(t) + \ell(t)) \leq d(t)$  for each task  $t \in T$ .

The inventory problem deals with the storage of an item over a set  $M$  of  $n$  periods. Initially, there are  $v$  units of the item. In each period  $m \in M$ , the item has a unit purchase price of  $b(m)$  and a unit sale price of  $s(m)$ . The goal is to determine the amount purchased  $x(m_i)$  and the amount sold  $y(m_i)$  in each period  $m_i$ ,  $1 \leq i \leq n$ , which will maximize  $\sum_{i=1}^n (x(m_i)s(m_i) - y(m_i)b(m_i))$ . There can at most be  $B$  units of the item at any time, so  $y + \sum_{i=1}^n (x(m_i) - y(m_i)) \leq B$  for all  $j \in \{1, \dots, n\}$ . Amounts purchased are delivered at the end of each month, so the amount sold in each month must

not exceed the stock at the end of the previous month, and  $y(m_i) \leq v + \sum_{j=1}^{i-1} (x(m_j) - y(m_j))$  for all  $i \in [1, n]$ .

Although these three problems involve quite different domains, they are quite similar in form (Table 3.1). Each problem description has one set of entities. A set of decisions must be made on a set of entities and a solution is found when decisions have been made for every entity. Each description contains constraints which define the allowable sets of decisions. Each description also contains an optimization criterion which defines the desired optimal solution. This is either the maximization or the minimization of a objective expression. The objective expression and the constraints are arithmetic expressions.

There are three types of symbols in the descriptions: *decision symbols*, *entity variables*, and *problem constants*. Decision symbols are associated with the set of decisions. Entity variables are associated with the set of entities. The entity variables may be subdivided into *return properties*, *resource properties*, and *limit properties*. The return property appears in the optimization criteria. Resource properties and limit properties appear in the constraints. Resource properties are usually combined with decision symbols in the constraints. Limit properties establish the limits in the constraints. It is possible for an entity variable to be both a return property and a resource property. The remaining symbols in the descriptions are problem constants. Problem constants include set sizes, initial values, and global limits.

Before a search problem can be solved efficiently, there must be a bookkeeping scheme that allows efficient orderly exploration of all possible solutions. A key in such a definition of the search space is the transformation of the set of decisions into a form suitable for efficient sequential decision making. Let this new set be known as the set of working decisions. The subset  $U'$  in the 0/1-knapsack problem has to be transformed into the set of membership decisions  $\{x(u) : u \in U\}$ . In scheduling to minimize weighted completion time, the schedule time for task  $t$ ,  $\sigma(t)$ , is usually simplified into the position of task  $t$  in the schedule,  $\rho(t)$ . These transformations often require deep domain knowledge. For instance, the replacement of  $\sigma(t)$  by  $\rho(t)$  is possible only after knowing that there is no waits between consecutive tasks. The optimal schedule is "as soon as possible". This requires specific values to decision variables.

Let the problem description actually used by the problem solver be known as a working description. The first step in forming a working description from

Table 3.1 Similarities between the problems descriptions of the 0/1-knapsack, scheduling to minimize weighted completion time, and inventory problems.

Problems	0/1-KNAPSACK	SCHEDULING	INVENTORY
Entity	item $u_i$ , set $U$ of items	task $t_j$ , set $T$ of tasks	period $m$ , set $M$ of periods
Decision	subset $U'$	schedule $\sigma$	amount purchased $x_{im}$ , amount sold $y_{im}$ , $\forall m \in M$
Constraints	$\sum u_i \leq K$	$\forall t \in T$	$y_{im} \leq v + \sum_{j=1}^{i-1} x_{jm} - y_{jm}, \forall j \in \{1, \dots, n\}$
Objective	maximize $\sum u_i$	minimize $\sum_{j=1}^n [c_j(t_j) + d_j(t_j)] \cdot \sigma_j$	maximize $\sum_{m=1}^M [k_m(m) x_{im} - y_{im}]$
Criteria			
Symbols	Decision Symbols	$U'$	$x_{im}, x_{im}$
Entity Properties	$(u_i, w_i)$	$d_j(t_j), d_j$	$b_{im}, s_{im}$
System Constants	$K, n$	$n$	$v, a$
Working Decision	$x_{ij}$	$\sigma_j$	$x_{im}, y_{im}$

The 0/1-knapsack problem has a value-description as well as an identity-description. Decision variable,  $v_j$ , is the membership status of item  $u_j$ ;  $v_j$  has value 0 if item  $u_j$  is to be excluded from the knapsack and has value 1 if  $u_j$  is to be included in the knapsack. In contrast, decision variable,  $j_p$ , represents the identity of an item that is to be included in the knapsack.

Sequencing to minimize weighted completion time also has a value-description as well as an identity-description. Decision variable,  $v_j$ , is the position of task  $t_j$  in the schedule. Decision variable,  $v_j$ , is the identity of the task occupying position  $j$  in the schedule.

The inventory problem only has a value-description. The lack of a suitable implicit value for the  $v$ -tuple prevents the inventory problem from having an identity-description. For even values of  $j$ ,  $v_j$  is the amount purchased in month  $j/2$ . When  $j$  is odd,  $v_j$  is the amount sold in month  $[j/2]$ .

The similarities between the different problem descriptions strongly indicate the presence of a domain-independent structure in the problem descriptions. Each problem description appears to have the same set of domain-independent elements: the entities, the decision variables, the state

the problem description is to replace the decision symbols by decision variables. Working descriptions use state variables to assist in the assignment of values to the decision variables. State variables record the effects of previous decisions.

A complete specification of a working decision requires a two-tuple ( $e$ -tuple,  $v$ -tuple) where  $e$ -tuple identifies its entity and  $v$ -tuple specifies its value. Decision variables in a general-description explicitly specify both the  $e$ -tuple and the  $v$ -tuple. A general description is not necessary for most problems. The decision variable,  $v_j$ , in a value-description determines the  $v$ -tuple of the  $(E_j, v_j)$ . A value-description produces a fixed height search tree in which all solution nodes are at a fixed distance from the root. In contrast, the decision variable,  $j_p$ , in an identity-description only specifies the value of the  $e$ -tuple. One common type of identity-description assumes a constant value for the  $v$ -tuple of the decision whose  $e$ -tuple is given by the decision variable. This assumption is often combined with the assumption of some other constant value for the  $v$ -tuples of all unassigned decisions. Another type of identity-description assumes that the decision made by decision variable  $j_p$  is  $[j_p, j]$ . The identity-description results in a search tree that may either be a fixed-height tree or a variable-height tree.

variables, the entity variables, the problem constants, the optimization function, and the constraints. Furthermore, this set of elements is complete, that is, the problem description can always be completely described in terms of the basic elements.

This structure can be exposed by a frame representation [Min75]. The slots in the frame are fixed and represent the different elements of the problem description. The slot filters describe the actual values of these elements for the particular problem being described.

Figure 3.1 shows the description frames for the 0/1-knapsack problem. The problem description frames for the other problems studied in this thesis are similar. The slots in the description frame are domain-independent. Hence, we can expect this frame representation to be applicable to all number optimization problems.

### 3.5 Constraints in the Problem Domain

As indicated in Section 3.2, the first step in exposing constraints in the problem domain is to find a suitable characterization of the problem domain. One possible classification of dominance relations is a classification based upon the possible positions of the dominating and dominated nodes. Each type of dominance relations is characterized by a three tuple  $\{T_1, T_2, T_3\}$ .  $T_1$  is the set of permissible levels for the dominating node,  $T_2$  is the set of permissible differences between the levels of the dominating and dominated nodes, and  $T_3$  is the relation between the dominating and dominated nodes.

The possible symbolic values for  $T_1$  are term and all. The dominating node must be a terminal node when  $T_1$  is term but can be any terminal or non-terminal node when  $T_1$  is all.

The value of  $T_2$  can be same, specific, or all. If  $T_2$  has the value same, the dominating node and dominated node must be on the same level, that is, they are equivalent nodes. When  $T_2$  is specific, the difference between the level of the dominating node and the level of the dominated node must be a symbolic constant. There is no restriction on the position of the dominated node when  $T_2$  is all.

$T_3$  has a value of either sib or all. The dominated node must be a sibling of the dominating node if  $T_3$  has the value sib but can be any arbitrary node if  $T_3$  has the value all.  $T_3$  can take on a value of sib only when  $T_2$  has a value of same; otherwise, the resulting classification does not make any sense.

ENTITY	$u_1, \dots, u_n$
DECISION VARIABLE	$x(u_i) = \{0,1\}, i=1, \dots, n$
STATE VARIABLES	$R(\text{apply}(x(u_i), \text{NODE}))$ $= R(\text{NODE}) + x(u_i) * r(u_i);$
CONSTRAINT	$W(\text{apply}(x(u_i), \text{NODE}))$ $= W(\text{NODE}) + x(u_i) * w(u_i);$ $R(\text{ROOT}) = 0; \quad W(\text{ROOT}) = 0$
ENTITY VARIABLES	$r(u_i), w(u_i)$
PROBLEM CONSTANTS	$K, n$
OPTIMIZATION	maximize $R(S);$ $S = \text{apply}(x(u_1), \dots, x(u_n), \text{ROOT})$
CONSTRAINT	$W(\text{NODE}) \leq K$

(a) Value description

ENTITY	$u_1, \dots, u_n$
DECISION VARIABLE	$x(u_i) = \{u_1, \dots, u_n\} \quad i=1, \dots, n$
STATE VARIABLES	$R(\text{apply}(x(u_i), \text{NODE}))$ $= R(\text{NODE}) + r(x(u_i))$
CONSTRAINT	$W(\text{apply}(x(u_i), \text{NODE}))$ $= W(\text{NODE}) + w(x(u_i))$

(b) Identity description

ENTITY VARIABLES	$r(u_i), w(u_i)$
PROBLEM CONSTANTS	$K, n$
OPTIMIZATION	maximize $R(\text{NODE})$
CONSTRAINT	$W(\text{NODE}) \leq K$

Figure 3.1 Description frames for the 0/1-knapsack problem.

Table 3.2 lists examples of the different types of dominance relations. A relation of type {term, same, sib or all} is a trivial relation as it is a relation between terminal or solution nodes. The value of the solution represented by the terminal nodes are known, so it is trivial to determine which is the better solution.

The lower-bound test, used in B&B algorithms as the elimination rule, has type {term, all, all}. The dominating node in the lower-bound test is the best solution node found so far and it can dominate any node whose lower bound is worse than its value.

A Greedy algorithm always chooses the optimal value for each decision variable. This implies that the search tree generated is a skewed tree in which one node dominates all its siblings at each level in the search tree and that greedy algorithms are dominance relations of type {all, same, sib}.

The upper-bound test is an example of a dominance relation of type {all, all, all}. In the upper bound test, a node dominates another node if the upper-bound of the first node is not greater than the lower-bound of the second node. No restrictions are placed on the positions of the dominating and dominated nodes.

Dominance relations of type {all, specific, all} are usually trivial relations or relations that correspond to well-behaved relations in other descriptions. For example, each node in the problem space generated by the identity-description for the 0/1-knapsack problem is a feasible solution and dominates its parent node. This relation corresponds to the trivial relation between terminal nodes in the corresponding value-description.

This interpretation of dominance relations with type {all, specific, all} implies that there is no need to explicitly look for dominance relations of this form as they will either be trivial relations or be found in another form.

Another constraint on the form of the useful dominance relations comes from studying the implementation overheads of the different types of dominance relations. The implementation overhead can be measured by the number of nodes to which each newly-generated node must be compared in order to determine if it either dominates another active node or is dominated by an active node (Table 3.3). The number of comparisons required by dominance relations between nodes on arbitrarily different levels is on the order of the number of active nodes. Unless the dominance relation is very powerful, the number of active nodes grows exponentially as the search proceeds. It is very likely that the implementation overhead of such dominance relations will exceed the savings resulting from the nodes that they eliminate. Thus, there is

Table 3.2 Examples of the different classes of dominance relations.

T <sub>1</sub>	T <sub>2</sub>	T <sub>3</sub>	EXAMPLE
term	same	sib	trivial relations
specific	all	all	trivial relations
all	all	all	lower-bound test
all	same	sib	greedy algorithm
all	all	all	dominance relations in dynamic programming
specific	all	all	
all	all	all	upper-bound test

also no need to look for dominance relations of types {*all*, *all*, *all*}.

Furthermore, dominance relations of types {*term*, *same*, *sib*} and {*term*, *same*, *all*} are trivial relations. Hence, a study of methods to find dominance relation may be restricted to dominance relations belonging to the classes {*all*, *same*, *sib*} and {*all*, *same*, *all*}.

### 3.6 Applicable Learning Mechanisms

Learning is a very general term denoting the way people (and computers) increase their knowledge and improve their skills. Figure 3.2 shows a simplified model of a machine learning system [CoF82]. The environment supplies some information to the learning element; the learning element uses this information to make improvements in the explicit knowledge base, and the performance element uses this knowledge base to perform its task. Information gained during attempts to perform the task provide feedback to the learning element.

The key factor in the design of learning systems is the information that it receives from the environment - particularly its level and its quality. The level of information refers to its degree of generality relative to the needs of the performance element. High-level information is abstract information that is relevant to a broad class of problems. Low-level information is detailed information that is relevant to only a few problems. The learning element bridges this gap between the level of the information provided by the environment and the level required by the performance element. Thus, if the learning system is given very abstract (high-level) information, it must fill in the missing details that allow the performance element to interpret the information in particular situations. Correspondingly, if the system is given very specific (low-level) information about how to perform in particular situations, the learning element must generalize this information by ignoring unimportant details into a rule that can be used by the performance element in other related situations.

The learning element does not know in advance exactly how to fill in the missing details or ignore unimportant details because its knowledge is imperfect. Consequently, it must guess, that is, form hypotheses, about how the gap between the levels should be bridged. After guessing, the system must receive some feedback that allows it to evaluate its hypotheses and revise them if necessary. Thus, a learning system learns by trial and error.

The level of information provided by the environment determines the kinds of hypotheses that the system must generate and determines the type of

Table 3.3 Upper bounds on the implementation overheads of different classes of dominance relations [ $\text{LEV}(A)$  is the level of node A].

$T_i$	$T_j$	$T_k$	Dominated Node	Dominating Node
		sib	$O(1)$	$O(1)$
same		all	$O(1)$	$O(1)$
	all		$O( \text{active nodes} )$	$O( \text{nodes on level specific} )$
term	spec.	all	$O( \text{terminal nodes} )$	$O( \text{active nodes} )$
	all	all	$O( \text{active terminal nodes} )$	$O( \text{active nodes} )$
		sib	$O( \text{siblings} )$	$O( \text{siblings} )$
same		all	$O( \text{nodes on the same level} )$	$O( \text{nodes on the same level} )$
all	spec.	all	$O( \text{nodes on level } (\text{LEV}(A) - \text{specific}) )$	$O( \text{nodes on level } (\text{LEV}(A) + \text{specific}) )$
	all	all	$O( \text{active nodes} )$	$O( \text{active nodes} )$

learning strategy used [Cor82, CM83, Mic86]. Several basic strategies have been distinguished: *role learning*, *learning by instruction*, *learning by deduction*, *learning by analog*, and *learning by induction*. The latter subdivides into *learning from examples* and *learning by observation and discovery*. These strategies are ordered by the increasing complexity of the hypothesis (inference) and reflects increasing effort on the part of the learning element and correspondingly decreasing effort on the part of the environment (teacher).

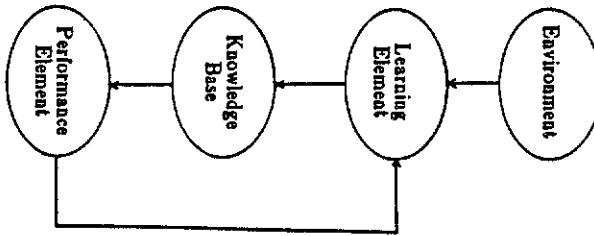
In role learning, the environment supplies knowledge in a form that can be used directly. The learning system need not understand or interpret the information supplied by the environment. All it does is to memorize the information for later use. One major concern here is how to index the stored knowledge for future retrieval. Role learning may be viewed as a form of data reduction [HK79], that is, it reduces a calculation to an access task by saving the input/output details of some computation and so bypass any future need for the intermediate computation process. Role learning was used in Samuel's famous checker-playing program [Sam59, Sam69]; the program was able to improve its performance by memorizing board positions and their associated minimax scores. It uses previously memorized board positions to improve the speed and depth of its look-ahead search during subsequent games.

Role learning can only be used for problems with known dominance relations and so cannot learn new dominance relations.

In *learning from instruction*, knowledge is acquired from a teacher or an organized source such as a textbook. Here, the system is given general-purpose knowledge or advice. It transforms this high-level knowledge into a form that can be used readily by the performance element. This transformation is called *operationalization*. The system must understand and interpret the high-level knowledge and relate it to what it already knows. Operationalization is an active process that can involve such activities as deducing the consequences of what it has been told, making assumptions and filling in the details, and deciding when to ask for more advice. Work in this area includes McCarthy's proposal for an 'advice-taker', the TERRASIS program [Pal82], and the program FOO [McH79].

At present, there is neither an organized source of knowledge about dominance relations nor an general expert in the derivation of dominance relations. Thus, learning by instruction is of little use in learning dominance relations although this assessment may change as more is known about dominance relations in the future.

Figure 3.2 A simplified model of a machine learning system.



In learning by deduction or deductive learning, the learner draws deductive, truth-preserving inferences from the knowledge given and stores useful conclusions. Learning by deduction includes knowledge reformulation, knowledge compilation, creation of macro-operators, caching, chunking, equivalence-preserving operationalization, and other truth-preserving transformations.

Learning by deduction includes aspects of automatic deduction or mechanical theorem proving [Nil80, WOL84]. An important, early advance in automatic deduction was the development of resolution procedures based upon the *resolution principle* [Rob55]. The resolution procedures constitute a logically complete method for proving theorems in first-order predicate calculus. However, the search space generated by the resolution method grows exponentially with the number of formulas used to describe a problem, so that problems of even moderate complexity cannot be solved in a reasonable time. Several domain-independent heuristics, e.g., set of support [WRC65], were proposed to deal with this issue, but they proved to be too weak to produce satisfactory results. There are too many possible inferences that can be drawn at any one time, and finding the relevant inferences can be impossible without domain-specific guidance. Furthermore, the efficiency of deductive reasoning is greatly affected by the way in which the body of knowledge is formalized.

Nonetheless there is an important class of problems that resist solution by any other method and must be solved by deduction. Deduction permits logically complex queries of a knowledge base containing generalizations even when we cannot evaluate a query directly. Recent work in deduction include work on common-sense reasoning [McD98, Doy99] and nonmonotonic logic [Bob90]. Nonmonotonic logic can retract a conclusion in the face of further evidence, so its set of inferable conclusions does not increase monotonically with the set of premises as in conventional deductive logics. While many procedures have been implemented that support this type of reasoning, their theoretical foundations are questionable.

Dominance relations can be learned by deduction, in particular, by automated deduction. A dominance relation is easily expressed in logic as an implication statement. Dominance relations are true theorems. Thus, if a candidate relation which have the prescribed form can be proved to be logically true, then it must be a dominance relation for the problem under study. Automated deduction is very good at finding proofs. Section 3.6.1 discusses in further detail the use of automated deduction to find dominance relations.

In learning by analogy, the learning system recognizes analogies in a related performance task and transform relevant knowledge in the knowledge base of the related task into a form suitable for use by its performance element. The study of learning by analogy has been mainly on the use of analogy in automated problem solving [Eva68, Kli71, MoN73, BrT79, McD79, Win79, Car83, Bur86], and the use of analogy to guide the modification of programs [MaW75, DeM77, Am86].

The close relationship between everyday notions of analogy and similarity has led to several models of analogical reasoning based on forms of partial mapping [Eva88, Win79]. These algorithms assume that the best partial match can be found by accumulating evidence for each of a number of possible object-to-object mappings between the two tasks and then choosing the one that scores the highest. In these systems, evidence for a match consists essentially of the number of relational connections preserved between the corresponding objects for a given alignment of objects. This approach presupposes that well-defined, bounded representational models of the tasks in both the base (that is, familiar) domain and the target domains are available as inputs. This presumption is often not true in learning situations, particularly if the new domain is totally unfamiliar. Matching cannot be used to construct such a representation where there was none before. Moreover, as the complexity of the descriptions increases, the conceptual representations for many real situations may contain many objects that do not take part in a specified analogy.

An alternative approach [Bur86, Car83] is based on analogical mapping. Such an approach uses a set of heuristics to delimit what is to be "imported" from a base to a target domain at a given time. One such heuristic involves mapping previously formed abstractions, such as underlying causal and planning rules. Focusing on such structures and their associated special cases and known problems allows for a much more top-down form of analogical reasoning.

Analogy is capable of learning new dominance relations. If the learning system has available to it a pool of problems with known dominance relations, it may be able to recognize analogies between the new problem domain and known problem domains. The common substructure between the two descriptions can then guide the transformation the dominance relations in the known domains into relations in the new domain. If the common substructure has been correctly identified and the transformation is sound, then the derived relation are likely to be dominance relations in the new domain. Section 3.6.2

describes learning dominance relations by analogy in greater detail.

Learning by induction involves the generalization of input information and selection of the most plausible or desirable result. Inductive learning can be subdivided into learning from examples and learning by observation and discovery. In learning from examples, the system is taught how to perform a task by presenting it with examples of how it should behave (positive instances) and examples of how it should not behave (negative instances). Examples can be viewed as pieces of very specific knowledge that cannot be used efficiently by the performance element. The system generalizes the examples to find higher level rules that can be applied to guide the performance element. The examples are provided by a source of information, which can be a teacher that knows the concept [And84, Buc78, DM83, FHNT2, Mic83, Ming84, Mit82, Ver78, Win75] or the environment on which the student performs experiments and from which it receives feedback [Wat70, AnSt9, Ob83, Lan83, Ara84]. The latter case is called *learning by experimentation* (this includes *learning by doing* and *learning by problem solving*).

In learning by observation and discovery (also called *descriptive generalization*), one searches without help of a teacher, for regularities and general rules explaining all or at least most observations. This form of learning includes *conceptual clustering* (form object classes describable by simple concepts), constructing classifications, fitting equations to data, discovering laws explaining a set of observations [Lan81, LBS83], and formulating theories accounting for the behavior of a system [Dal82, Len82, Len83a, Len83b].

Simon and Lea [SL74] describe the problem of inductive learning as the problem of using training instances, selected from some space of possible instances, to guide the search for general rules. They call the space of possible training instances the *instance space* and the space of possible general rules the *rule space*. An intelligent program may also select its own training instances by actively searching the instance space in order to resolve some ambiguity about the rules in the rule space. A learning system moves back and forth between an instance space and a rule space until it converges on the desired rule. When the training instances are in a form far removed from the form of the rules, special processes are needed to interpret the raw training instances. Similarly, when the program needs some new training instances, experimental planning routines use the current high-level hypotheses to guide the search of the instance space.

One issue involving the instance space is the *quality* of the training instances. High-quality training instances are unambiguous and provide

reliable guidance to the search of the rule space. Low-quality training instances invite multiple, conflicting interpretations and only provide tentative guidance to the rule-space search.

Ambiguity also arises if the program must learn from *unclassified training instances*. In these so-called *unsupervised learning situations*, the program is given heuristic information that it must use to classify the training instances itself. If this heuristic knowledge is weak and imperfect, the rule-space search must treat the resulting classifications as being potentially incorrect.

Another factor affecting the quality of the training instances is the *order* in which they are presented. A good training sequence systematically varies the relevant features to determine which features are important.

With unambiguous training instances, the program can embody a whole set of constraining assumptions about the instances that permit rapid location of the appropriate rules. Ambiguous instances cause the program to consider a much larger space of hypotheses. If the training instances can contain errors, the program must consider that any given incorrect instance is due to either measurement error or classification error.

Another design issue is how to search the instance space. Most work in inductive learning have avoided this issue by assuming either that the instances are presented all at once or else that the program has no control over their selection. Programs that update their hypotheses as additional training instances are made available perform *incremental learning*. Programs that explicitly search the instance space perform *active instance selection*.

Most methods of searching the instance space use a set,  $H$ , of hypotheses in the rule space that are currently believed by the program to be the most plausible. One approach tries to discriminate as much as possible among the alternatives within  $H$ . A training instance can be chosen that "splits  $H$  in half", such that half of the hypotheses can be ruled out when the new instance is obtained. Another approach is to choose the most likely hypothesis in  $H$  and try to confirm it by checking additional training instances (particularly instances with extreme characteristics). This permits the learning system to determine the limits of applicability of the hypothesis under consideration. A third approach, called *expectation-based filtering*, selects training instances that contradict the hypotheses in  $H$ . The learning program can then focus its attention on those instances in which its current hypotheses break down.

The training instances may need to be transformed before they can be used to guide the search of the rule space. This transformation process can be quite difficult, especially in perceptual learning tasks [Win75].

The rule space is usually defined by specifying the kinds of operators and terms that can be used to represent a rule. The designer of a learning system seeks to choose a rule space that is easy to search and that contains the desired rule or rules. Two factors that influence the choice of a representation language for the rule space are the kinds of inference supported by the representation and the *single-representation trick*. The single-representation trick uses the same representation for both the rule space and the instance space. Training instances can then be viewed literally as highly specific pieces of acquired knowledge.

The four main methods that have been used to search the rule space differ primarily in how they refine,  $H$ , the set of the currently most plausible rules. The presentation of the training instances drive the search in *data-driven methods*, while *a priori* model guides the search in *model-driven methods*.

The first data-driven method is the *version-space method* (and several related techniques). This approach uses the single-representation trick to represent training instances as very specific points in the rule space. The set  $H$  is initialized to contain all hypotheses consistent with the first positive training instance. New training instances are examined one at a time and pattern-matched against  $H$  to determine whether the hypotheses in  $H$  should be generalized or specialized.

The second data-driven method does not use the single-representation trick but uses special procedures (or knowledge) to examine the set of training instances and decide how to refine,  $H$ , the current set of hypotheses. The program can be viewed as having a set of *hypothesis-refinement operators*. In each cycle, it uses the data to choose one of the operators and then applies it. This approach is used in Langley's BACON system [Lang81, LBS83], Lenat's AM system [Dal82], and Lenat's EURISKO system [Len83b].

The third approach is *model-driven generate and test*. This method repeatedly generates and test hypotheses from the rule space against the training instances. Model-based knowledge is used to constrain the hypothesis generator to generate only plausible hypotheses. Meta-DENDRAL [Bab78] and INDUCE 1.2 [DIN83] use this approach.

The final approach is *model-driven schema instantiation*. It uses a set of rule schemas, which are abstract skeletal rules, to provide general constraints on the form of plausible rules. The method attempts to instantiate these schemas from the current set of training instances. The instantiated schema that best fits the training instances is considered the most plausible rule. Schema instantiation methods are useful when a system has many constraints

that can be grouped together to form a schema. The SPARC system [Die80] uses schema instantiation to solve a learning problem that arises in the card game Eleusis.

Data-driven techniques generally support incremental learning. The version space method can easily modify its set  $H$  to account for new training instances without any backtracking by the learning program. In contrast, model-driven methods, which test and reject hypotheses based on an examination of the whole body of data, must either backtrack or search the rule space again, because the criteria by which hypotheses were originally tested (or schemas instantiated) have changed.

However, model-driven methods tend to have good noise immunity. Since the whole set of training instances is available, the program can use statistical measures of how well a proposed hypothesis accounts for the data. When a set of hypotheses,  $H$ , is tested against noisy training instances, the model-driven methods need not reject a hypothesis on the basis of one or two counterexamples. In data-driven methods,  $H$  is revised each time on the basis of the current training instance. A single erroneous instance can cause a large perturbation in  $H$  (from which it may never recover).

Learning systems often use the *closed-world assumption*, that is, the rule space contains the desired concept. The closed-world assumption allows programs to locate the desired concept by progressively excluding concepts that are known to be incorrect. In some learning problems, the program can assume that the desired rule or rules exist somewhere in the rule space. Consequently, the search has a well-defined goal. In many situations, however, there is no such guarantee, and the learning program must confront the possibility that its representation of the rule space is inadequate and should be expanded. This is called the problem of *new terms*. One approach to expanding the rule space is to add new terms to the representation. The problem of defining new terms is quite difficult to solve. It is fairly easy to incorporate operators that create new terms in the hypothesis-refinement operator approach. Both BACON and AM programs create new terms by combining and refining existing terms.

In *single-concept learning*, the learning element is presented with positive and negative instances of one concept. It finds some concept that effectively partitions the space of all instances into two regions: positive and negative. In *multiple-concept learning*, there are training instances of several concepts. For each concept description, there is a corresponding region in the instance space. Overlapping concept descriptions can lead to integration problems. Hence, it may be necessary to modify existing rules when adding a new rule or concept.

to the knowledge base. Michalski's AQ11 [MC80] is a system that learns multiple concepts.

In *single-step tasks*, one rule or a set of independent rules can be applied in one step to accomplish the performance tasks. In *multiple-step tasks*, several rules must be chained together into a sequence. The learning element must now consider possible interactions among rules when modifying the knowledge base. There is also the problem of credit assignment. In a single-step task, the application of the performance standard immediately determines whether or not the rule is correct. In a multiple-step task, feedback is not available until the problem is completely solved or the game is complete. Moreover, the program can only determine if the entire sequence of rules is good and bad. The overall performance standard must be converted into performance standards for each rule. Learning systems which deal with multiple-step tasks include Samuel's checkers player [Sam67], Waterman's poker player [Wat70], Sussman's HACKER planning systems [Sus75], and Mitchell's LEX system for symbolic interpretation [MUB83].

Learning by examples can be used to learn dominance relations. A weak problem solver can generate the complete search tree of any sample instance of the given problem. In a complete search tree, the values of the best solutions that can be derived from each node are always known. Thus, the nodes in the search tree can provide examples of how the dominance relations, if any, should and should not behave. However, in the absence of an oracle, examples of dominance relations cannot be identified among the nodes in the problem trees without knowledge of the relations. The learning system cannot expect to be presented with a set of examples by a source of information. It must perform experiments and use the feedback from the environment to determine the desired examples. Section 3.8.3 considers learning dominance relations by experimentation in greater detail.

### 3.8.1 Learning Dominance Relations by Deduction

In the prototypical automated deduction problem we have a set,  $S$ , of well-formed formulas (wffs) from which we wish to prove some goal wff,  $W$ . A resolution-based system produces proofs by contradiction. We negate the goal wff and add the negated goal wff to the set of axioms. The resolution-based system attempts to derive a contradiction represented by the empty clause NIL. Derivation of the empty clause proves the goal wff.

As examples, we apply automated deduction to candidate relations from a value-description of the 0/1-knapsack problem. For each node,  $u$ , in the search tree, let  $D_u(u)$  represent the items included in the knapsack by the decisions which have already been made. The state variables are  $R(u)$ , the cumulative value of the items in  $D_u(u)$ , and  $W(u)$ , the cumulative weight of the items in  $D_u(u)$ . In addition, let  $L(u)$  be the level of node  $u$  in the search tree, and let  $B(u)$  be the value of the best solution contained in node  $u$ . Let the candidate relation and goal wff be

$$[=](L(x_1),L(y)) \& [>](R(x),R(y)) \& [=](W(x),W(y)) \rightarrow [>](B(x),B(y))$$

where  $=\{u,v\}$ ,  $<\{u,v\}$ , and  $>\{u,v\}$  are the predicates denoting the equal to, less than, and greater than relationships between  $u$  and  $v$ .

Negation of the goal wff produces the clauses:

$$\begin{aligned} C_1: &= (L(x_1),L(y)) \\ C_2: &> (R(x_2),R(y_2)) \\ C_3: &< (W(x_3),W(y_3)) \\ C_4: &\neg>(B(x_4),B(y_4)) \end{aligned}$$

The set of axioms consists of clauses that describe pertinent aspects of problem solving and the problem description. First, we tell the automated deduction program that the value  $B(u)$  is the sum of  $G(u)$  and  $H(u)$ , the value of the best solution of the remaining decisions.

$$C_5: = (B(x_5), \text{SUM}(R(x_5), H(x_5)))$$

In a value-description, equilevel nodes have identical sets of undecided items. D. Thus, if  $W(u) < W(v)$  for equilevel nodes  $u$  and  $v$  then for every solution derived from  $v$  which includes the subset  $D' \subseteq D$  in the knapsack, there is a solution derived from  $u$  which includes the subset  $D'' \subseteq D$  in the knapsack such that  $D' \subseteq D''$ . This implies that  $H(u) > H(v)$  when  $W(u) < W(v)$  and  $L(u) = L(v)$ .

$$C_6: \neg = (L(x_6),L(y_6)) \& \neg < (W(x_6),W(y_6)) \& \neg > (B(x_6),B(y_6))$$

Furthermore, we include clauses that describe a pertinent property of the sum of two quantities.

$$C_7: \neg >(t_1,t_2) \& \neg >(t_3,t_4) \& > (\text{SUM}(t_1,t_3), \text{SUM}(t_2,t_4))$$

Finally, we tell the automated deduction program about a relationship between the functions  $>\{u,v\}$ ,  $=\{u,v\}$ , and  $<\{u,v\}$ .

$$C_6: = \{>(u,v), \text{po} <(u,v) \mid \neg = (u,v)\}$$

Figure 3.3 shows the resulting resolution tree for the goal  $wff$ . The contradiction of the negated goal  $wff$  proves the goal  $wff$  and so the candidate relation is a dominance relation.

In contrast, automated deduction is unable to derive the empty clause for the slightly different candidate relation:

$$= \{L(x), L(y)\} \& >(R(x), R(y)) \& >(W(x), W(y)) \rightarrow \geq(B(x), B(y))$$

Figure 3.4 shows the resolution tree for this relation. The automated deduction system concludes that either the candidate relation is not a dominance relation or we are ignorant of some step in its proof.

These examples illustrate the strengths, weaknesses, and limitations of automated deduction as a tool for finding dominance relations. The use of sound inference rules ensure that an automated deduction program only derives dominance relations which are true when the premises are true. This leads to two very desirable properties for the results from automated deduction: certainty and generality. If we ensure that the set of premises consists only clauses known to be definitely true, then any relation found by the automated deduction program is definitely a dominance relation. Moreover, as the axioms describe the general problem instance, these results are dominance relations for the general problem as opposed to just a subset of problem instances.

Automated deduction makes explicit all required deduction steps, and requires that all information about the problem be explicitly given to the program, including the implicit information that a person tacitly assumes. This can expose questionable or even incorrect assumptions. More importantly, this makes explicit the reasons why a particular relation is a dominance relation and the processes by which the conclusion is reached. Such knowledge is very valuable as it can lead to an understanding and formalization of the deep causal knowledge about the derivation of dominance relations.

On the other hand, being forced to supply such minute and mundane information can often be a nuisance as well as a serious handicap. The absence of a single necessary axiom will prevent the theorem prover from finding a proof of a true goal  $wff$ . This can seriously hinder the ability of automated deduction to find dominance relations in learning situations where there is little formalized knowledge about the domain to which the learning system is being applied.

$$C_0: = \{L(x), L(y)\}$$

$$C_1: \sim =\{L(x), L(y)\} \mid \sim <(W(x), W(y)) \mid >(R(x), R(y))$$

$$C_2: <(W(x), W(y)) \mid \sim >(H(x), H(y))$$

$$C_3: <(W(x), W(y))$$

$$C_4: >(H(x), H(y))$$

$$C_5: \sim >(B(x), B(y))$$

$$C_6: =\{B(x), SUM(R(x), H(x)), SUM(R(y), H(y))\}$$

$$C_7: =\{B(x), SUM(R(x), H(x)), H(x)\}$$

$$C_8: =\{B(x), SUM(R(y), H(y)), H(y)\}$$

$$C_9: >(SUM(R(x), H(x)), SUM(R(y), H(y)))$$

$$C_{10}: >(SUM(R(x), H(x)), H(x))$$

$$C_{11}: >(SUM(R(y), H(y)), H(y))$$

$$C_{12}: >(SUM(R(x), H(x)), SUM(R(y), H(y)))$$

$$C_{13}: >(SUM(R(x), H(x)), H(x))$$

$$C_{14}: >(SUM(R(y), H(y)), H(y))$$

$$C_{15}: >(R(x), R(y)) \mid \sim >(H(x), H(y))$$

$$C_{16}: >(R(x), R(y))$$

$$C_{17}: >(R(x), R(y))$$

$$C_{18}: >(R(x), R(y))$$

$$C_{19}: >(R(x), R(y))$$

$$C_{20}: >(R(x), R(y))$$

$$C_{21}: >(R(x), R(y))$$

$$C_{22}: >(R(x), R(y))$$

$$C_{23}: >(R(x), R(y))$$

$$C_{24}: >(R(x), R(y))$$

$$C_{25}: >(R(x), R(y))$$

$$C_{26}: >(R(x), R(y))$$

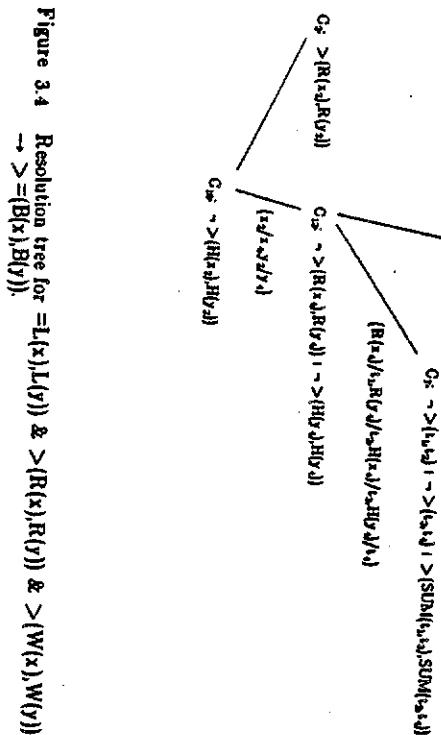
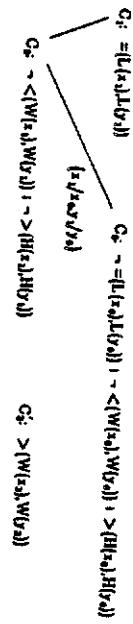
$$C_{27}: >(R(x), R(y))$$

$$C_{28}: >(R(x), R(y))$$

$$C_{29}: >(R(x), R(y))$$

$$C_{30}: >(R(x), R(y))$$

Figure 3.3 Resolution tree for  $=\{L(x), L(y)\} \& >(R(x), R(y)) \& <(W(x), W(y)) \rightarrow >(B(x), B(y))$ .



The user of an automated deduction system must supply three types of information: the inference rules, the axioms representing the problem, and a strategy. The set of inference rule is well-defined and well-known. Thus, the issues that have to be addressed before dominance relations can be learned by deduction are the derivation of the axioms and the strategy. The axioms supplied by the user must include the explicit problem description as well as implicit knowledge about the problem domain and the search algorithm. Deriving the relevant implicit knowledge may be difficult as it often requires deep knowledge about a new, possibly ill-understood problem domain. The strategy is determined by the preferences assigned to the different axioms and inference rules. These preferences are problem-dependent and their derivation is subject to the same difficulties as the derivation of the implicit domain knowledge. These facts include axioms representing the problem description, axioms representing implicit knowledge about the problem domain and general problem solving, and the preferences among the various axioms and the inference mechanisms.

### 3.6.2 Learning Dominance Relations by Analogy

An analogy system finds a known problem domain (base domain) similar to the new problem domain (target domain) and uses analogical mappings between the two domains to transform dominance relations in the base domain

An automated deduction system cannot modify the set of axioms supplied to it. It is necessary to restart the theorem-prover for each new candidate relation. A more serious handicap is that the automated deduction system is restricted to whatever knowledge the program designer has about dominance relation and the problem domain. There is very little causal knowledge about dominance relations, and knowledge about new problem domains may be limited.

It is clear that automated deduction is not sufficient by itself to tackle the problem of finding dominance relations. For one it needs an external mechanism to supply it with candidate relations. It also needs to be coupled with other mechanisms which are able to explore the environment of dominance relations and perhaps automatically modify the set of initial axioms. This should include the ability to innovate and come up with new symbols and not just be constrained to the set of symbols in the initial problem description. The known problem description is often optimized towards state-space methods of solving the problem and may obscure the presence of dominance relations.

Information: the inference rules, the axioms representing the problem, and a strategy. The set of inference rule is well-defined and well-known. Thus, the issues that have to be addressed before dominance relations can be learned by deduction are the derivation of the axioms and the strategy. The axioms supplied by the user must include the explicit problem description as well as implicit knowledge about the problem domain and the search algorithm. Deriving the relevant implicit knowledge may be difficult as it often requires deep knowledge about a new, possibly ill-understood problem domain. The strategy is determined by the preferences assigned to the different axioms and inference rules. These preferences are problem-dependent and their derivation is subject to the same difficulties as the derivation of the implicit domain knowledge. These facts include axioms representing the problem description, axioms representing implicit knowledge about the problem domain and general problem solving, and the preferences among the various axioms and the inference mechanisms.

Figure 3.4 Resolution tree for  $=L(x), L(y) \&gt; R(x), R(y) \&gt; W(x), W(y)$

$\rightarrow \> = (B(x), B(y))$

into relations in the target domain. We can illustrate this process by two single machine scheduling problems: scheduling to minimize weighted completion time and scheduling to minimize cumulative lateness.

Scheduling to minimize the cumulative lateness involves a set of  $n$  tasks:  $t_1, \dots, t_n$ . Each task,  $t_i$ , has a length,  $\ell(t_i)$ , and a deadline,  $d(t_i)$ . The goal in this problem is to find the one-processor schedule  $\sigma()$  that minimizes  $\sum_{i=1}^n [\rho(t_i) + \ell(t_i) - d(t_i)]$ .

Each task,  $t'_i : i = 1, \dots, n$ , in scheduling to minimize weighted completion time has length,  $\ell'(t'_i)$ , and weight,  $w'(t'_i)$ . This problem seeks the one-processor schedule,  $\sigma'()$ , that minimizes  $\sum_{i=1}^n [\rho'(t'_i) + \ell'(t'_i)]w'(t'_i)$ .

Table 3.4 shows the identity-descriptions for these two problems. The decision variable,  $\rho'_i$ , is the identity of the  $i$ -th task in the schedule  $\sigma'()$ . There are two state variables:  $T(N)$  and  $L(N)$ . Let  $s(N)$  be the set of assigned tasks in the partial schedule represented by node,  $N$ .  $T(N)$  is completion time of the most recently scheduled task in the partial schedule represented by node  $N$  while  $L(N)$  is the sum of the set of tasks in  $s(N)$ . Hence,  $\rho'_i$  and  $T(N)$  are identical to  $\rho_i$  and  $T(N)$ .  $C(N)$  is the sum of the weighted completion of the tasks in  $s(N)$ .

Let scheduling to minimize cumulative lateness be the base domain and let scheduling to minimize weighted completion times be the target domain. The matching of the problem descriptions shown in Table 3.4 leads to the following analogical mapping:

$$\begin{array}{lcl} t_i & \rightarrow & t'_i \\ \rho_i & \rightarrow & \text{rho}'_i \\ T(N) & \rightarrow & T'(N) \\ L(N) & \rightarrow & C'(N) \\ \ell(t_i) & \rightarrow & \ell'(t'_i) \\ d(t_i) & \rightarrow & w'(t'_i) \end{array}$$

Scheduling to minimize cumulative lateness is known to have a greedy solution which always schedules the unscheduled task with the smallest lateness when there is only one task with this value of lateness. This greedy solution corresponds to the dominance relation, IF [SIB(x,y)] [L(x) < L(y)] THEN [B(x) ≤ B(y)], where  $B(x)$  is the value of the best solution derived from node  $x$ , and  $SIB(x,y)$  is a predicate which tests whether nodes  $x$  and  $y$  are siblings.

Table 3.4 Identity-descriptions for scheduling to minimize cumulative completion time and scheduling to minimize cumulative lateness.

ENTITY	Scheduling to Minimize Cumulative Lateness	Scheduling to Minimize Weighted Completion Times
DECISION VARIABLE	$\rho_i : \{t_1, \dots, t_n\} \rightarrow \{t'_1, \dots, t'_n\}$	$\rho'_i : \{t'_1, \dots, t'_n\} \rightarrow \{t_1, \dots, t_n\}$
STATE VARIABLES	$T(P) = T(N) + \ell(\rho_i)$ where $P = \text{apply}(\rho_i, N)$	$T'(P) = T'(N) + \ell'(\rho'_i)$ where $P = \text{apply}(\rho'_i, N)$
	$L(P) = L(N) - T(P) + d(\rho_i)$ where $P = \text{apply}(\rho_i, N)$	$C'(P) = C'(N) + T'(P) * w'(\rho'_i)$ where $P = \text{apply}(\rho'_i, N)$
ENTITY VARIABLES	$t(i)$	$t'(i)$
PROBLEM CONSTANTS	$\mathbf{n}$	$\mathbf{n}$
OPTIMIZATION	minimize $L(S)$	minimize $C'(S)$
CONSTRAINT		

Application of the analogical mapping leads to the replacement of  $L(x)$  by  $C'(x)$  and the rule,  $\text{IF } [SIB(x,y)] \text{ [C}(x) < C'(y)] \text{ THEN } [B](x) \leq B(y)$ . This rule is another way of stating the known Greedy scheduling algorithm that minimizes weighted completion times when only one unscheduled task has the smallest completion time.

Analogy can be a very powerful and efficient way of deriving dominance relations when there is a good match between the target domain and the base domain. The known dominance relation in the base domain provides an excellent starting point for the derivation of a new relation. Other learning methods start with much less knowledge. Moreover, the analogical mapping provides powerful heuristic knowledge that drastically reduces the amount of search during the symbolic transformation of the relations.

The most important factor affecting the quality of the process of learning by analogy is the quality of the match between the chosen base domain and the target domain. A poor match requires more "guessing" by the transformation routines and increases the likelihood that the derived rule is incorrect.

Even when there is a good match between the two domains, the rules derived by analogy are heuristics, which may or may not be dominance relations. We cannot be certain that the derived rules are definitely dominance relations because heuristic knowledge is used to find the "best" base domain and to transform the relations in the base domain. Testing of the derived rules with sample problem instance can increase our confidence in the derived rules but not total confidence as testing cannot test all possible problem instances.

The basic mechanism in analogy is the transformation of known relations in an analogous base domain. This mechanism limits the use of analogy to learn dominance relations. Analogy cannot derive rules which are unrelated to known dominance relations so it cannot derive dominance heuristics dissimilar to known dominance relations. Moreover, analogy does not work when there is no suitable base domain. Analogy is not useful in domains which are different from the domains with known dominance relations.

A system that learns dominance relations by analogy must have three types of information: a good set of domains with known dominance relations, routines to measure similarities between domains, and routines to transform relations in the base domain into relations in the target domain. At present, there are only a small number of problems with known dominance relations.

However, this problem will be alleviated as more work is done on dominance relations. Both the issues of similarities between domains and transformations between domains will have to be addressed before analogy can be used to derive dominance relations.

### 3.6.3 Learning Dominance Relations by Experimentation

Learning by experimentation learns by performing experiments on the environment and using the results of these experiments to drive the learning process. This feedback from the environment allows the learning system to adapt its search for dominance relations to the domain under study. Thus, learning by experimentation is not constrained to domains and relations that are similar to known domains and known dominance relations even though it is likely to be more efficient and successful in familiar domains than in strange domains. This graceful degradation permits experimentation to be applied successfully to a much wider set of domains than the other learning mechanisms.

The task of learning dominance relations is a multiple-concept learning task as it is possible for a problem to have more than one dominance relation. To learn dominance relations by experimentation, the system must be able to generate sample problem instances, solve these sample problem instances, and induce possible dominance relations from instances in the search space. A weak problem solver can solve the problem instances. The other two topics are both open questions and must be solved before experimentation can learn dominance relations.

As indicated in Section 3.6, in the absence of an oracle, it is impossible to identify examples of dominance relations among the nodes of the search tree without explicit knowledge of the proposed relation. It is necessary to propose a relation before the system can find examples of this relation. This restriction rules out the use of data-driven methods. Model-driven schema instantiation is also inappropriate because there is no good schema for the dominate relations. Hence, a system that learns dominance relations by experimentation must use model-driven generate-and-test search the space of possible relations (rule space). Model-driven generate-and-test considers one candidate relation at a time, thereby reducing the task of learning dominance relations into a sequence of single-concept learning tasks.

Clearly, heuristic methods have to be used to generate sample problem instances as well as to find dominance relations in the sample search trees. Any relation derived by experimentation is guaranteed to be true only for the set of sample problem instances used in its derivation. The use of fallible heuristic knowledge to generate this set of sample problem instances implies

that experimentation can only find dominance heuristics. There will always be some doubts whether or not the relations found by experimentation are true dominance relations. Some other method of proving the correctness of the dominance heuristics will be necessary.

### 3.8.4 Combining the Different Learning Mechanisms

In the previous sections, we have seen that both analogy and deduction may not be useful for all problems. Analogy is only useful in problem domains similar to domains with known dominance relations. Deduction is only useful in well-known problems because of the need for the designer of the learning system to supply all information relevant to the derivation of the dominance relations. A learning system will have to work in an unfamiliar environment.

Only experimentation works well in these situations. Thus, experimentation must be the primary mechanism that drives a system for learning dominance relations. Analogy and deduction are secondary mechanisms; they are necessary as they can be very effective in the proper situations.

Figure 3.5 shows the different learning mechanisms. Both analogy and experimentation derive dominance heuristics. Deduction derives dominance relations. Any relation proved by deduction logically follows from the axioms describing the problem as a true dominance relation. Hence, one way of verifying the dominance heuristic is to use deduction to try to prove the dominance heuristics. This scheme also provides a possible solution to the problem of proposing suitable candidate relations for automated deduction.

The process of learning dominance relations by each learning mechanisms may be viewed as the execution of a plan by the learning mechanism. The efficiency of the learning mechanisms can be enhanced by augmenting their general plans with domain specific knowledge. One way is to use analogy on similar domains that possess good plans. Obviously, this aspect of the learning system has to wait until each of the learning mechanisms has been sufficiently well developed so that there is an adequate pool of plans. Not all the steps in a plan are useful so it will be necessary to devise mechanisms to remove those steps which are not useful before adding the plan to the pool of known plans.

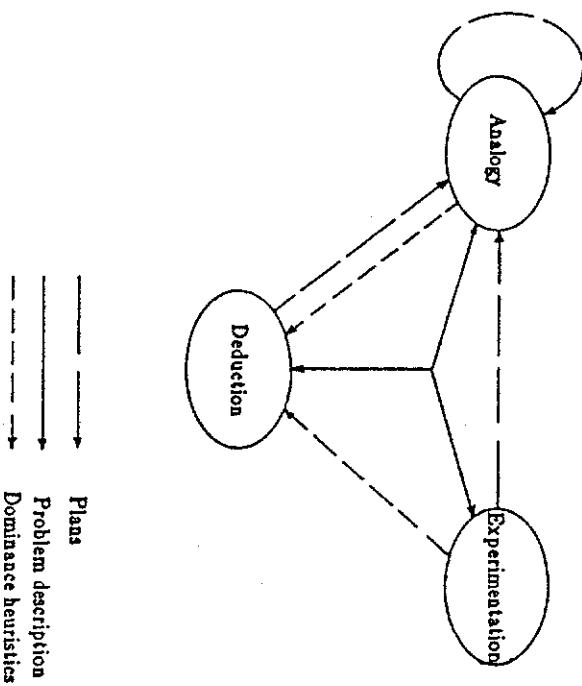


Figure 3.5 Interactions between the learning mechanisms

This chapter presents the methodology used to design a system for learning dominance relations as well as some results of applying this methodology. Among the results are the representations for the dominance relations and the number optimization problems, a classification of dominance relations, and the applicable learning mechanisms.

The dominance relation can be naturally expressed as either a production rule or the implication statement of predicate calculus. The similarities between the descriptions of different number optimization problems indicate the presence of a regular structure in the problem descriptions. This structure may be expressed in terms of a complete set of domain-independent elements. A suitable representation for the number optimization problems is the frame. Slots in the frame represent the set of elements. The slot fillers represent the values of these elements in the particular problem description.

Dominance relations may be classified by the permissible levels of the dominating and dominated nodes. This classification is able to expose those dominance relations which are either trivial relations or relations requiring excessive implementation overheads. These considerations suggest that the most useful types of dominance relations are those between nodes on the same level in the search tree, and so the study of learning dominance relations can be restricted to such dominance relations.

Dominance relations can be learned by analogy, deduction, and experimentation. Both analogy and deduction perform poorly in domains which are dissimilar to the domains used to derive their knowledge bases. Experimentation receives feedback from the environment, and can be effective in unfamiliar domains. Hence, experimentation is the primary learning mechanism, while analogy and deduction are secondary learning mechanisms. Nonetheless, both analogy and deduction are important. Analogy can be very effective in domains similar to known domains, while deduction is the only mechanism that guarantees the correctness of its learned relations.

## CHAPTER 4 LEARNING DOMINANCE RELATIONS BY EXPERIMENTATION

### 4.1 Introduction

As indicated in the previous chapter, dominance relations may be learned by analogy, deduction and experimentation. The next step in our methodology is to refine the individual learning mechanisms by developing the procedures and knowledge that implement the learning mechanism.

Learning by deduction only learns relations that can be deduced from the facts known to the deduction program. It is most useful in domains where the relevant facts are known. Learning by analogy transforms a known dominance relation into a relation in the new problem domain. It is most useful when the new problem domain is similar to a domain with known dominance relations. Learning by experimentation is the most general mechanism as the feedback that it receives from its experiments on the environment permit it to learn relations in unfamiliar problem domains and unfamiliar types of relations. Thus, learning by experimentation is the primary learning mechanism for dominance relations.

In the work of this thesis, we have studied the refinement of learning dominance relations by experimentation. This study has led to the development of MUFT-1: a prototype program that is able to learn dominance relations by experimentation. MUFT-1 uses the model-driven generate-and-test paradigm because it is not possible to identify examples of the dominance relation without explicit knowledge of the dominance relation. Generate-and-test is a well-known general paradigm in artificial intelligence [Wing81] and it has been applied to the task of learning cleavage rules for DENDRAL's mass-spectrometer simulator [Bau78] and to the task of learning structural descriptions [Dik83]. However, the key to a successful generate-and-test scheme is the domain-specific model which holds the domain knowledge. The models used for learning cleavage rules and structural descriptions cannot be used for learning dominance relations.

### 3.7 Summary

The design of MUTT-I is described in Sections 4.2 to 4.7. Section 4.8 summarizes the results of applying MUTT-I to several test problems.

#### 4.2 MUTT-I

Figure 4.1 shows the organization of MUTT-I. The searcher implements the generate-and-test search of the space of candidate relations. A candidate relation is accepted as a dominance heuristic if the searcher determines that this relation behaves like a dominance relation during tests with sample search trees. The forester is responsible for the supply of appropriate sample search trees.

The searcher continues to test candidate relations until it either exhausts its resource allocation or finds a dominance heuristic; it then suspends itself and the scheduler takes over. The scheduler decides whether to stop or to continue searching for dominance heuristics. If the scheduler makes the latter decision, it will use the description manager to evaluate the effectiveness of the current description and if necessary make a description switch before restarting the searcher with a new resource allocation.

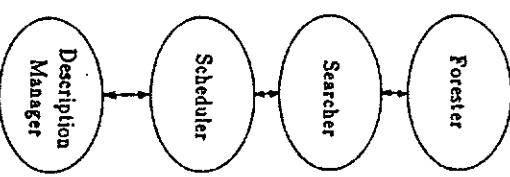
Details of the individual components will be found in the sections that follow this section. Section 4.3 describes the scheduler, Section 4.4 describes the description manager, and Section 4.5 describes the forester. The design of the searcher depends on the representation of the candidate relations so we will discuss the representation of candidate relations in Section 4.6 before describing the searcher in Section 4.7. These descriptions are illustrated with examples drawn mainly from the application of MUTT-I to the 0/1-knapsack problem. Examples from other problems are used only when there is no appropriate examples for the 0/1-knapsack, and these cases will be explicitly noted in the descriptions.

#### 4.3 Scheduler

The task of the scheduler of the scheduler is to coordinate the activities of the other components of the system. It incorporates the top-level control structures of the system.

One of its duties is to decide if there is sufficient justification for MUTT-I to continue to search for dominance heuristics. One event that leads to the activation of the scheduler is the discovery of a dominance heuristic by the searcher. Whenever this event occurs, the scheduler tests whether the

Figure 4.1 Organization of MUTT-I



dominance heuristic is a complete solution to the problem under study, that is, it corresponds to a greedy solution of the problem. An affirmative answer causes the scheduler to stop; otherwise the search for dominance heuristics continues.

When the scheduler is activated by the exhaustion of the searcher's resource allocation, the searcher makes its decision on the recent progress of the searcher. The scheduler measures the searcher's progress by  $\tau$ , the amount of processing time consumed by the searcher since the most recent dominance heuristic was found or the start of the program if no dominance heuristic has yet been found. The scheduler also has a preset threshold,  $\tau_0$ , which represents the maximum amount of processing that the user is willing to expand without finding a dominance heuristic. If  $\tau$  is less than  $\tau_0$ , the scheduler will decide to continue searching for dominance heuristics. If  $\tau$  is greater than or equal to  $\tau_0$ , the scheduler will query the user for guidance on whether to continue or stop. Values between 10 seconds and 100 seconds are suitable.

If the searcher decides to continue searching, it will allocate a new resource slice to the searcher after the description manager ensures that the current description is the correct. The size of the time slice determines the frequency with which the description manager evaluates the effectiveness of the current description. MUTT-1 uses fixed sized time slice of size  $\kappa$ . Let  $\eta$  be the number of working descriptions for the problem and let  $\kappa_0$  be a preset value. Then,  $\kappa$  has the definition:

$$\kappa = \frac{\tau_0}{\eta \times \kappa_0} \quad (4.1)$$

$\kappa_0$  represents the average number of evaluations for each working description during a time period of duration  $\tau_0$ . Two contradictory factors influence the choice of  $\kappa_0$ 's value.  $\kappa_0$  should be small to maximize the time actually spent on searching for dominance heuristics. However,  $\kappa_0$  should be large to minimize the likelihood of wasting processing time on an inappropriate working description. Values between 5 and 10 are suitable for  $\kappa_0$ .

The time slice in MUTT-1 is a soft limit; the searcher relinquishes control at the earliest convenient moment after it exceeds the time limit. This simplifies the design and implementation of the time slice mechanism.

#### 4.4 Description Manager

The ideal description manager is able to automatically derive the optimal working description from the problem statement. This capability requires deep causal knowledge, about problem descriptions and the problem domain, and is a major open research topic in artificial intelligence.

The design of MUTT-1 does not address this issue. MUTT-1 uses a a dumb description manager. The user is required to supply the appropriate working descriptions. Furthermore, the user must ensure that these working descriptions are correct and complete. The use of a dumb manager that relies totally on the user is acceptable for a prototype program used as knowledge test-bed, but not for a delivery system. Automated derivation of working descriptions is an issue that must be addressed during the refinement of this prototype program.

The user supplies a working description to MUTT-1 in the form of a w-frame. The w-frame is an extension of the description frame introduced in Section 3.4. It augments the basic problem description with information that assists MUTT-1 in manipulating the problem description. Figure 4.2 shows a stylized rendition of the frame that describes the value-description for the 0/1-knapsack problem. The Typical slot defines typical values of each entity variable and each problem constant. It includes a LISP function that returns a typical value. These LISP functions may use one or more of the additional LISP functions stored in the LISP-Fn slot of the w-frame.

The set of w-frames representing the working descriptions,  $\{\psi_i : i = 1, \dots, n\}$ , are linked together into a description list which is ordered by increasing values of  $\epsilon(\psi_i)$ .  $\epsilon(\psi_i)$  is a heuristic measure that estimates the average time required to find a dominance relation in description  $\psi_i$ . It is defined by the formula:

$$\epsilon(\psi_i) = \begin{cases} \frac{T(\psi_i)}{N(\psi_i)} & N(\psi_i) > 0 \\ T(\psi_i) & N(\psi_i) = 0 \end{cases} \quad (4.2)$$

where  $T(\psi_i)$  is the total computation time consumed by the searcher while using working description  $\psi_i$ , and  $N(\psi_i)$  is the number of dominance heuristic found by the searcher in working description  $\psi_i$ .

When the description manager is activated by the scheduler, it updates the  $\epsilon$ -value of the current description and compares this  $\epsilon$ -value with the  $\epsilon$ -value of the next description in the description list. If the current description

```

Name: w1
Type: w-frame
Problem: 0/1-knapsack
Description-Type: value-representation

Entities:
  Name: r1
  Number: 8
  Entity Variables:
    Entity Variable 1:
      Name: r1
      Value Type: integer
      Range: >0
      Number: 8
      Typical:
        Value: [1..10]
        Code: {+ 1 {.. (mod 9)}}
  Problem Constants:
    Name: a
    Value Type: integer
    Range: >0
    Typical:
      Value: [2..10]
      Code: {+ 2 {.. (mod 8)}}
  State Variables:
    State Variable 1:
      Name: R1
      Definition:
        Value: R(Np) = R(Np) +
           $\sum_{i=1}^r x_i * w_i$ 
        Code: lispf01
      Initial Value: 0
      Constraints:
        Value: R(Np) ≤ K
        Code: lispf02
      Optimization:
        Expression: R(Nmax)
        Extreme Value: maximum
        Code: lispf03
      Decision Variable:
        Name: x1
        Value: [0..1]
        LISP Functions: lispf00, lispf01, lispf02, lispf03, lispf04

```

The forester supplies test trees to the searcher. The same test tree may be used to test many different relations; thus, the forester retains all previously-generated test trees in a tree pool. The forester always tries to satisfy a request from its tree pool. It will have to generate new trees if either the pool has been exhausted or the trees in the pool do not meet the searcher's demands. These new trees are added to the tree pool for use with future candidate relations.

The first step in generating a new tree is to generate a new problem instance. A new problem instance is generated by creating an instance of the p-frame (problem instance frame) and filling in the slots. The slots in the p-frame describe the structure of each problem instance for the problem under study and how to fill in these slots. The forester derives the p-frame from the w-frame by extracting the relevant sub-slots from the slots describing the problem constants and the entity variables in the w-frame. The p-frame specifies the manner in which values can be assigned to the slots of an instance of the p-frame. Figure 4.3(a) shows the p-frame for the 0/1-knapsack problem. Figure 4.3(b) shows an instance of the p-frame representing the 0/1-knapsack instance,  $\Pi_1$ .

Once a new problem instance has been generated, the forester uses a weak problem solver to solve the problem instance. Table 4.1 shows the search tree,  $T_1$ , generated by the solution of the problem instance,  $\Pi_1$ , shown in Figure 4.3(b). This search tree is stored in the Search-Tree slot of  $\Pi_1$ .

#### 4.6 Forester

The description manager initiates a description switch. The first step in a description switch is to save the current state of the searcher into an appropriate file. The description manager then reorganizes the description list to maintain the ordering by c-values. Finally, it restores the searcher to the state stored in the file corresponding to the new current description.

Figure 4.2 Frame describing value-description of 0/1-knapsack problem.

#### 4.6 Representation of Candidate Relations

MUTT-I represents candidate relations by c-rules which are production rules of the following form:

Table 4.1 Search tree  $r_1$ .

$N$  UNI RNL WIN RNL Sizing Punkt Objekt

$N$	UNI	RNL	WIN	RNL	Sizing	Punkt	Objekt
1	z	z	z	z	z	z	$N_1, N_2$
2	z	z	z	z	z	z	$N_2, N_3$
3	z	z	z	z	z	z	$N_3, N_4$
4	z	z	z	z	z	z	$N_4, N_5$
5	z	z	z	z	z	z	$N_5, N_6$
6	z	z	z	z	z	z	$N_6, N_7$
7	z	z	z	z	z	z	$N_7, N_8$
8	z	z	z	z	z	z	$N_8, N_9$
9	z	z	z	z	z	z	$N_9, N_{10}$
10	z	z	z	z	z	z	$N_{10}, N_{11}$
11	z	z	z	z	z	z	$N_{11}, N_{12}$
12	z	z	z	z	z	z	$N_{12}, N_{13}$
13	z	z	z	z	z	z	$N_{13}, N_{14}$
14	z	z	z	z	z	z	$N_{14}, N_{15}$
15	z	z	z	z	z	z	$N_{15}, N_{16}$
16	z	z	z	z	z	z	$N_{16}, N_{17}$
17	z	z	z	z	z	z	$N_{17}, N_{18}$
18	z	z	z	z	z	z	$N_{18}, N_{19}$
19	z	z	z	z	z	z	$N_{19}, N_{20}$
20	z	z	z	z	z	z	$N_{20}, N_{21}$
21	z	z	z	z	z	z	$N_{21}, N_{22}$
22	z	z	z	z	z	z	$N_{22}, N_{23}$
23	z	z	z	z	z	z	$N_{23}, N_{24}$
24	z	z	z	z	z	z	$N_{24}, N_{25}$
25	z	z	z	z	z	z	$N_{25}, N_{26}$
26	z	z	z	z	z	z	$N_{26}, N_{27}$
27	z	z	z	z	z	z	$N_{27}, N_{28}$
28	z	z	z	z	z	z	$N_{28}, N_{29}$
29	z	z	z	z	z	z	$N_{29}, N_{30}$
30	z	z	z	z	z	z	$N_{30}, N_{31}$
31	z	z	z	z	z	z	$N_{31}, N_{32}$
32	z	z	z	z	z	z	$N_{32}, N_{33}$
33	z	z	z	z	z	z	$N_{33}, N_{34}$
34	z	z	z	z	z	z	$N_{34}, N_{35}$
35	z	z	z	z	z	z	$N_{35}, N_{36}$
36	z	z	z	z	z	z	$N_{36}, N_{37}$
37	z	z	z	z	z	z	$N_{37}, N_{38}$
38	z	z	z	z	z	z	$N_{38}, N_{39}$
39	z	z	z	z	z	z	$N_{39}, N_{40}$
40	z	z	z	z	z	z	$N_{40}, N_{41}$
41	z	z	z	z	z	z	$N_{41}, N_{42}$
42	z	z	z	z	z	z	$N_{42}, N_{43}$
43	z	z	z	z	z	z	$N_{43}, N_{44}$
44	z	z	z	z	z	z	$N_{44}, N_{45}$
45	z	z	z	z	z	z	$N_{45}, N_{46}$
46	z	z	z	z	z	z	$N_{46}, N_{47}$
47	z	z	z	z	z	z	$N_{47}, N_{48}$
48	z	z	z	z	z	z	$N_{48}, N_{49}$
49	z	z	z	z	z	z	$N_{49}, N_{50}$
50	z	z	z	z	z	z	$N_{50}, N_{51}$
51	z	z	z	z	z	z	$N_{51}, N_{52}$
52	z	z	z	z	z	z	$N_{52}, N_{53}$
53	z	z	z	z	z	z	$N_{53}, N_{54}$
54	z	z	z	z	z	z	$N_{54}, N_{55}$
55	z	z	z	z	z	z	$N_{55}, N_{56}$
56	z	z	z	z	z	z	$N_{56}, N_{57}$
57	z	z	z	z	z	z	$N_{57}, N_{58}$
58	z	z	z	z	z	z	$N_{58}, N_{59}$
59	z	z	z	z	z	z	$N_{59}, N_{60}$
60	z	z	z	z	z	z	$N_{60}, N_{61}$
61	z	z	z	z	z	z	$N_{61}, N_{62}$
62	z	z	z	z	z	z	$N_{62}, N_{63}$
63	z	z	z	z	z	z	$N_{63}, N_{64}$
64	z	z	z	z	z	z	$N_{64}, N_{65}$
65	z	z	z	z	z	z	$N_{65}, N_{66}$
66	z	z	z	z	z	z	$N_{66}, N_{67}$
67	z	z	z	z	z	z	$N_{67}, N_{68}$
68	z	z	z	z	z	z	$N_{68}, N_{69}$
69	z	z	z	z	z	z	$N_{69}, N_{70}$
70	z	z	z	z	z	z	$N_{70}, N_{71}$
71	z	z	z	z	z	z	$N_{71}, N_{72}$
72	z	z	z	z	z	z	$N_{72}, N_{73}$
73	z	z	z	z	z	z	$N_{73}, N_{74}$
74	z	z	z	z	z	z	$N_{74}, N_{75}$
75	z	z	z	z	z	z	$N_{75}, N_{76}$
76	z	z	z	z	z	z	$N_{76}, N_{77}$
77	z	z	z	z	z	z	$N_{77}, N_{78}$
78	z	z	z	z	z	z	$N_{78}, N_{79}$
79	z	z	z	z	z	z	$N_{79}, N_{80}$
80	z	z	z	z	z	z	$N_{80}, N_{81}$
81	z	z	z	z	z	z	$N_{81}, N_{82}$
82	z	z	z	z	z	z	$N_{82}, N_{83}$
83	z	z	z	z	z	z	$N_{83}, N_{84}$
84	z	z	z	z	z	z	$N_{84}, N_{85}$
85	z	z	z	z	z	z	$N_{85}, N_{86}$
86	z	z	z	z	z	z	$N_{86}, N_{87}$
87	z	z	z	z	z	z	$N_{87}, N_{88}$
88	z	z	z	z	z	z	$N_{88}, N_{89}$
89	z	z	z	z	z	z	$N_{89}, N_{90}$
90	z	z	z	z	z	z	$N_{90}, N_{91}$
91	z	z	z	z	z	z	$N_{91}, N_{92}$
92	z	z	z	z	z	z	$N_{92}, N_{93}$
93	z	z	z	z	z	z	$N_{93}, N_{94}$
94	z	z	z	z	z	z	$N_{94}, N_{95}$
95	z	z	z	z	z	z	$N_{95}, N_{96}$
96	z	z	z	z	z	z	$N_{96}, N_{97}$
97	z	z	z	z	z	z	$N_{97}, N_{98}$
98	z	z	z	z	z	z	$N_{98}, N_{99}$
99	z	z	z	z	z	z	$N_{99}, N_{100}$

Figure 4.3

(a) The p-frame for the 0/1-knapsack problem, and (b) the p-

Name:  $\Pi_1$   
Type: p-frame  
Problem: 0/1-knapsack

n:

Value:

Code: lispfn10

w:

Value:

Code: lispfn11

K:

Value:

Code: lispfn12

Search-Tree: nil

Lisp Functions: lispfn10, lispfn11, lispfn12

(a)

(b)

IF      t-condition  
d-condition<sub>i</sub>  
d-condition<sub>m</sub>

THEN    x dominates y

A candidate relation is a hypothesis that if two nodes satisfy the conditions described in the hypothesis then we can always identify the node with the better solution. This hypothesis can only be accepted if all nodes which satisfy the conditions in the hypothesis also agrees with its identification of the node with the better solution.

The left-hand side of the c-rule identifies the conditions which the nodes must satisfy in order for node x to dominate node y. It consists of a t-condition (type condition) and one or more d-conditions (dominance conditions). The t-condition identifies the type of the candidate relation represented by the rule, that is, the set of nodes to which this c-rule may be applied. Let L(x) be the level of node x and SIB(x,y) be a logical predicate which returns the value true if and only if x and y are siblings. The rule representing a candidate relation of type {all, same, sib} has [SIB(x,y)] as its t-condition. The term c-rule will be used to refer to a c-rule of this form. A candidate relation of type {all, same, all} has a c-rule with [L(x)=L(y)] as its t-condition. The term β-rule denotes a c-rule of this form.

The d-conditions establish the conditions which nodes x and y must satisfy in order for x to dominate y. Each d-condition is a comparison between two symbols. Possible choices for the comparison operator are = (equal to), > (greater than), and < (less than). Possible symbols include state variables, entity variables, problem constants, functions of other symbols, and numbers.

The initial set of symbols for the 0/1-knapsack problem is {R(x), R(y), W(x), W(y), K, r(x), r(y), w(x), w(y)}. State variables have different values for different nodes, so it is necessary to use the symbol representing the state variable value of a particular node. Each state variable gives rise to two symbols, representing the values of the state variable for the two nodes involved in the dominance relation. Hence, the two state variables, R<sub>1</sub> and W<sub>1</sub>, contributes four symbols, R(x), R(y), W(x), and W(y), to the initial set of symbols. It is possible to include the entity variable in a node by associating a node with the entity involved in the decision leading to the generation of the node. This convention leads to the symbols, r(x), r(y), w(x), and w(y). Problem constants are independent of the nodes. There is no need to associate

them with the nodes, so problem constant K can be directly added to the list of symbols.

The right-hand side of each c-rule states that node x dominates node y. A c-rule casts this statement as a relationship between the values of the best solutions in the two nodes, i.e. B(x) and B(y). By definition, B(x) is either smaller than or equal to B(y) for a minimization problem, that is, [B(x) ≤ B(y)].

A c-rule is an acceptable rule if it represents a dominance relation. Those c-rules which do not represent a dominance relation may be divided into unacceptable rules and impossible rules. An unacceptable rule is a c-rule whose set of conditions (left-hand side) cannot consistently identify the dominating node and the dominated node. An impossible rule is an c-rule which is impossible, that is, no search tree contains a pair of nodes which matches its left hand side.

A node pair (N<sub>i</sub>, N<sub>j</sub>) is an instance of a c-rule if the substitutions, x/N<sub>i</sub> and y/N<sub>j</sub>, cause the left-hand side of the c-rule to be true. Instance (N<sub>i</sub>, N<sub>j</sub>) is a positive instance if B(N<sub>i</sub>) ≥ B(N<sub>j</sub>) and a negative instance if B(N<sub>i</sub>) < B(N<sub>j</sub>). A positive instance supports the c-rule. A negative instance contradicts the c-rule.

Associated with each c-rule is a δ-value which reflects its suitability for use as a dominance relation. Initially, the δ-values of the c-rules are not known. The searcher finds the δ-values of c-rules by testing with search trees. If testing finds only positive instances of the c-rule, then the c-rule has a δ-value of acceptable. If testing finds one or more negative instances of the c-rule, then the c-rule has a δ-value of unacceptable. If testing fails to find any instance of the c-rule, then the c-rule has a δ-value of impossible.

Consider the c-rule, IF [SIB(x,y)] [R(x)>R(y)] THEN [B(x)≥B(y)] (rule α<sub>2</sub>' in Figure 4.4). (N<sub>2</sub>, N<sub>1</sub>) in search tree r<sub>1</sub> (Table 4.1) is an instance of this c-rule because N<sub>2</sub> and N<sub>1</sub> are siblings and R(N<sub>2</sub>) is greater than R(N<sub>1</sub>). However, B(N<sub>2</sub>) is smaller than B(N<sub>1</sub>), so (N<sub>2</sub>, N<sub>1</sub>) is a negative instance of α<sub>2</sub>'. In contrast, B(N<sub>3</sub>) exceeds B(N<sub>1</sub>), and instance (N<sub>1</sub>, N<sub>3</sub>) is a positive instance of α<sub>2</sub>'. The presence of both positive and negative instances of α<sub>2</sub>' results in the assignment of a δ-value of unacceptable to α<sub>2</sub>. A c-rule with a δ-value of unacceptable has demonstrated that its left-hand side is unable to consistently identify the node with the better solution, and is an unacceptable c-rule.

Testing will not find any instances of the c-rule, IF [SIB(x,y)] [R(x)=R(y)] THE N [B(x)≥B(y)] (c-rule α<sub>4</sub>), and its δ-value will be impossible. A c-rule with

a  $\delta$ -value of impossible is likely to be an impossible rule, but it may also be an acceptable rule or an unacceptable rule. The ambiguity in the actual nature of the rule comes from the use of a finite number of search trees during testing; it is impossible to test the c-rule against all possible search trees. Testing may have missed the search trees that contain positive and negative instances of the c-rule. An example of this will be the true c-rule, IF  $[L(x)=L(y)] \wedge [R(x)=R(y)] \wedge [W(x)=W(y)]$  THEN  $[B(x) \geq B(y)]$  (rule  $\beta'_3$ ).  $\beta'_3$  is an acceptable rule, but search tree  $r_1$  contains no instance of this rule.

An example of a c-rule with a  $\delta$ -value of acceptable is the rule, IF  $[L(x)=L(y)] \wedge [R(x)>R(y)] \wedge [W(x)<W(y)]$  THEN  $[B(x) \geq B(y)]$  (rule  $\beta'_6$ ).  $\beta'_6$  is able to consistently identify the node with the better solution in the test set of search trees. Thus, it is most likely to be an acceptable c-rule, but may also be an unacceptable rule. As before, this ambiguity is due to the use of a limited sample of search trees during testing. It is possible that there are search trees containing negative instances but are not included in the sample set.

Learning by experimentation can only find the  $\delta$ -values of the c-rules. It is unable to find the actual nature of the c-rule. The design of MURT-1 assumes that an adequate sample set of test trees is sufficient justification for accepting the  $\delta$ -value of a c-rule as the actual nature of the c-rule. Henceforth, we will use interchangeably the terms, acceptable rule and rule with a  $\delta$ -value of acceptable. The terms, unacceptable (impossible) rule and rule with a  $\delta$ -value of unacceptable (impossible), will be treated similarly.

#### 4.6.1 Rule Space

The process of finding the  $\delta$ -values of c-rules may be viewed as a search in the space of possible c-rules. One possible organization of this rule space is a simple directed graph. There are no loops in a simple directed graph so we can use the terminology commonly associated with trees such as parents, children, ancestors, and descendants. Figure 4.4 shows a small portion of the rule graph for the 0/1-knapsack problem. The root of this graph is the null rule, IF [true] THEN  $[B(x) \geq B(y)]$ , and is a dummy vertex inserted to complete the graph. All other vertices represent c-rules. Each directed edge represents the addition of a condition. The terminal vertex of each edge represents the c-rule formed by adding the condition represented by the edge to the c-rule represented by its initial vertex.

The children of a c-rule in this rule graph represent specializations of the c-rule and so there is inheritance of  $\delta$ -values by descendant c-rules and ancestor

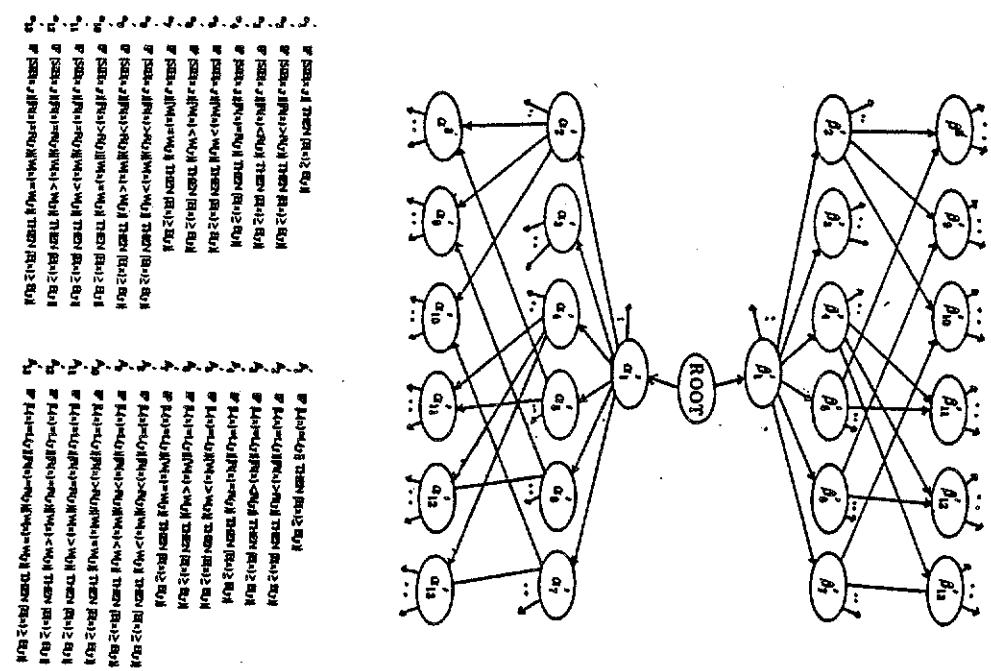


Figure 4.4 A fragment of the rule graph for the 0/1-knapsack problem.

c-rules. Section 4.6.2 discusses this topic in greater detail.

The rule graph contains both  $\alpha$ -rules and  $\beta$ -rules. The  $\alpha$ -rule is a c-rule with  $[SIB(x,y)]$  as its t-condition, and the  $\beta$ -rule is a c-rule with  $[L(x)=L(y)]$  as its t-condition. For each  $\alpha$ -rule, there is a  $\beta$ -rule with an identical right-hand side and an identical set of d-conditions, and vice-versa. We refer to these two rules as parallel rules. The c-rules in Figure 4.4 have been labeled such that rule  $\alpha_i$  and rule  $\beta'_i$  are parallel rules. Parallel rules only differ in their t-conditions which implies a possible linkage of their  $\delta$ -values. Section 4.6.3 explores parallelism in greater detail.

#### 4.6.2 Value Inheritance

The parent of a c-rule represents a generalization of the c-rule. Dropping a d-condition from a c-rule transforms the c-rule into one of its parent rules. For instance, the elimination of condition  $[R(x)>R(y)]$  from the c-rule  $\alpha_3'$  produces the c-rule  $\alpha_3$ , while the elimination of condition  $[W(x)>W(y)]$  produces the c-rule  $\alpha_3''$ . Both  $\alpha_3'$  and  $\alpha_3''$  are parents of  $\alpha_3$ .

Suppose that  $(N_i, N_j)$  is a positive instance of the  $\alpha_3$ . This assumption implies that the clause  $\{[SIB(N_i, N_j)] \& [R(N_i)>R(N_j)] \& [W(N_i)>W(N_j)]\}$  is true. As a conjunction of predicates is true if and only if all the predicates are true, the clause formed by removing one or more conditions from the above clause will still be true. Removing the predicate  $[W(N_i)>W(N_j)]$  results in the true clause  $\{[SIB(N_i, N_j)] \& [R(N_i)>R(N_j)] \& [B(N_i)\geq B(N_j)]\}$ . This new clause implies that  $(N_i, N_j)$  will also be a positive instance of  $\alpha_3$ . By a similar argument, we can show that  $(N_i, N_j)$  is also a positive instance of  $\alpha_3''$ . This example illustrates the inheritance of instances by the parents of a c-rule. Every positive (negative) instance of a c-rule will be a positive (negative) instance of its parent c-rule.

By definition, there are no instances of an impossible rule. For this condition to be satisfied, the impossible rule cannot inherit any instance from its children, that is, none of its children can have any instances. Hence, all children of an impossible rule must also be impossible rules.

An acceptable rule has positive instances but no negative instance. The latter condition implies that none of its children can have negative instances. Each child of an acceptable rule is either another acceptable rule or an impossible rule. However, there must be at least one acceptable rule among the children of an acceptable rule. There will be no positive instances of a c-rule if none of its children has any positive instances.

There are negative instances of an unacceptable rule. Inheritance of instances implies that one or more of the children of the unacceptable rule have negative instances. There may also be positive instances of the unacceptable rule. Thus, it is possible for a child of an unacceptable rule to have no instances, only positive instances, only negative instances, or both positive and negative instances; so a child of an unacceptable rule can be another unacceptable rule, an acceptable rule, or an impossible rule.

#### 4.6.3 Parallelism

Siblings are nodes on the same level with the same parent, and  $[SIB(N_i, N_j)]$  implies  $[L(N_i)=L(N_j)]$ . Parallel rules have the same set of d-conditions and right-hand side. A positive (negative) instance of an  $\alpha$ -rule is also a positive (negative) instance of its parallel  $\beta$ -rule. The node pair  $(N_2, N_1)$  is a positive instance of the  $\alpha$ -rule  $\alpha_1'$ , so the clause  $\{[SIB(N_2, N_1)] \& [R(N_2)>R(N_1)] \& [B(N_2)\geq B(N_1)]\}$  is a true clause. Replacing  $[SIB(N_2, N_1)]$  by its implied condition  $\{[L(N_2)=L(N_1)]\}$  yields  $\{[R(N_2)>R(N_1)] \& [B(N_2)>=B(N_1)]\}$ . Hence,  $(N_2, N_1)$  is also a positive instance of the  $\beta$ -rule  $\beta_1'$ .

A  $\beta$ -rule is a c-rule between two equilevel nodes. Node pair  $(N_i, N_j)$  is a sibling pair if  $N_i$  and  $N_j$  are siblings and a non-sibling pair if  $N_i$  and  $N_j$  are not siblings. Siblings have a common parent, while non-sibling equilevel nodes do not. As a result, siblings exhibit greater correlation in the values of their state variables than do non-siblings. Hence, the absence of positive (negative) instances among the sibling pairs does not imply the absence of positive (negative) instances among the non-sibling pairs. For example, there are no instances of the rule IF  $[SIB(x,y)]$   $[R(x)=R(y)]$  THEN  $[B(x)\geq B(y)]$  (rule  $\alpha_4'$ ) but there are positive and negative instances of the rule IF  $[L(x)=L(y)]$   $[R(x)=R(y)]$  THEN  $[B(x)\geq B(y)]$  (rule  $\beta_4'$ ). The absence of instances of  $\alpha_4'$  implies the absence of instances of  $\beta_4'$  among the sibling pairs. We know that there are instances of  $\beta_4'$ , so there must be positive instances and negative instances of  $\beta_4'$  among the non-sibling pairs.

Nonetheless, it is still possible to deduce something about the possible instances of a c-rule from the instances of its parallel c-rule. Table 4.2 summarizes the relationships between the  $\delta$ -values of parallel rules. Let  $\alpha_i'$  and  $\beta'_i$  be a pair of parallel rules, where  $\alpha_i'$  is an  $\alpha$ -rule and  $\beta'_i$  is a  $\beta$ -rule. Suppose that  $\alpha_i'$  is an unacceptable c-rule. Its negative instances are also negative instances of  $\beta'_i$ . Any c-rule with one or more negative instances in an unacceptable c-rule, so  $\beta'_i$  must be an unacceptable rule.

Suppose that  $\alpha'_i$  is an acceptable c-rule. Then, there will be positive instances of  $\alpha'_i$  and no negative instances of  $\alpha'_i$ . The positive instances of  $\alpha'_i$  are also positive instances of  $\beta'_i$ . However, the absence of negative instances of  $\alpha'_i$  does not imply the absence of negative instances of  $\beta'_i$ . There may or may not be negative instances of  $\beta'_i$ . Hence,  $\beta'_i$  may either be an acceptable c-rule or an unacceptable c-rule.

Suppose that  $\alpha_i$  is an impossible c-rule. There are no positive instances of  $\alpha'_i$  and no negative instances of  $\alpha'_i$ . Neither of this facts provide any useful information regarding the instances of  $\beta'_i$ , and  $\beta'_i$  may be an acceptable rule, an unacceptable rule, or an impossible rule. Rule  $\alpha_i$  has a  $\delta$ -value of impossible, but its parallel rule,  $\beta$ , is an unacceptable rule. In contrast, rule  $\alpha_i$  is impossible, while rule  $\alpha'_i$  is acceptable.

#### 4.7 Searcher

As indicated in Section 4.2, the searcher implements the generate-and-test scheme used by MUTH-I to search the rule space. A generate-and-test scheme consists of two basic modules: a generator and a tester. The generator enumerates possible relations. The tester evaluates each proposed relation, either accepting it or rejecting it. There are two types of generators: the exhaustive generator and the informed generator. The simpler exhaustive generator generates all possible rules while an informed generator uses possibility-limiting information to restrict the relations which it proposes.

A c-rule has a unique set of symbol pairs if each of the d-conditions compares two different symbols and no two d-conditions compares the same pair of symbols. Consider a c-rule with m d-conditions. Suppose that the symbols in the d-conditions are chosen from a pool of n different symbols. The number of possible symbol pairs is  $C_2^m = m(m-1)/2$ . There are m d-conditions in the c-rule so there are  $C_m^{m(m-1)/2}$  ways of choosing the m symbol pairs for the c-rule. Each set of m symbol pairs corresponds to  $3^m$  possible rules because there are three possible choices for the comparison operator. Hence, the number of possible c-rules,  $n_c$ , is

$$n_c = C_m^{m(m-1)/2} \times 3^m$$

$$= \frac{[m(m-1)/2]!}{\{m(m-1)/2 - m\} m!} 3^m$$

Table 4.2  $\delta$ -values of parallel rules.

(a) Parallel $\beta$ -rule.	
$\delta$ -value of $\alpha$ -rule	$\delta$ -value of parallel $\beta$ -rule
acceptable	acceptable or unacceptable
impossible	acceptable or impossible or unacceptable
unacceptable	unacceptable

(b) Parallel $\alpha$ -rule.	
$\delta$ -value of $\beta$ -rule	$\delta$ -value of parallel $\alpha$ -rule
acceptable	acceptable or impossible
impossible	impossible
unacceptable	acceptable or impossible or unacceptable

$$= \left[ \prod_{i=0}^{n-1} \frac{n^2 - n - i}{2(n-i)} \right] 3^n \quad (4.3)$$

$= O(n^6)$  since  $n > m$

Table 4.3 shows the results of evaluating Equation (4.3) for several values of  $m$  and  $n$ . The number of possible c-rules increases rapidly as  $n$  increases and is very large even for relatively small values of  $m$  and  $n$ . These results emphasize the need for an informed generator when using generate-and-test to search the space of possible c-rules.

The searcher (Figure 4.5) uses a hierarchical approach to the problem of limiting the number of proposed relations: the cartographer identifies a 'promising' region in rule space, and the generate-test-scrounge process explores the individual c-rules in this region. The cartographer uses a map to describe the 'promising' region which it has identified. The derivator and assayer help the cartographer generate the maps.

The generate-test-scrounge process is a modification of the basic generate-and-test scheme. The generator proposes possible c-rules (in the current map) and the tester evaluates the proposed rule. There is also a scrounger which deduces the implications of the tester's decision for other c-rules; thereby permitting decisions to be made on more than one c-rule in a generate-test-scrounge cycle.

The searcher employs a simple control strategy. The generate-test-scrounge process proposes and evaluates c-rules until it has completely explored the current maps. It then activates the cartographer. The cartographer updates the maps by adding new c-rules to the maps. Upon completion of the update, the cartographer deactivates itself, and the generate-test-scrounge process proceeds to evaluate the new c-rules.

As indicated in Section 4.2, the searcher works under a time limit imposed by the scheduler. The searcher checks for violation of this limit at specific instances during its execution, that is, whenever any component of the searcher deactivates itself.

Sections 4.7.1 to 4.7.6 describe in detail the individual components of the searcher.

Table 4.3 Results of evaluating  $n_t$  for several values of  $m$  and  $n$ .

$m$	$n_t$				
	$n=4$	$n=5$	$n=6$	$n=7$	$n=8$
1	18	30	45	63	84
2	135	405	945	1890	3102
3	540	3240	12285	35910	88452
4	1215	17010	110565	481785	1658175
5	1438	61236	272729	1044807	2388200
6	729	153000	3648845	39588450	271643460

#### 4.7.1 Cartographer

The role of the cartographer is to find regions in the  $\alpha$ -graph ( $\beta$ -graph) which are likely to contain acceptable c-rules and to produce an  $\alpha$ -map ( $\beta$ -map) that describes these regions to the generate-test-scouting process. The cartographer operates in an incremental fashion: each time the cartographer is activated, it updates the  $\alpha$ -map ( $\beta$ -map) by adding new 'promising' c-rules to the  $\alpha$ -map ( $\beta$ -map).

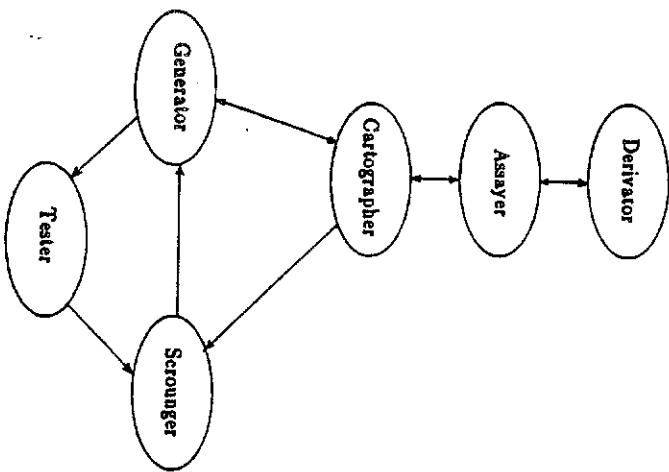
The cartographer's concept of a 'promising' c-rule is based upon its belief that an acceptable c-rule must be a semantically sensible construct, that is, its d-conditions are semantically sensible comparisons between semantically sensible symbols. In the current design of MUXT-1, the cartographer delegates much of the responsibility of ensuring semantic sensibility to the derivator and the assayer. When the cartographer becomes active, it uses the assayer to find new d-conditions. The assayer, in turn, uses the derivator to find new variables. The derivator contains heuristic rules that discard nonsensical symbols, and the assayer contains heuristic rules that generate sensible d-conditions from the symbols acceptable to the derivator.

If the assayer cannot find any new d-conditions, the cartographer causes MUXT-1 to stop; otherwise, it adds the operators associated with these new d-conditions to its pool of operators. Operators in the maps represent the addition of a d-condition to the c-rule represented by the node to which the operator is applied. The cartographer then adds the links and vertices resulting from these new operators to the maps.

Not all the c-rules that can be formed from the new operators will be included by the cartographer in the maps. Value-inheritance permits the cartographer to exclude all descendants of any c-rule known to have a  $\delta$ -value of either impossible or acceptable. Every descendant of an impossible c-rule is an impossible c-rule and can never lead to a dominance heuristic. Every descendant of an acceptable c-rule has a  $\delta$ -value of either acceptable or impossible but the descendant of a c-rule is a specialization of the c-rule; hence, descendants of an acceptable c-rule can never lead to any dominance heuristic better than that represented by the c-rule.

A dominance relation is a statement about the best solutions in the two nodes that satisfies its conditions. The value of the best solution in a node  $u$  is the sum of its current return,  $R(u)$ , and the best possible return from its remaining decisions. Thus, an acceptable rule has to contain statements about the relationship between the current returns of node  $x$  and node  $y$  as well as the relationship between the future returns of the two nodes. Instances of a

Figure 4.5 Organization of the searcher



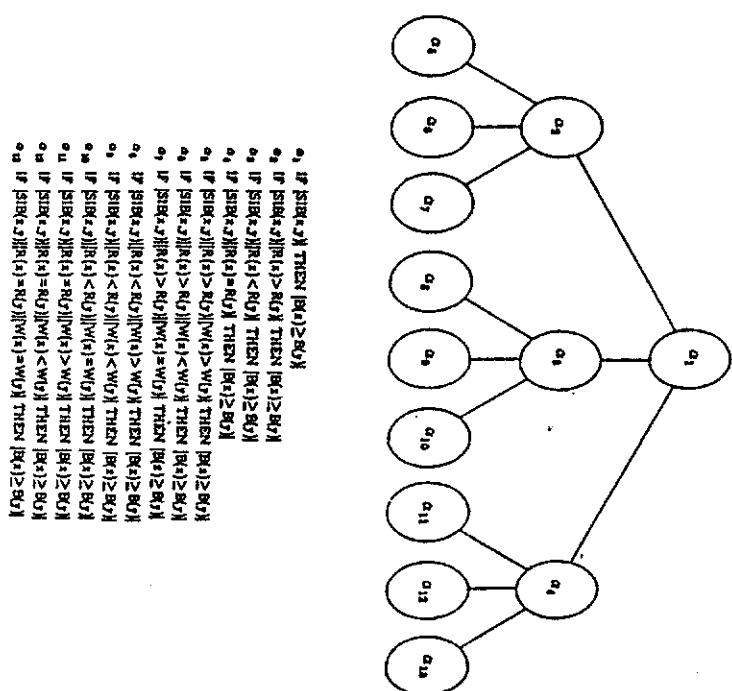
state variable in siblings show greater correlation than instances of a state variable in non-sibling equilevel nodes. Thus,  $\beta$ -rules must express the relationship between the current returns explicitly as a d-condition between  $R(x)$  and  $R(y)$ , while it may be possible for  $\alpha$ -rules to express the relationship implicitly by d-conditions between other state variables. However, we can always add the corresponding d-condition between  $R(x)$  and  $R(y)$  without altering the meaning of the set of d-conditions. These arguments suggest that it is sufficient to restrict the maps to those c-rules which contain a d-condition between  $R(x)$  and  $R(y)$ . A simple way of ensuring that all c-rules in the graph contain a d-condition between  $R(x)$  and  $R(y)$  is to only apply operators containing  $R(x)$  and  $R(y)$  to the root of the map.

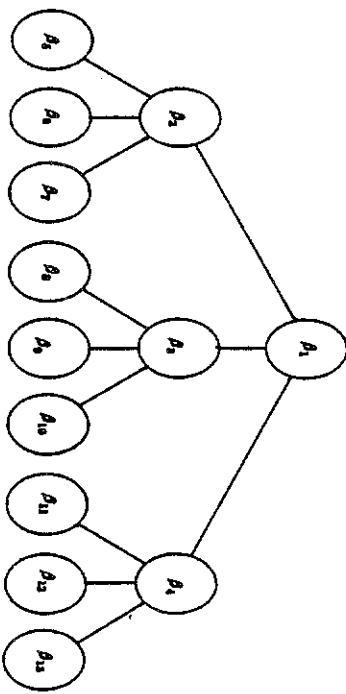
Even with the above measures, it is easy for the numbers of possible c-rules in the maps to become unmanageably large. To prevent this from happening, the MUTT-1's cartographer imposes a preset limit on the number of d-conditions which a c-rule in the maps can have.

When the cartographer is activated for the first time in a value-description of the 0/1-knapsack problem, the initial  $\alpha$ -map  $G_\alpha^0$  consists of the c-rule,  $\alpha_0$ : IF  $[SBR(x,y)]$  THEN  $[B(x) \geq B(y)]$ , and the initial  $\beta$ -map  $G_\beta^0$  consists of the c-rule,  $\beta_0$ : IF  $[L(x) = L(y)]$  THEN  $[B(x) \geq B(y)]$ . Suppose that the assayer returns the set of d-conditions: { $[R(x) < R(y)]$ ,  $[R(x) = R(y)]$ ,  $[R(x) > R(y)]$ ,  $[W(x) < W(y)]$ ,  $[W(x) = W(y)]$ ,  $[W(x) > W(y)]$ }. Only the operators involving d-conditions between  $R(x)$  and  $R(y)$  are applied to  $\alpha_0$  and  $\beta_0$ . Both  $\alpha_0$  and  $\beta_0$  have  $\delta$ -values of unacceptable, so no nodes are eliminated by value-inheritance of  $\delta$ -values. The resulting maps,  $G_\alpha^1$  and  $G_\beta^1$ , are shown in Figures 4.6 and 4.7.

When the generate-test-scrounge process completes its exploration of  $G_\alpha^1$  and  $G_\beta^1$ , the cartographer sets to work on updating the two maps. It will initiate the derivation of d-conditions. Suppose that this results in the assayer augmenting the set of d-conditions with  $[K\Delta W(x) < RW(x)]$ ,  $[K\Delta W(x) = RW(x)]$ ,  $[K\Delta W(x) > RW(x)]$ ,  $[K\Delta W(y) < RW(y)]$ ,  $[K\Delta W(y) = RW(y)]$ , and  $[K\Delta W(y) > RW(y)]$ , where  $K\Delta W(x)$  and  $RW()$  are new state variables derived by the derivator.  $K\Delta W()$  is the difference between the knapsack capacity  $K$  and the state variable  $W()$ .  $RW()$  is the sum of the weights of the objects for which decisions have yet to be made. Suppose that the d-limit is three. Figure 4.8 shows the new  $\alpha$ -map  $G_\alpha^2$ , and Figure 4.9 shows the new  $\beta$ -map  $G_\beta^2$ .  $G_\alpha^2$  differs from  $G_\alpha^1$  due to the different  $\delta$ -values of parallel rules in  $G_\alpha^1$  and  $G_\beta^1$ .

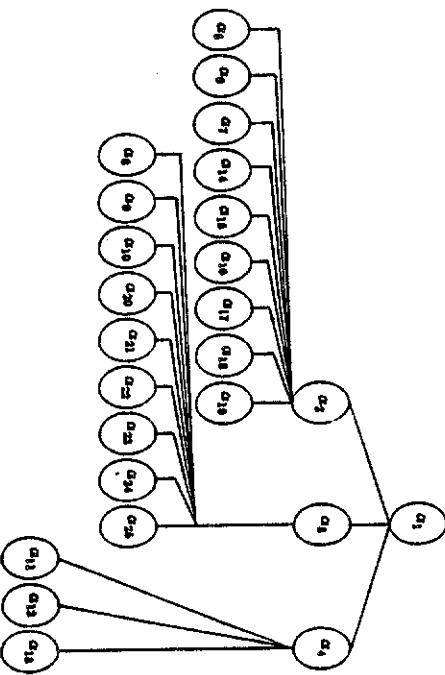
Figure 4.8  $G_\alpha^1$





$A_1$  IF  $|x| > 80$  THEN  $|x| \geq 80$   
 $A_2$  IF  $|x| < 100$  THEN  $|x| \geq 80$   
 $A_3$  IF  $|x| < 100$  AND  $x < 0$  THEN  $|x| \geq 80$   
 $A_4$  IF  $|x| = 100$  THEN  $|x| \geq 80$   
 $A_5$  IF  $|x| = 100$  AND  $x < 0$  THEN  $|x| \geq 80$   
 $A_6$  IF  $|x| < 100$  AND  $x > 0$  THEN  $|x| \geq 80$   
 $A_7$  IF  $|x| < 100$  AND  $x > 0$  AND  $x < 80$  THEN  $|x| \geq 80$   
 $A_8$  IF  $|x| < 100$  AND  $x > 0$  AND  $x > 80$  THEN  $|x| \geq 80$   
 $A_9$  IF  $|x| < 100$  AND  $x > 0$  AND  $x = 80$  THEN  $|x| \geq 80$   
 $A_{10}$  IF  $|x| < 100$  AND  $x < 0$  AND  $x < -80$  THEN  $|x| \geq 80$   
 $A_{11}$  IF  $|x| < 100$  AND  $x < 0$  AND  $x > -80$  THEN  $|x| \geq 80$   
 $A_{12}$  IF  $|x| < 100$  AND  $x > 0$  AND  $x > 80$  THEN  $|x| \geq 80$

Figure 4.7 G1



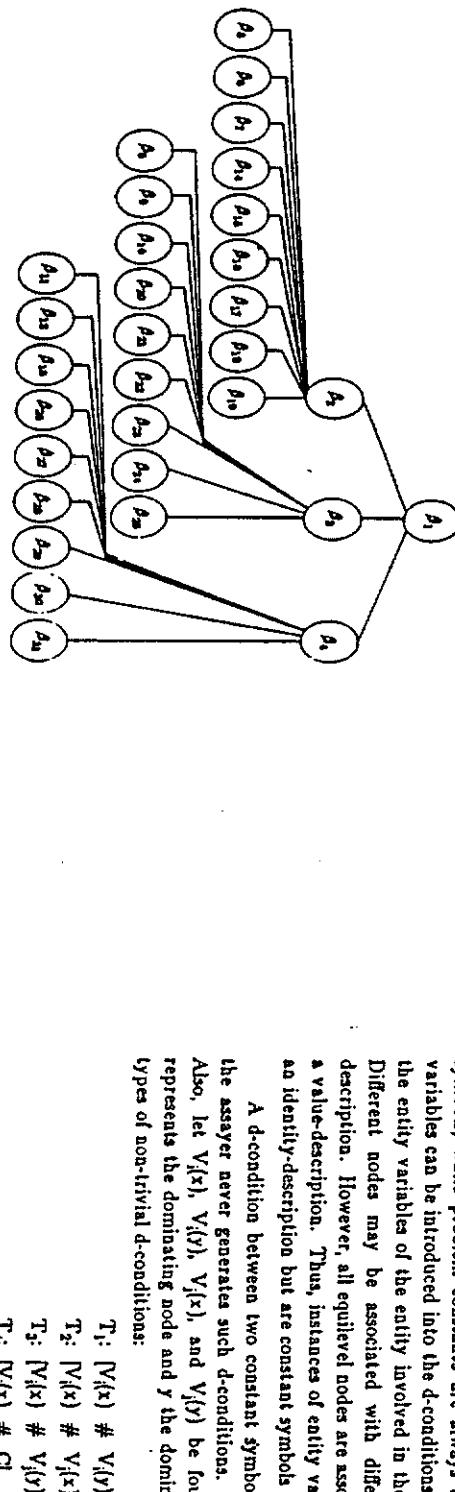
$D_1$  IF  $|x| < 100$  THEN  $|x| \geq 80$   
 $D_2$  IF  $|x| < 100$  AND  $x < 0$  THEN  $|x| \geq 80$   
 $D_3$  IF  $|x| < 100$  AND  $x > 0$  THEN  $|x| \geq 80$   
 $D_4$  IF  $|x| = 100$  THEN  $|x| \geq 80$   
 $D_5$  IF  $|x| = 100$  AND  $x < 0$  THEN  $|x| \geq 80$   
 $D_6$  IF  $|x| = 100$  AND  $x > 0$  THEN  $|x| \geq 80$   
 $D_7$  IF  $|x| < 100$  AND  $x < 0$  AND  $x < -80$  THEN  $|x| \geq 80$   
 $D_8$  IF  $|x| < 100$  AND  $x < 0$  AND  $x > -80$  THEN  $|x| \geq 80$   
 $D_9$  IF  $|x| < 100$  AND  $x < 0$  AND  $x = -80$  THEN  $|x| \geq 80$   
 $D_{10}$  IF  $|x| < 100$  AND  $x > 0$  AND  $x < 80$  THEN  $|x| \geq 80$   
 $D_{11}$  IF  $|x| < 100$  AND  $x > 0$  AND  $x > 80$  THEN  $|x| \geq 80$   
 $D_{12}$  IF  $|x| < 100$  AND  $x > 0$  AND  $x = 80$  THEN  $|x| \geq 80$   
 $D_{13}$  IF  $|x| < 100$  AND  $x < 0$  AND  $x < -100$  THEN  $|x| \geq 80$   
 $D_{14}$  IF  $|x| < 100$  AND  $x < 0$  AND  $x > -100$  THEN  $|x| \geq 80$   
 $D_{15}$  IF  $|x| < 100$  AND  $x < 0$  AND  $x = -100$  THEN  $|x| \geq 80$   
 $D_{16}$  IF  $|x| < 100$  AND  $x > 0$  AND  $x < 100$  THEN  $|x| \geq 80$   
 $D_{17}$  IF  $|x| < 100$  AND  $x > 0$  AND  $x > 100$  THEN  $|x| \geq 80$   
 $D_{18}$  IF  $|x| < 100$  AND  $x > 0$  AND  $x = 100$  THEN  $|x| \geq 80$

Figure 4.8 G2

#### 4.7.2 Assayer

A symbol is either a variable symbol or a constant symbol. Variable symbols may have different values for different nodes; constant symbols have the same value for all nodes. Instances of state variables are always variable symbols, while problem constants are always constant symbols. The entity variables can be introduced into the d-conditions by associating each node with the entity variables of the entity involved in the decision leading to the node. Different nodes may be associated with different entities in an identity-description. However, all equivalent nodes are associated with the same entity in a value-description. Thus, instances of entity variables are variable symbols in an identity-description but are constant symbols in a value-description.

A d-condition between two constant symbols is a trivial d-condition, and the assayer never generates such d-conditions. Let  $C$  be a constant symbol. Also, let  $V_i(x)$ ,  $V_j(y)$ ,  $V_i(x)$ , and  $V_j(y)$  be four variable symbols, where  $x$  represents the dominating node and  $y$  the dominated node. There are four types of non-trivial d-conditions:



A d-condition of type  $T_1$  is a condition between two instances of the same state variable and is always semantically sensible if the variable  $V_i(x)$  is semantically sensible; thus, the assayer always generates d-conditions of type  $T_1$  from the set of variables supplied by the derivator.

In the 0/1-knapsack problem, both  $R()$  and  $W()$  are semantically sensible variables because they have semantic meaning.  $R()$  is the return from the items included in the knapsack, and  $W()$  is the weight of the item in the knapsack. Among the possibilities considered by the derivator when defining new state variables are combinations between  $W()$  and  $K$ . One possible new state variable is  $K\Delta W()$ , the difference between  $K$  and  $W()$ . This state variable is semantically sensible; it is the capacity remaining in the knapsack. In contrast, the variable  $K\Sigma W()$ , the sum of  $K$  and  $W()$ , has no semantic meaning. In contrast, d-conditions of type  $T_2$  and  $T_3$  compare instances of different variables while d-conditions of type  $T_4$  compare an instance of a variable with a constant symbol. The assayer generates a condition belonging to one of these three types only when it is able to deduce from the working description that

Figure 4.9  $C^{\beta}$

the symbols in the d-conditions are semantically related. By definition, two symbols are semantically related if they are connected either explicitly or implicitly by a construct in the working description.

A state variable is related to a constant if the constant is an associated problem constant or a symbolic constant within the range of the state variable. In the 0/1-knapsack problem, state variable  $W()$  is semantically related to the problem constant  $K$  as  $K$  is associated with  $R()$  by the constraint  $W(N) \leq K$ .

Two state variables are related if they share a common construct in the working description, if their definitions share a common symbol, or if they are related to a common variable. One new state variable generated by the derivator for the 0/1-knapsack problem is  $RW()$ , the sum of the weights of the unassigned items. Its definition is the formula,  $RW(N_p) = RW(N_p) - w(\theta)$ , where  $N_p$  is the node in the search tree generated by the application of operator  $\theta$  to node  $N_p$ .  $RW()$  is semantically related by  $W()$  as the entity variable  $w()$  appears in both their definitions. Recall that, the definition of  $W()$  is  $W(N_p) = W(N_p) + w(\theta)$ . Another new state variable derived by the derivator for the 0/1-knapsack problem is  $K\Delta W()$ , the capacity remaining in the knapsack. This state variable is defined by the formula,  $K\Delta W(N) = K - W(N)$ .  $RW()$  is semantically related to  $K\Delta W()$  since both of them are related to  $W()$ .

Whenever the assayer in MUFTT-1 becomes active, its first action is to call upon the derivator to find new symbols. There is a global symbol list which consists of two sublists; the a-list and the b-list. The a-list holds the variables and constants generated by the derivator during its most recent activation, while the b-list holds the variables and constants generated by the derivator prior to its most recent activation. The cartographer then looks for sensible d-conditions using the rules outlined earlier in this section. It looks for d-conditions between two symbols from the a-list and d-conditions between one symbol from the a-list and one symbol from the b-list. Useful d-conditions between two symbols from the b-list have already been considered by previous activations of the assayer. To facilitate future design changes in the assayer, the list of d-conditions derived by the assayer is also partitioned into two sublists: a u-list and a v-list. The u-list holds the d-conditions derived during the most recent activation of the assayer, while the v-list holds those derived during previous activations.

When the assayer becomes active for the first time in a value-description of the 0/1-knapsack problem, both the a-list and b-list are empty. The derivator will put the symbols  $r()$ ,  $w()$ ,  $K()$ ,  $W()$ , and  $K$  in the a-list. The

assayer considers  $r()$ ,  $w()$ , and  $K$  to be constant symbols, so it does not generate any d-conditions between them. Moreover,  $r()$  and  $w()$  do not appear in any constraint with either  $R()$  or  $W()$ , so no d-conditions between them will be generated.  $K$ , the knapsack capacity, and  $W()$  appear in a constraint, but the assayer ignores this constraint as it is an implicit condition satisfied by all feasible nodes in a problem tree. State variables  $R()$  and  $W()$  do not appear together in any constraint. The assayer will only generate the d-conditions:  $[R(x) < R(y)]$ ,  $[R(x) = R(y)]$ ,  $[R(x) > R(y)]$ ,  $[W(x) < W(y)]$ ,  $[W(x) = W(y)]$ , and  $[W(x) > W(y)]$ .

During the second activation of the assayer, the set of symbols known to the derivator has been augmented by the entity variable,  $r\#w()$ , and the state variables,  $RR()$ ,  $RW()$ , and  $K\Delta W()$ .  $r\#w()$  is the ratio of  $r()$  to  $w()$ .  $RR()$  is the sum of the return of the unassigned items. The assayer will not generate any d-condition involving  $r\#w()$  as it is not related to any variable symbol. As indicated previously,  $W()$ ,  $R()$ ,  $RW()$ , and  $K\Delta W()$  are all semantically related.  $RW()$  and  $K\Delta W()$  are forward variables while  $W()$  is a backward variable, so the assayer uses  $RW()$  and  $K\Delta W()$  instead of  $RW()$  and  $W()$  and augments the set of 'sensible' d-conditions with  $[RW(x) < K\Delta W(x)]$ ,  $[RW(x) = K\Delta W(x)]$ , and  $[RW(x) > K\Delta W(x)]$ . Similarly,  $RR()$  is related to  $R()$  as both state variables have  $r()$  in their definitions. However, the assayer does not generate any d-conditions between the two variables as  $R()$  is a backward variable while  $RR()$  is a forward variable. Furthermore, all nodes on the same level have identical values for  $RR()$  and  $RW()$  and so the assayer does not generate d-conditions between  $RR()$  and  $RW()$ .

#### 4.7.3 Derivator

The derivator derives new state variables and entity variables in an incremental fashion. Each time the derivator is activated, it generates new variables by combining existing variables and constants. It includes these new variables in its set of known variables and constants. These new variables will participate in the derivation of other variables during the derivator's next activation.

New entity variables are derived by the formula

$$E_0 = E_1 \# E_2 \quad (4.4)$$

where  $E_0$  is the new entity variable,  $E_1$  and  $E_2$  are two existing entity variables, and  $\#$  is one of the arithmetic operators:  $+$  (plus),  $-$  (minus),  $\times$

(times), and  $\div$  (divide). Only arithmetic operators are needed as the domain of MTT-1 has been restricted to number optimization problems. Indiscriminate application of Eq. (4.4) often results in new entity variables which are semantic nonsense. The derivator uses heuristic rules to restrict its suggestions to combinations that are likely to have semantic meaning. The derivator combines two entity variables only when it finds evidence in the problem description that they are related.

In Section 3.4, we have subdivided the entity variables in the problem description into return properties, resource properties, and limit properties.

Entity variables are related if they share a common construct in the problem description. Two common shared constructs are the objective function and the constraint. It is also possible that two or more constraints have a common variable. For these problems, the derivator includes the constraint obtained by substituting for the common variable in the set of constraints. When combining two variables with a common construct, the derivator uses either the arithmetic operator linking the variables in the common construct or its inverse. This principle guides the derivation of new variables between two return properties, two resource properties, two limit properties, a resource variable and a limit property, and a return property and a limit property.

The entity variables in the 0/1-knapsack problem do not share any common construct. The entity variables in sequencing to minimize weighted completion time do share common construct. Recall that sequencing to minimizes weighted completion time involves a set  $T$  of  $n$  tasks. Each task  $t \in T$  has a length  $\ell(t)$ , a weight  $w(t)$ , and a deadline  $d(t)$ . The desired solution is the one-processor schedule  $\sigma$  for  $T$  which minimizes  $\sum_{t \in T} [\sigma(t) + \ell(t)]w(t)$  and satisfies  $[\sigma(t) + \ell(t)] \leq d(t)$  for each task  $t \in T$ . The entity variables in this problem are  $w(t)$ ,  $\ell(t)$ , and  $d(t)$ .  $\ell(t)$  and  $d(t)$  are related by the constraint,  $[\sigma(t) + \ell(t)] \leq d(t)$ . The derivator groups together these two variables by moving  $\ell(t)$  to the right-hand side of the inequality, resulting in the expression,  $(\sigma(t) \leq [d(t) - \ell(t)])$ .  $\ell(t)$  and  $d(t)$  are now related by the minus operator, and the derivator generates the new entity variable,  $d\Delta\ell(t)$ , which is the difference between  $d(t)$  and  $\ell(t)$ . Similarly, the optimization function,  $\sum_{t \in T} [\sigma(t) + \ell(t)]w(t)$  causes the derivator to derive the entity variable,  $\ell \Delta w(t)$ , which is the product of  $\ell(t)$  and  $w(t)$ .

Another good candidate for a new semantically meaningful entity variable is the ratio of a return property and a resource property. This new variable measure for the per unit return of this entity. In the 0/1-knapsack problem,  $r(t)$

is a return property, while  $w(t)$  is a resource property. The derivator will derive the new entity variable,  $r\Delta w(t)$ , the ratio of  $r(t)$  to  $w(t)$ . It is the return per unit weight of the entity.

When the derivator combines two entity variables of the same type, it regards the new variable as a variable of the same type as the variables used in its definition. The derivator generates a hybrid property when the two entity variables in Eq. (4.4) are of different types. The new entity variable  $r\Omega w(t)$  is a hybrid property as it results from an arithmetic combination of a return property with a resource property. Sensible combination of hybrid variables often requires background knowledge which cannot be gleaned from the problem description. The current design of the derivator does not combine hybrid properties with other entity variables.

The derivator always tries to generate sensible state variables. It is much easier to define a nonsensical state variable than a nonsensical entity variable as a state variable is more complex than an entity variable. A state is a data structure which gives a snapshot of the condition of the problem solving process at each stage of its solution and the state variables are elements of this data structure. Thus, the state variable is a summary of an aspect of the decision sequence leading to the node. The derivator's heuristic rules are based upon our experiences in working with search algorithms and can produce some 'promising' state variables.

The first type of state variables known to the derivator are the *summary variables*. Summary variables are summarize the values of entity variables of the set of unassigned entities.

Each item in the 0/1-knapsack problem has two entity variables,  $r(t)$  and  $w(t)$ . The derivator derives the summary variables,  $RR(t)$  and  $RW(t)$ , to correspond to these entity variables.  $RR(t)$  is the sum of  $r(t)$  for the unassigned items, and  $RW(t)$  is the sum of  $w(t)$  for the unassigned items.

The second type of state variables are the *alternate variables*. An alternate variable is an alternative view of another state variable. Although the two variables are equivalent, one variable is often more suitable than the other variable for expressing 'some fact'. The derivator generates an alternate variable when it finds a state variable with a limit; The difference between the limit and the state variable is an alternative view of the state variable.

In the 0/1-knapsack problem, state variable  $W(t)$  is associated with the problem constant  $K$  by the constraint  $W(t) \leq K$ . The derivator will derive the new variable  $K\Delta W(t)$ , where  $K\Delta W(t) = K - W(t)$ .  $K\Delta W(t)$  represents the unused capacity in the knapsack and is an alternative view of  $W(t)$ . However,

$K\Delta W()$  is more convenient than  $W()$  for expressing the condition that all the unassigned items can fit into the knapsack. This condition is  $|K\Delta W(x)| \geq RW(x)$  when  $K\Delta W()$  is used, and  $|SUM(W(x), RW(x))| \leq K$  when  $W()$  is used. The latter form is more awkward than the former, and is much more difficult to learn.

As indicated in the previous section, the symbol list in MUTT-1 consists of an a-list and a b-list. The b-list, in turn, consists of two sublists: the ba-list and the bb-list. When the assayer first activates the derivator, the a-list, the ba-list, and the bac-list are all empty so the derivator copies the set of symbols in the working description into the a-list.

In subsequent activations, the derivator first appends the contents of the ba-list to the bb-list and moves the contents of the a-list into the ba-list. It then derives sensible variables from the symbols in the b-list according to the rules discussed earlier in this section. There are no restrictions on the type of the variables from the ba-list. The derivator derives a variable that involves a symbol in the bb-list only when it is accompanied by a symbol from the ba-list. Variables involving just symbols in the bb-list have been previously generated.

When the cartographer first activates the derivator for the value-description of the 0-1-knapsack problem, the a-list, ba-list, and bb-list are all empty. The derivator copies the set of symbols in the working description into its a-list. This consists of the two entity variables,  $w()$  and  $r()$ , the two state variables,  $W()$  and  $R()$ , and the problem constant B. The a-list is now  $\{R(), W(), r(), K\}$ .

When the cartographer activates the derivator for the second time, the derivator copies the ba-list into the bb-list and moves the contents of the a-list into the ba-list. Thus, the ba-list is now  $\{R(), W(), r(), w()\}$ , while the bb-list is still  $\{\}$ . There are no restrictions on the form of the variables derived by the derivator from the ba-list. The derivator will derive the entity variable,  $r\Omega w()$ , the summary variables,  $RW()$  and  $RR()$ , and the alternative variable,  $K\Delta W()$ .

#### 4.7.4 Tester

The tester forms a set of test patterns from the proposed rule: the positive pattern [ $t$ -condition] [ $d$ -condition]<sub>1</sub> ... [ $d$ -condition]<sub>m</sub>] [ $B(x) > B(y)$ ], the null pattern [ $t$ -condition] [ $d$ -condition]<sub>1</sub> ... [ $d$ -condition]<sub>m</sub>] [ $B(x) = B(y)$ ], and the negative pattern [ $t$ -condition] [ $d$ -condition]<sub>1</sub> ... [ $d$ -condition]<sub>m</sub>] [ $B(x) < B(y)$ ]. Let the term *positive example* denote a node pair that matches the positive

pattern. Similarly, a *negative example* matches the negative pattern and a *null example* matches the null pattern.

Let the *antirule* be the rule formed by interchanging the symbols  $x$  and  $y$  in the conditions of a rule. The negative (positive) pattern of the antirule is identical to the positive (negative) pattern of the proposed rule. The null patterns are identical. The sets of test patterns produced from the rule and antirule are isomorphic. A positive example for any rule is a negative example for its antirule and vice-versa.

The rule, IF  $[SIB(x,y)] [R(x) > R(y)]$  THEN  $[B(x) \geq B(y)]$  (rule  $\alpha_2$ ), is the first rule in map  $G_o$  (Figure 4.4) selected for testing by the generator. The tester will create the positive pattern  $[SIB(x,y)] [R(x) > R(y)] [B(x) > B(y)]$ , the null pattern  $[SIB(x,y)] [R(x) > R(y)] [B(x) = B(y)]$ , and the negative pattern  $[SIB(x,y)] [R(x) > R(y)] [B(x) < B(y)]$  from  $\alpha_2$ .

Interchanging the symbols  $x$  and  $y$  in the left-hand side of  $\alpha_2$  produces IF  $[SIB(x,y)] [R(x) < R(y)]$  THEN  $[B(x) \geq B(y)]$ , which is the c-rule,  $\alpha_c$ . The positive pattern for  $\alpha_3$  is  $[SIB(x,y)] [R(x) < R(y)] [B(x) > B(y)]$ , which is identical to the negative pattern of  $\alpha_2$  once the symbols have been renamed. Similarly, the negative and null patterns for  $\alpha_3$ ,  $[SIB(x,y)] [R(x) < R(y)] [B(x) < B(y)]$  and  $[SIB(x,y)] [R(x) < R(y)] [B(x) = B(y)]$ , are respectively equivalent to the positive and null patterns of  $\alpha_2$ .

The tester gets one or more test trees from the forester and matches the set of patterns against node pairs in the test trees. Table 4.4 lists the possible outcomes of this matching process and their implications. The presence of both positive and negative examples implies that the set of d-conditions is unable to consistently identify the node with the better solution; thus, the rule and its antirule are both unacceptable rules.

The presence of positive (negative) examples and the absence of negative (positive) examples suggests possible acceptability of the rule (antirule) and unacceptability of the antirule (rule). However, any rule can only be accepted only after it has been subjected to adequate testing. MUTT-1 uses a very simple criteria: the testing is adequate if the number of positive and null examples exceeds  $s_1$  and the number of test trees exceeds  $s_2$ . Suitable values for  $s_1$  and  $s_2$  are 20 and 10 respectively. The tester resumes the matching process if it finds the test set to be inadequate. This simple criteria is sufficient for a prototype program where any errors due to inadequate testing will be apparent to the human experimenter. Statistically sound criteria will have to be developed in the future.

Table 4.4 Possible outcomes of the matching process.

Positive Examples	Null Examples	Negative Examples	Implication	
			Rule	Antirule
yes	no	no	acceptable	unacceptable
yes	yes	no	acceptable	unacceptable
yes	no	yes	unacceptable	unacceptable
yes	yes	yes	unacceptable	unacceptable
no	yes	yes	unacceptable	acceptable
no	no	yes	unacceptable	acceptable
no	no	no	impossible	impossible
no	yes	no	acceptable	acceptable

If there are no examples then either the two rules are impossible or the set of test trees does not contain the proper examples. The tester uses the size of the test set to decide between these two possibilities. If the number of test trees is less than  $\delta_3$ , then it resumes the matching process with new test trees. A suitable value for  $\delta_3$  is 10.

Suppose that the tester will match its set of patterns against the nodes in search tree,  $r_1$  (Table 4.1).  $\{N_{1a}, N_{1b}\}$ ,  $\{N_{1a}, N_{1b}, N_{1c}, N_{1d}\}$ ,  $\{N_{1a}, N_{1b}, N_{2a}, N_{2b}\}$ ,  $\{N_{2a}, N_{2b}\}$ , and  $\{N_{2a}, N_{2b}\}$  match the positive pattern;  $\{N_{1a}, N_{1b}\}$  matches the null pattern; and  $\{N_{2a}, N_{1b}\}$ ,  $\{N_{1a}, N_{2b}\}$ ,  $\{N_{1a}, N_{1b}\}$ , and  $\{N_{1a}, N_{1b}\}$  match the negative pattern. The presence of both positive and negative examples causes the tester to reject rule  $\alpha_2$  and its antirule  $\alpha_3$ .

#### 4.7.5 Scronger

The scronger investigates the effects of a decision on other rules. The decision made by the tester can imply decisions for other rules due to the inheritance of decision values and the parallelism between the  $\alpha$ -graph and the  $\beta$ -graph.

Suppose that rule  $c_i$  is an unacceptable rule. There is at least one unacceptable rule among  $c_i$ 's children, but there may also be acceptable rules or impossible rules. The scronger attempts to identify the decision values of the children rules by matching the examples of  $c_i$  against d-conditions that connect  $c_i$  to its children in the map. Let  $d_j$  be the d-condition that connects  $c_i$  to  $c_j$  in the map. An example of  $c_i$  that satisfies a d-condition is an example of the new rule formed by adding the d-condition to  $c_i$ . The subset of example set of  $c_i$  which match the  $d_j$  forms a example set for  $c_j$ . If this subset contains both positive and negative examples, then both  $c_j$  and its antirule may be rejected.

A positive example of rule  $\alpha_2$  is a node pair that satisfies the conditions  $[SIB(x,y)]$ ,  $[R(x)>R(y)]$ , and  $[B(x)>B(y)]$ . A positive example of  $\alpha_2$  that satisfies the condition  $[W(x)>W(y)]$  is a node pair that satisfies the conditions  $[SIB(x,y)]$ ,  $[R(x)>R(y)]$ ,  $[W(x)>W(y)]$ , and  $[B(x)>B(y)]$ , that is, a positive example of the resulting child rule  $\alpha_5$ . Similarly, a null (negative) example of rule  $\alpha_2$  that satisfies  $[W(x)>W(y)]$  is also a null (negative) example of rule  $\alpha_5$ .

An empty subset implies that there are no examples of both the  $\alpha_i$  and its antirule among the test trees that generated the example set of  $\alpha_i$ . The stronger does not know whether or not the set of test trees used on  $\alpha_i$  as an adequate test set. Thus, it labels  $\alpha_i$  and its antirule as potentially impossible

rules and waits for the tester to make the final decision on the rules.

Additional testing by the tester may reveal that the potentially impossible rule is actually an acceptable rule, an unacceptable rule, or an impossible rule.

Similarly, if the subset does not contain negative (positive) examples of  $\alpha_i$ , then  $\alpha_i$  is a potential acceptable (impossible) rule and its antirule is a potential impossible (acceptable) rule.

The generate-test-scrounge process gives a higher priority to rules which have been identified as potential impossible or acceptable rules.

For example, the rejection of  $\alpha_2$  causes the scrounger to reject the parallel rules  $B_2$  and  $\beta_2$ . The scrounger then tries to tentatively identify the decision values of  $\alpha_2$ 's children. In tree  $r_1$  (Table 4.1), node pairs,  $(N_{\alpha}, N_j)$ ,  $(N_{\alpha}, N_{i^*})$ ,  $(N_{\beta}, N_{j^*})$ ,  $(N_{\alpha}, N_{2^*})$ ,  $(N_{\alpha}, N_{2^*})$ ,  $(N_{\alpha}, N_{2^*})$  and  $(N_{\alpha}, N_{2^*})$  are positive examples of  $\alpha_2$ .  $(N_{\alpha}, N_{2^*})$  is a null example of  $\alpha_2$ , and node pairs,  $(N_2, N_i)$ ,  $(N_4, N_3)$ ,  $(N_4, N_5)$ ,  $(N_2, N_1)$ , and  $(N_{10}, N_{10})$  are negative examples. The scrounger matches these examples against the condition  $[W(x) > W(y)]$  and finds that all of the examples successfully matches the d-condition. There are positive, null, and negative examples of  $\alpha_5$  in tree  $r_1$ . The scrounger rejects rule  $\alpha_5$  which leads to the rejection of  $\alpha_9$ ,  $\beta_9$ , and  $\beta_9$ .

Next, the scrounger considers the condition  $[W(x) < W(y)]$ . Only  $(N_{\alpha}, N_{12})$  satisfies this condition. This result implies that there are only positive examples of the resulting rule,  $\alpha_6$ , in test tree,  $r_1$ . Rule  $\alpha_6$  is a potentially acceptable rule. Additional testing by the tester will reveal that  $\alpha_6$  is an acceptable rule.

Finally, the scrounger considers the condition  $[W(x) = W(y)]$ . No example satisfies this condition, so rule  $\alpha_7$  is a potentially impossible rule. Further testing by the tester will determine that  $\alpha_7$  is actually an acceptable  $\alpha$ -rule.

The acceptance of a rule may change the order in which rules are proposed by the generator. An acceptable rule indicates the presence of a region of consistent projections in the symbol space. Priority should then be given to determining the actual boundaries of this consistent region because all rules within its boundaries are acceptable rules. This implies priority for the testing of siblings of the acceptable rule and their descendants.

#### 4.7.6 Generator

MUTT-1 keeps a set of supporting reasons for each candidate rule. The generator always proposes the rule with the best supporting reasons. The acceptable reasons are:

- (1) The rule is the rule selected by the breadth-first criterion.
- (2) The rule is either a sibling of an accepted rule or a descendant of a sibling of an accepted rule.
- (3) The scrounger has tentatively identified this rule as either an impossible rule or an acceptable rule.
- (4) The parallel  $\beta$ -rule is an acceptable rule, or the parallel  $\alpha$ -rule is either an acceptable rule or an impossible rule.

These reasons have been listed in order of increasing priority. The generator tries to select the rule whose most compelling reason has the highest priority level. In case of ties, the generator will continue looking at the priority level of the next most compelling reason until all ties are resolved. The breadth-first criterion ensures that there will always be one 'best' rule. The motivation behind reasons (2), (3), (4), and (6) have been discussed elsewhere in this chapter.

#### 4.8 Verification

Hand-simulations have demonstrated the ability of MUTT-1 to learn dominance relations. A prototype of the searcher has also been implemented. Appendix A contains a listing of the searcher. The implemented searcher uses a static set of user-supplied test search trees as the forester was not implemented. Also, the variable descriptions must be supplied to the implemented searcher in the appropriate form as the description manager was not implemented. Nonetheless, this implementation of the searcher is useful as it demonstrates the validity of three vital aspects of the design, that is, the rule graph description of the rule space, the process of deriving new symbols, and the process of deriving new d-conditions.

When MUTT-1 is applied to the 0/1-knapsack problem, it learns the relation,  $\text{IF } [L(x)=L(y)] \text{ } [R(x) \geq R(y)] \text{ } [W(x) \leq W(y)] \text{ THEN } [B(x) \geq B(y)]$ , which is the well-known dominance relation for the 0/1-knapsack problem. It is also able to learn the relation,  $\text{IF } [l(x)=l(y)] \text{ } [R(x) \geq R(y)] \text{ } [\Delta W(x) \geq RW(l)] \text{ THEN } [B(x) \geq B(y)]$ , which leads to the problem solver to include all of the unassigned items when the capacity of the knapsack exceeds the total weight of

the unassigned items. Another relation learned by MU<sup>T</sup>T-1 for the 0/1-knapsack problem is IF  $\{B(x,y)\} [R(x) \geq R(y)]$  THEN  $[B(x) \geq B(y)]$  where  $RERR()$  is defined as the sum of  $R()$  and  $RR()$ .  $RERR()$  is an upper bound on the return from any node. Appendix B shows the unmodified trace produced by our implementation of the searcher for the value-description of the 0/1-knapsack problem when supplied with the accompanying input data.

MU<sup>T</sup>T-1 has also been able to learn the well-known dominance relation found in sequencing to minimize weighted completion time, sequencing to minimize weighted lateness, a simplified inventory problem, a simplified two-machine flowshop problem, sequencing to minimize latest completion time, and sequencing to minimize task weight.

#### 4.9 Summary

MU<sup>T</sup>T-1 is a prototype program that is able to learn dominance relations by experimentation. It uses the model-driven generate-and-test paradigm, and it models the space of possible candidates by a rule graph. MU<sup>T</sup>T-1 uses both the properties of this rule graph and semantics of the problem to limit the number of possible candidates proposed and tested by the generate-and-test-process. An integral part of this process is the derivation of new variables. MU<sup>T</sup>T-1 has been applied to several problems with known dominance relations and has been able to learn those dominance relations.

## CHAPTER 5 CONCLUSIONS AND FUTURE WORK

Ways of improving the efficiency of the search process includes reducing the number of nodes searched by the process and reducing the average time to search each node. The number of nodes searched by the search process may be reduced by using an appropriate problem representation, an appropriate selection scheme, and pruning rules. The average time to search each node may be reduced by reducing the mismatch between the algorithm and the architecture. In this thesis, we have studied two of these problems: the effect of the architecture on the selection scheme and the derivation of pruning rules.

Even though the best-first B&B algorithm generally expands the least number of nodes among the different B&B algorithms, it is not always the best B&B algorithm. Most computer systems with large memories use a multi-level hierarchy of at least two levels. On a two-level memory system, the cost of accessing an item of information depends on the level in which the item is located. The secondary memory is significantly slower than the main memory. The best-first B&B algorithm has exponential space and time complexities, while the depth-first B&B algorithm has a linear space complexity and an exponential time complexity. For many problem instances, the use of the best-first B&B algorithm will require the use of both main memory and secondary memory, while the depth-first B&B algorithm often only requires the use of main memory. As a result, the average time required to expand a node in the best-first B&B algorithm is greater than that required by the depth-first B&B algorithm. The best B&B algorithm will depend on the characteristics of the problem and the characteristics of the memory hierarchy. When the disparity between the numbers of nodes expanded by the best-first and depth-first B&B algorithms is small, the depth-first B&B algorithm is generally preferable. On the other hand, when the best-first B&B algorithm expands much less nodes than the depth-first B&B algorithm, the best-first B&B algorithm should be used. For intermediate situations, the algorithm chosen will depend on the relative speeds of the different levels of the hierarchy: it is

the depth-first B&B algorithm when the secondary memory is very slow, and the modified B&B algorithm otherwise.

The modified B&B algorithm partitions the range of possible lower bounds into several distinct subranges. Nodes failing within each subrange are stored in a LIFO stack associated with the subrange. Each stack may partly reside in the secondary memory. The modified B&B strategy does not wait for nodes to be transferred from the secondary memory to the main memory but continues to expand nodes which are currently available in the main memory. The modified B&B algorithm demonstrates the effectiveness of including architectural considerations in the design of an algorithm.

Dominance relations are pruning rules between two nodes in the search tree. Even though dominance relations are problem dependent, it is possible to find them by machine learning techniques. In this thesis, we have studied

machine learning of dominance relations for a number of optimization problems. A natural representation for dominance relations is the production rule. The production rule can be easily converted into either a logical implication or a frame. The representation of the optimization problems embeds semantic information into the problem description. This frame representation permits a systematic manipulation of semantic information about the problem domain and facilitates the transfer of information between domains.

Dominance relations can be learned by experimentation, deduction, and analogy. Learning by deduction only learns relations which can be deduced from facts known to the deduction program. Not all of the relevant facts are known about every problem, especially unfamiliar problems. Learning by analogy transforms a known dominance relation into a relation which may be suitable for the new problem. It is most useful when the new problem domain is similar to a domain with known dominance relations. Learning by experimentation is the most general mechanism as it can learn relations in unfamiliar problem domains and learn new types of relations. Thus, learning by experimentation is the primary learning mechanism studied in this thesis. Both analogy and deduction are also necessary as they are very effective when used in the correct context.

We have studied learning dominance relations by experimentation in greater detail and implemented a prototype program. This program uses model-driven generate-and-test due to the impossibility of identifying training instances needed to drive a data-driven method paradigm based upon heuristic constraints derived from manipulations of the standard descriptions as well as

constraints in the space of relations to reduce the number of proposed candidate relations. In addition, the program is able to automatically define and manipulate new variable symbols.

There is much left to do in machine learning of dominance relations before the goal of an integrated system incorporating multiple learning mechanisms is achieved. We have only taken a small step in the process of refining the learning mechanisms. MUTT-1 is a prototype and has to be extended to other types of number optimization problems. We do not anticipate major change in the design of MUTT-1, only modifications to the heuristic knowledge driving the assayer and derivator to reflect new problem domains. Work has also to be done on deriving the knowledge needed to drive the analogy and deduction mechanisms.

This work can be extended into a study of dominance relations for decision problems. There are no simple return variables in decision problems and it is necessary to deal with different types of heuristic measure. There is also a question of the utility of dominance relations for these problems.

In addition, there is the intriguing possibility of using learning by discovery to automate the process of finding the knowledge which drives learning by experimentation, learning by analogy, and learning by deduction.

Another possible avenue of investigation will be to study pruning heuristics comprised of candidate relations which usually give good results. Some questions in this direction are the statistical measures of a good heuristic and the effects of this relaxation on the learning mechanisms.

## LIST OF REFERENCES

- [AEH84] T. O. Alanko, H. H. A. Erik, and I. I. Halikala, "Virtual Memory Behavior of some Sorting Algorithms," *IEEE Trans. on Software Engineering*, vol. SE-10, no. 4, pp. 422-431, July 1984.
- [AHU74] A. V. Aho, J. E. Hopcroft, and J. D. Ullman, *The Design and Analysis of Computer Algorithms*, Addison-Wesley, 1974.
- [Am86] S. Amarel, "Program Synthesis as a Theory Formation Task: Problem Representations and Solution Methods," *Machine Learning: An Artificial Intelligence Approach - Volume II*, R. S. Michalski, J. G. Carbonell, and T. M. Mitchell (Eds.), Morgan Kaufmann, Los Altos, 1986.
- [And84] P. M. Andras, "Constraint-Limited Generalization: Acquiring Procedures from Examples," *AAAI-84*, pp. 6-10, 1984.
- [AnS79] Y. Anzai and H. Simon, "The Theory of Learning by Doing," *Psychological Review*, Vol. 36, No. 2, pp. 124-140, 1979.
- [Ara84] A. A. Araya, "Learning Problem Classes by Means of Experimentation and Generalization," *AAAI-84*, pp. 11-15, 1984.
- [BGM70] B. S. Brown, F. G. Gustavson, and E. S. Mankin, "Sorting in a Paging Environment," *CACM*, vol. 13, no. 8, pp. 438-444, Aug. 1970.
- [Bas76] J. L. Baer and G. R. Sager, "Dynamic Improvement of Locality in Virtual Memory Systems," *IEEE Trans. on Software Engineering*, vol. SE-2, no. 1, pp. 54-62, March 1976.
- [Bel62] R. Bellman and S. Dreyfus, *Applied Dynamic Programming*, Princeton University Press, NJ, 1962.
- [Bel71] M. Belmore and J. C. Malone, "Pathology of Travelling Salesman Subtour-Elimination Algorithms," *Operations Research*, vol. 19, no. 10, pp. 278-307, 1971.

- [BeNe88] M. Bellmore and G. L. Nemhauser, "The Traveling Salesman Problem: A Review," *Operations Research*, vol. 16, pp. 538-558, 1968.
- [Ber70] H. Berliner, "The B' Tree Search Algorithm: A Best-First Proof Procedure," *Artificial Intelligence*, vol. 12, pp. 23-40, 1970.
- [Bob80] D. G. Bobrow (Ed.), "Special Issue on Non-Monotonic Logic," *Artificial Intelligence*, vol. 13, 1980.
- [BrP81] C. A. Brown and P. W. Purdom, Jr., "An Average Time Analysis of Backtracking," *SIAM J. on Computing*, Vol. 10, No. 3, pp. 583-593, 1981.
- [BrT79] F. M. Brown and S. A. Tarnlund, "Inductive Reasoning on Recursive Equations," *Artificial Intelligence*, vol. 12, pp. 207-220, 1979.
- [BuM78] B. G. Buchanan and T. M. Mitchell, "Model-Directed Learning of Production Rules," *Pattern-Directed Inference Systems*, D. A. Waterman and F. Hayes-Roth (Eds.), Academic Press, 1978.
- [Bur86] M. H. Burstein, "Concept Formation by Incremental Analogical Reasoning and Debugging," *Machine Learning: An Artificial Intelligence Approach . Volume II*, R. S. Michalski, J. G. Carbonell, and T. M. Mitchell (Eds.), Morgan Kaufmann, Los Altos, 1986.
- [CMMS83] J. G. Carbonell, R. S. Michalski, and T. M. Mitchell, "An Overview of Machine Learning," *Machine Learning*, R. S. Michalski, J. G. Carbonell, and T. M. Mitchell (Eds.), Tioga, 1983.
- [Car83] J. G. Carbonell, "Learning by Analogy: Formulating and Generalizing Plans from Past Experiences," *Machine Learning*, R. S. Michalski, J. G. Carbonell, and T. M. Mitchell (Eds.), Tioga, 1983.
- [CoF82] P. R. Cohen and E. A. Feigenbaum (Eds.), *The Handbook of Artificial Intelligence*, William Kaufmann, Los Altos, 1982.
- [Com79] D. Comer, "The Ubiquitous B-tree," *Computing Surveys*, vol. 11, no. 2, pp. 121-157, June 1979.
- [Dal82] R. Davis and D. B. Lenat, *Knowledge-Based Systems in Artificial Intelligence*, McGraw-Hill, 1982.
- [Dav78] R. Davis, "Knowledge Acquisition in Rule-based Systems: Knowledge About Representations as a Basis for System Construction and Maintenance," *Pattern-directed Inference Systems*, D. A. Waterman and F. Hayes-Roth (Eds.), Academic Press, New York, 1978.
- [Dem77] N. Dershowitz and Z. Manas, "The Evolution of Programs: Automatic Program Modification," *IEEE Trans. on Software Engineering*, vol. SE-3, no. 6, pp. 377-385, 1977.
- [Den70] P. J. Denning, "Virtual Memory," *Computing Surveys*, vol. 2, pp. 153-189, 1970.
- [DivM83] T. G. Dietterich and R. S. Michalski, "A Comparative Review of Selected Methods for Learning from Examples," *Machine Learning*, R. S. Michalski, J. G. Carbonell, and T. M. Mitchell (Eds.), Tioga, 1983.
- [Die80] T. G. Dietterich, "Applying General Induction Methods to the Card Game Eleusis," *AAAI/4*, pp. 218-220, 1980.
- [Doyle79] J. Doyle, "A Truth Maintenance System," *Artificial Intelligence*, vol. 12, pp. 231-272, 1979.
- [Dri77] S. E. Dreyfus and A. M. Law, *The Art and Theory of Dynamic Programming*, Academic Press, New York, 1977.
- [Eas58] W. L. Eastman, "A Solution to the Traveling-Salesman Problem," presented at the *American Summer Meeting of the Econometric Society*, Cambridge, Mass., Aug. 1958.
- [ER86] M. A. Elfrimson and T. C. Ray, "A Branch-and-Bound Algorithm for Plant Location," *Operations Research*, Vol. 14, pp. 361-368, 1986.
- [Eva88] T. G. Evans, "A Program for the Solution of Geometric Analogy Intelligence Test Questions," *Semantic Information Processing*, M. L. Minsky (Ed.), MIT Press, Cambridge, 1988.
- [FHN72] R. E. Fikes, P. E. Hart, and N. J. Nilsson, "Learning and Executing Generalized Robot Plans," *Artificial Intelligence*, Vol. 3, pp. 251-288, 1972.
- [Fel71] W. Feller, *An Introduction to Probability Theory and its Applications*, Vol. II, 2nd edition, John Wiley & Sons, Inc., 1971.
- [Fel74] D. Ferrati, "Improving Locality by Critical Working Sets," *CACM*, vol. 17, no. 11, pp. 614-620, Nov. 1974.

- [Fer75] D. Ferrari, "Tailoring Programs to Models of Program Behavior," *IBM Journal of Research and Development*, vol. 19, no. 3, pp. 244-251, May 1975.
- [Fer76] D. Ferrari, "The Improvement of Program Behavior," *IEEE Computer*, pp. 38-47, Nov. 1976.
- [FIP79] P. C. Fisher and R. L. Probert, "Storage Reorganization Techniques for Matrix Computation in a Paging Environment," *CACM*, vol. 22, no. 7, pp. 405-415, July 1979.
- [Gar79] M. R. Garey and D. S. Johnson, *Computers and Intractability: A Guide to the Theory of NP-Completeness*, W. H. Freeman and Company, San Francisco, 1979.
- [GarN72] R. S. Garfinkel and G. L. Nemhauser, *Integer Programming*, Wiley, New York, 1972.
- [Gar73] R. Garfinkel, "On Partitioning the Feasible Set in a Branch-and-Bound Algorithm for the Asymmetric Traveling-Salesman Problem," *Operations Research*, Vol. 21, No. 1, pp. 340-342, 1973.
- [Geo72] A. M. Geoffrion and R. E. Marsten, "Integer Programming Algorithms: A Framework and State-of-the-Art Survey," *Management Science*, Vol. 18, No. 9, pp. 485-491, May 1972.
- [Get70] H. Greenberg and R. Hegerich, "A Branch Search Algorithm for the Knapsack Problem," *Management Science*, vol. 16, no. 5, pp. 327-332, 1970.
- [HaG71] D. J. Hatfield and J. Gerald, "Program Restructuring for Virtual Memory," *IBM Systems Journal*, vol. 10, no. 3, pp. 168-192, 1971.
- [HoS74] E. Horowitz and S. Sahni, "Computing Partitions with Applications to the Knapsack Problem," *J. ACM*, vol. 21, no. 2, pp. 277-292, 1974.
- [HoS78] E. Horowitz and S. Sahni, *Fundamentals of Computer Algorithms*, Computer Science Press, 1978.
- [Iba73] T. Ibaraki, "Solvable Classes of Discrete Dynamic Programming," *J. of Math. Analysis and Applications*, vol. 43, pp. 642-663, 1973.
- [Iba76a] T. Ibaraki, "Computational Efficiency of Approximate Branch-and-Bound Algorithms," *Math. of Oper. Research*, Vol. 1, No. 3, pp. 287-298, 1976.

- [Iba76b] T. Ibaraki, "Theoretical Comparisons of Search Strategies in Branch-and-Bound Algorithms," *Int. Jt. of Comp. and Inf. Sci.*, Vol. 5, No. 4, pp. 315-344, 1976.
- [Iba77] T. Ibaraki, "The Power of Dominance Relations in Branch-and-Bound Algorithms," *JACM*, Vol. 24, No. 2, pp. 264-270, 1977.
- [Iba78] T. Ibaraki, "Depth-n Search in Branch-and-Bound Algorithms," *Int. J. of Comp. and Inf. Sci.*, Vol. 7, No. 4, pp. 315-343, 1978.
- [IgS65] E. Ignall and L. Schrage, "Application of the Branch and Bound Technique to Some Flow-shop Scheduling Problems," *Operations Research*, vol. 11, pp. 400-412, 1963.
- [InK77] G. Ingagiola and J. Korsch, "A General Algorithm for One Dimensional Knapsack Problems," *Operations Research*, Vol. 25, No. 5, pp. 752-759, 1977.
- [KarH77] R. Karp and M. Held, "Finite State Processes and Dynamic Programming," *SIAM J. on Applied Math.*, vol. 15, pp. 803-818, 1967.
- [Kar76] R. M. Karp, *Algorithms and Complexity: New Direction and Recent Results*, J. F. Traub (Ed.), Academic Press, pp. 1-18, 1976.
- [Kli71] R. Kling, "A Paradigm for Reasoning by Analogy," *Artificial Intelligence*, vol. 2, pp. 147-178, 1971.
- [KnM75] D. E. Knuth and R. W. Moore, "An Analysis of Alpha-Beta Pruning," *Artificial Intelligence*, Vol. 6, pp. 283-326, 1975.
- [KoS74] W. Kohler and K. Steiglitz, "Characterization and Theoretical Comparison of Branch-and-Bound Algorithms for Permutation Problems," *JACM*, Vol. 21, No. 1, pp. 140-156, 1974.
- [Kor80] R. E. Korf, "Toward a Model of Representation Changes," *Artificial Intelligence*, Vol. 14, pp. 41-78, 1980.
- [KuK83] V. Kumar and L. Kanal, "A General Branch-and-Bound Formulation for Understanding and Synthesizing And/Or Tree Search Procedures," *Artificial Intelligence*, Vol. 21, pp. 179-195, 1983.
- [LBS83] P. Langley, C. L. Bradshaw, and H. A. Simon, "Rediscovering Chemistry with the BACON System," *Machine Learning*, R. S. Michalski, J. G. Carbonell, and T. M. Mitchell (Eds.), Morgan, 1983.

- [HK79] D. B. Lenat, P. Hayes-Roth, and P. Khatri, "Cognitive Economy in Artificial Intelligence Systems," *IJCAI 6*, pp. 531-536, 1979.
- [LaD60] A. H. Land and A. Doig, "An Automatic Method for Solving Discrete Programming Problems," *Econometrica*, Vol. 28, pp. 497-520, 1960.
- [LaW66] E. L. Lawler and D. W. Wood, "Branch-and-Bound Methods: A Survey," *Operations Research*, Vol. 14, pp. 699-719, 1966.
- [Len81] P. Langley, "Data-Driven Discovery of Physical Laws," *Cognitive Science*, Vol. 5, pp. 31-54, 1981.
- [Len76] J. Lenstra, "Sequencing by Enumerative Methods," *Math. Centre Tract 89*, Mathematisch Centrum, Amsterdam, 1976.
- [Len82] D. B. Lenat, "The Nature of Heuristics," *Artificial Intelligence*, Vol. 19, pp. 189-249, 1982.
- [Len83a] D. B. Lenat, "Theory Formation by Heuristic Search: The Nature of Heuristics II: Background and Examples," *Artificial Intelligence*, Vol. 21, pp. 31-59, 1983.
- [Len83b] D. B. Lenat, "EURISKO: A Program that Learns New Heuristics and Domain Concepts: The Nature of Heuristics III: Program Design and Results," *Artificial Intelligence*, Vol. 21, pp. 61-98, 1983.
- [Li85] G.-J. Li, *Parallel Processing of Combinatorial Search Problems*, Ph. D. Thesis, Purdue University, W. Lafayette, 1985.
- [LiW84] G. J. Li and B. W. Wah, "Computational Efficiency of Parallel Approximate Branch-and-Bound Algorithms," *Proc. Int'l Conf. on Parallel Processing*, Aug. 1984, pp. 473-480.
- [LiW85] G. J. Li and B. W. Wah, "MANIP-2: A Multicomputer Architecture for Evaluating Logic Programs," *Proc. Int'l Conf. on Parallel Processing*, IEEE, pp. 122-130, June 1985.
- [MUB83] T. M. Mitchell, P. E. Uigoff, and R. B. Baerji, "Learning by Experimentation: Acquiring and Refining Problem-Solving Heuristics," *Machine Learning*, R. S. Michalski, J. G. Carbonell, and T. M. Mitchell [Eds.], Tioga, 1983.
- [MaM78] A. Martelli and U. Montanari, "Optimizing Decision Trees Through Heuristically Guided Search," *CACM*, Vol. 21, pp. 1025-1039, 1978.

- [MaW75] Z. Manya and R. J. Waldinger, "Knowledge and Reasoning in Program Synthesis," *Artificial Intelligence*, vol. 6, pp. 175-208, 1975.
- [McC89] A. C. McKellar and E. G. Coffman, Jr., "Organizing Matrices and Matrix Operations for Paged Memory Systems," *CACM*, vol. 12, no. 3, pp. 153-164, 1969.
- [McD78] D. McDermott, "Planning and Acting," *Cognitive Science*, vol. 2, pp. 71-109, 1978.
- [McD79] J. McDermott, "Learning to Use Analogies," *IJCAI 6*, pp. 568-576, 1979.
- [MiC80] R. S. Michalski and R. L. Chilausky, "Learning by Being Told and Learning From Examples: An Experimental Comparison of the Two Methods of Knowledge Acquisition in the Context of Developing an Expert System for Soybean Disease Diagnosis," *Int. J. of Policy Analysis and Information Systems*, vol. 4, pp. 125-161, 1980.
- [MiC83] R. S. Michalski, "A Theory and Methodology of Inductive Learning," *Machine Learning*, R. S. Michalski, J. G. Carbonell, and T. M. Mitchell [Eds.], Tioga, 1983.
- [Mic88] R. S. Michalski, "Understanding the Nature of Learning: Issues and Research Directions," *Machine Learning: An Artificial Intelligence Approach - Volume II*, R. S. Michalski, J. G. Carbonell, and T. M. Mitchell [Eds.], Morgan Kaufmann, Los Altos, 1988.
- [Min75] M. Minsky, "A Framework for Representing Knowledge," *The Psychology of Computer Vision*, P. H. Winston [Ed.], McGraw-Hill, New York, 1975.
- [Min84] S. Minton, "Constraint-Based Generalization: Learning Game Playing Plans from Single Examples," *AAAI-84*, pp. 251-254, 1984.
- [Mit70] L. Mitten, "Branch-and-Bound Methods: General Formulation and Properties," *Operations Research*, Vol. 18, pp. 24-34, 1970.
- [Mit82] T. M. Mitchell, "Generalization as Search," *Artificial Intelligence*, Vol. 18, pp. 203-228, 1982.
- [Mot70] D. J. Mostow and F. Hayes-Roth, "Operationalizing Heuristics: Some AI Methods for Assisting AI Programming," *IJCAI 6*, pp. 601-608, 1979.
- [MoN73] J. A. Moore and A. Newell, "How Can MERLIN Understand? Knowledge and Cognition, L. Gregg [Ed.], Erlbaum, Hillsdale, 1973.

- [Mor68] J. E. Morrison, "User Program Performance in Virtual Storage Systems," *IBM Systems Journal*, vol. 5, no. 2, pp. 216-223, 1966.
- [NeS56] A. Newell and H. A. Simon, "The Logic Theory Machine," *IRE Transactions on Information Theory*, vol. 2, pp. 61-70, 1956.
- [Nil71] N. J. Nilsson, *Problem Solving Methods in Artificial Intelligence*, McGraw Hill, New York, 1971.
- [Nil80] N. J. Nilsson, *Principles of Artificial Intelligence*, Tioga Publishing Company, Palo Alto, California, 1980.
- [Ohi83] S. Ohlsson, "A Constrained Mechanism for Procedural Learning," *IJCACI*, 83, pp. 426-428, 1983.
- [Pear84] J. Pearl, *Heuristics: Intelligent Search Strategies for Computer Problem Solving*, Addison-Wesley, 1984.
- [Rob85] J. A. Robinson, "A Machine-oriented Logic Based on the Resolution Principle," *J. ACM*, vol. 12, pp. 23-41, 1965.
- [Ros70] S. M. Ross, *Applied Probability Models with Optimization Applications*, Holden-Day, San Francisco, 1970.
- [Sah77] S. Sahni, "General Techniques for Combinatorial Approximation," *Operations Research*, vol. 25, no. 8, pp. 920-936, 1977.
- [Sam59] A. L. Samuel, "Some Studies in Machine Learning Using the Game of Checkers," *IBM J. Research and Development*, vol. 3, pp. 210-226, 1959.
- [Sam87] A. L. Samuel, "Some Studies in Machine Learning Using the Game of Checkers. II - Recent Progress," *IBM J. Research and Development*, vol. 31, pp. 601-617, 1987.
- [Set81] Y. Sekiguchi, "A Unifying Framework of Combinatorial Optimization Algorithms: Tree Programming and its Validity," *J. Operations Research of Japan*, vol. 24, no. 1, 1981.
- [SL74] H. A. Simon and G. Lea, "Problem Solving and Rule Induction: A Unified View," *Knowledge and Cognition*, L. Gregg (Ed.), Lawrence Erlbaum, Hillsdale, 1974.
- [Smi78] A. J. Smith, "Sequentiality and Prefetching in Database Systems," *ACM Trans. on Database Systems*, vol. 3, no. 3, pp. 223-247, Sept. 1978.
- [Smi84] D. R. Smith, "Random Trees and the Analysis of Branch-and-Bound Procedures," *Journal of the ACM*, vol. 31, no. 1, pp. 163-188, 1984.
- [StK83] G. C. Stockman and L. N. Kanal, "Problem Reduction Representation for the Linguistic Analysis of Waveforms," *IEEE Trans. on PAMI*, Vol. 5, No. 3, pp. 287-298, 1983.
- [Sto84] H. S. Stone, "The Average Complexity of Depth-First Search," IBM Research Report RC 10717, 1984.
- [Sus75] G. J. Sussman, *A Computer Model of Skill Acquisition*, American Elsevier, New York, 1975.
- [Ver78] S. A. Vere, "Inductive Learning of Relational Productions," *Pattern-Directed Inference Systems*, D. A. Waterman and F. Hayes-Roth (Eds.), Academic Press, 1978.
- [WLY85] B. W. Wah, G.-J. Li, and C. F. Yu, "Multiprocessing of Combinatorial Search Problems," *IEEE Computer*, vol. 18, no. 6, pp. 93-108, June 1985.
- [Wol84] P. H. Wos, R. Overbeek, E. Lusk, and J. Boyle, *Automated Reasoning: Introduction and Applications*, Prentice-Hall, New Jersey, 1984.
- [WRCo85] L. Wos, G. A. Robinson, and D. F. Carson, "Efficiency and Completeness of the Set of Support Strategy in Theorem Proving," *J. ACM*, vol. 32, pp. 538-541, 1985.
- [Wah84] B. W. Wah and Y. W. Ma, "MANIP-A Multicomputer Architecture for Solving Combinatorial Extremum-Search Problems," *IEEE Trans. on Computers*, vol. C-33, no. 5, pp. 377-390, May 1984.
- [WahY82] B. W. Wah and C. F. Yu, "Probabilistic Modeling of Branch-and-Bound Algorithms," *Proc. COMPSAC*, IEEE, pp. 647-653, 1982.
- [WahY85] B. W. Wah and C. F. Yu, "Stochastic Modeling of Branch-and-Bound Algorithms with Best-First Search," *IEEE Trans. on Software Engineering*, vol. SE-11, no. 9, pp. 922-934, Sept. 1985.
- [Wat70] D. A. Waterman, "Generalization Learning Techniques for Automating the Learning of Heuristics," *Artificial Intelligence*, Vol. 1, No. 1/2, pp. 121-170, 1970.
- [Whi60] D. White, *Dynamic Programming*, Oliver and Boyd, Edinburgh, 1960.

- [Win75] P. H. Winston, "Learning Structural Descriptions from Examples," *The Psychology of Computer Vision*, P. H. Winston [Ed.], McGraw Hill, New York 1975.
- [Win79] P. H. Winston, "Learning and Reasoning by Analogy," *CACM*, Vol. 23, pp. 689-703, 1979.
- [Win84] P. H. Winston, *Artificial Intelligence* Addison-Wesley, Reading, 1984.
- [YuW83] C. F. Yu and B. W. Wah, "Virtual Memory Support for Branch-and-Bound Algorithms," *Proc. of COMPSAC*, pp. 618-628, 1983.
- [YuW84] C. F. Yu and B. W. Wah, "Efficient Branch-and-Bound Algorithms on a Two-Level Memory System," *Proc. of COMPSAC*, pp. 504-514, 1984.

## APPENDICES

### Appendix A: A Prototype of the Searcher

```

;cond
;[null chcl] : best choice has been identified
; delete choice from Activelist
;[setq Activelist (delete choice Activelist)]
;[tramprop Rank choices]
; delete antichoice from Activelist if there is and antichoice
; and it is a member of Activelist
;[setq frame (eval choice)]
;[setq antichoice (eval frame chtlist)]
;[cond [[and antichoice (member antichoice Activelist)]]
;       ;[setq Activelist (delete antichoice Activelist)]
;       ;[tramprop Rank antichoice]]
;[cond
;  [princ "Generation chooses rule "][princ choice]
;  [princ " or rank "][princ choicerank][lterpr]
;  [return choice]
;  ; keep on comparing with other elements of Activelist
;  [setq rule (car chcl)]
;  [setq chcl (cdr chcl)]
;  [cond
;    [> (get 'Rank chcl) choicerank] ; rule has higher rank
;    ;[setq choice chcl] ; make chcl the new choice
;    ;[setq choicerank (get 'Rank chcl)][update the rank
;    ;[go next-crcl]]
;  ]prolog
;  ]prolog
;  ]defun
;***** functions that implement the tester *****
; tester : the tester
;
;[defun tester (frame)
;  ;[null frame] ; all frames crule
;  (prog (data tframe tcrule nrule dframs crule)
;    (setq data (read frame))
;    (setq crule (read frame chtlist))
;    (princ "Testing rule ")[princ crule]
;    (princ " with data ") [princ data][terpr]
;    (cond
;      [(equal data nil)
;       (prog ()]
;         [setq initial (length)
;          next-tree
;          (cond
;            [(null initlist) ; all trees done
;             ;[go available]] ; go evaluate test results
;            [cond
;              ; more search trees left, work on these
;              ;[setq init (car initlist)]]
;            ]next-crcl
;          )]
;        )]
;      )]
;    )]
;  )
;
```





```

[cond [[null brule], [return nil]] ; no parallel rule
; there is a parallel rule
[setq brule [eval brule]]
[setq adets [vref biframa cDelta]]
[setq atype [vref biframa cType]]
[cond [[and (equal adets 'unacceptable) (equal atype 'lhp)]
      [go set-value]]]
];cond
[cond [[and (equal adets 'impossible) (equal atype 'bav)]]
      [go set-value]]]
];cond
; remaining cases can only result in upgrading the ranking
[cond [[vref biframa cDelta] [feature nil]]; delta of brule known
; delta of brule not known so upgrade its ranking
];uprop [Rank 4 brule]
[print "scrubber: "][print brule][print " is parallel to "]
[print adets][print " <rule "]
[print adets][print "> increase rank of "][print brule][temp]
[return nil]; no new decision to scrubage on

; rest is code to handle setting of decision

; set-value
[cond [[null [vref biframa cDelta]]]; delta of brule not known
(vset biframa cDelta adets)
[set biframa cEx [list 'parallel 'o adets]]
[print "scrubber: "][print brule][print " is parallel to "]
[print adets][print " <rule "]
[print adets][print "> "][print brule][print " is "]
[print adets][print "Activelias delta brule Activelias"]
[temprop [Rank brule]
[return [list brule]]]
];cond
; delta of brule known ... set it confirmed by parallelism
[cond [[equal adets [vref biframa cDelta]]]; confirms delta
[return nil]
];contradiction
; contradiction
; child-delta : And implied delta values of children rules
; child-delta : And implied delta values of children rules
(datum child-delta (satd))
(prng [fifmns adets brule biframa nu-rules]
[setq biframa [eval adets]]
[setq adets [vref biframa cDelta]]
[setq blist [vref biframa cChildren]]; list of adets's children
[setq nu-rules all]
next-brule
[cond [[null blist] [return nu-rules]]; all children processed
[setq brule [car blist]]
[setq blist (cdr blist)]
[setq biframa [eval brule]]
[setq bdelta [vref biframa cDelta]]
[cond [[and (equal adets 'impossible) (full bdelta))]
; bdelta not known
[setq biframa cDelta 'impossible]; set implied decision
[defun parent-delta (adets)
; parent-delta : adjust parent rules
[defun parent-delta (adets)

```

```

[prng [fifmns brule biframa blist nu-rules]
[setq biframa eval brule]]
; nothing can be deduced if adets is not 'unacceptable'
[cond [[equal [vref biframa cDelta] 'unacceptable] [return nil]]]
; adets is unacceptable
[setq blist [vref biframa cParent]]
[setq nu-rules nil]; list of jobs to be added to joblist
];cond
[cond [[null blist] [return nu-rules]]];
; all parents done
[setq brule [car blist]]
[setq blist (cdr blist)]
[setq biframa [eval brule]]
[cond [[null [vref biframa cDelta]]]; delta of brule not known
[set biframa cDelta 'unacceptable]; set implied decision
(vset biframa cEx [list 'parent 'of adets])
[print "scrubber: "][print brule]
[print nu-rules][print "> "][print brule]
[print " is unacceptable" ] [terpri]
[seq activism [delta brule Activelias]]
[temprop [Rank brule]
[setq nu-rules (append nu-rules brule)]
[go next-brule]]
];cond
; delta of brule is known, so try to confirm
[cond [[equal [vref biframa cDelta] 'unacceptable]; confirms delta
[go next-brule]]
];contradiction
; child-delta : And implied delta values of children rules
; child-delta : And implied delta values of children rules
(datum child-delta (satd))
(prng [fifmns adets brule biframa nu-rules]
[setq biframa [eval adets]]
[setq adets [vref biframa cDelta]]
[setq blist [vref biframa cChildren]]; list of adets's children
[setq nu-rules all]
next-brule
[cond [[null blist] [return nu-rules]]; all children processed
[setq brule [car blist]]
[setq blist (cdr blist)]
[setq biframa [eval brule]]
[setq bdelta [vref biframa cDelta]]
[cond [[and (equal adets 'impossible) (full bdelta))]
; bdelta not known
[setq biframa cDelta 'impossible]; set implied decision
[defun parent-delta (adets)
; parent-delta : adjust parent rules
[defun parent-delta (adets)

```

```

;test bframe cEw [list 'parent' of rule]
(princ "scrugger: ") (princ brule)
(princ " is a child of impossible c-rule ")
(princ crule) (princ ">") (princ brule)
(princ " is impossible") (terpri)
;test Activists (delet brule Activists)
(terpri) "Rank brule"
;testq Activeists (appendd no-crules brule)
;testq Activeists (appendd no-crules brule)
;go next-brule
)cond
(cond [and [equal adelta 'impossible] [equal bdelta 'impossible]
; decision confirmed
; go next-brule]
)cond
; go next-brule
)cond
; equal adelta 'impossible'
; contradiction
;test bframe cEv [append [test bframe cEv]
; [list 'child' of rule 'implies a delta]]
(go next-brule)
)cond
; test bframe cEv [append [test bframe cEv]
; [list 'child' of rule 'implies a delta]]
;delta is nil
;test bframe 'C-Delta' 'acceptable'; set implied decision
;test bframe cEv [list 'child' of rule]
(princ "scrugger: ") (princ brule)
(princ " is a child of acceptable c-rule ")
(princ crule) (princ ">") (princ brule)
(princ crule) (princ ">") (princ brule)
(princ " is acceptable/impossible") (terpri)
;test Activists (delet brule Activists)
(terpri) "Rank brule"
;testq Activeists (appendd no-crules brule)
;go next-brule
)cond
;cond [and [equal adelta 'acceptable] [member bdelta 'acceptable 'impossible]]
; confirm delta
; go next-brule
)cond
; contradiction
;test bframe cEv [append [test bframe cEv]
; [list 'child' of rule 'implies 'acceptable' or 'unacceptable']]
;go next-brule
)cond
;cond [and [equal adelta 'acceptable]
; go next-brule]
;test bframe cEv [append [test bframe cEv]
; [list 'child' of rule 'implies 'acceptable' or 'unacceptable']]
;go next-brule
)cond
;cond [and [equal adelta 'unacceptable]
; go next-brule]
;test bframe cEv [append [test bframe cEv]
; [list 'child' of rule 'implies 'unacceptable']]
;go next-brule
)cond
;cond [and [equal adelta 'unacceptable]
; go next-brule]
;test bframe cEv [append [test bframe cEv]
; [list 'child' of rule 'implies 'unacceptable']]
;go next-brule
)cond
;cond [and [equal adelta 'unacceptable]
; go next-brule]
;test bframe cEv [append [test bframe cEv]
; [list 'child' of rule 'implies 'unacceptable']]
;go next-brule
)cond
;cond [and [equal adelta 'unacceptable] [null bdelta]]
;test bframe (no-delta bframe bframe)
;test bframe cEv [list 'child' of rule]
(princ "scrugger: ") (princ brule)
(princ " is a child of unacceptable c-rule ")

```

```

;princ crule) (princ "> partition shows ")
;princ brule) (princ " is ")
;print (test bframe cDelt) (terpri)
;cond [equal (test bframe cDelt) 'acceptable) (clump1 bframe)]
;testq Activists (delet brule Activists)
;terpri) "Rank brule"
;testq Activeists (appendd no-crules brule)
;go next-brule]
;cond
;go next-brule)
;go next-brule
)cond
; do-divide : find instances of parent that satisfy a crule
;defun do-divide (frame bframe)
;prop () 
; And instances of adrule
;seq link (test bframe cCond)
;test bframe cPosist (cons ' + (position link (cdr (test bframe cPosist)))))
;test bframes cTestList (cons ' + (position link (car (test bframe cPosist))))))
;test bframe cNullList (cons 0 (position link (car (test bframe cNullList))))
;test bframe cNumpos (1- (length (test bframe cPosist)))
;test bframe cNumneg (1- (length (test bframe cNegat)))
;test bframe cNumall (1- (length (test bframe cNulls)))
;evaluate instance set of rule
;not really the same as for tester
;include hudge
;cond [[> (test bframe cNumneg) 0] ; there are +ve instances
;return unacceptable]
;cond
;ao-va instances
; code here is a kludge
;return 'acceptable'
)cond
;code here is a kludge
;ao-va instances
;code here is a kludge
;return 'acceptable'
)cond
;ao-va and no-va instances
;cond [[> (test bframe cNumpos) 0] ; there are null instances
;code here also a kludge
;return 'acceptable']
)cond
;ao-va instances of any kind
;again, code here is a kludge
;return 'impossible'
)cond
;defun portion [link cRule]
;prop ([& y q q num alist]
;princ link) (terpri)
;princ link)

```





```

(setq alfa (vref afname cl,b))
(setq blist (vref corevalst lev))
next-brule
(cons (null blist) ;|| rule not in map - either 'acceptable' or 'impossible
; adjust ranking of rule
(purpdrop <Rank rank &rule)
(go next-level))
(setq bifrms (eval (car blist)))
(setq blist (cdr blist))
; check if car &rule) matches (cdr &rule)
; seq bcondlist (car (vref afname cl,b))
(next-brule)
(cons
  (null bcondlist) ;|| rule found
  (vref afname cParallel (vref bifrms cName))
  (vref bifrms cParallel (vref afname cName)))
(go next-level)
; null
; ; try next b-condition
(next-bcondition (car bcondlist))
(first bcondlist (cdr bcondlist))
(cons [[number bcondition alfa] (go next-bcondition)]
  nil
  (cons
    (cond
      (prog)
      )
    )))
; set cParallel for newly generated rules
; (implicit: if cParallel == nil implies parallel rule
; not interesting [i.e. has delta value of either
; acceptable or impossible
; acceptible or 'impossible
; defun setc()
; (prog ([lev marker])
; (setq lev 1)
; (setq marker lmarker)
; next-level
; (cond [(and (< lev marker) (vref bifrmList lev))
;        (cond [[and (< lev marker) (vref bifrmList lev))
;               ; below depth limit and
;               ; there are actually rules on this level
;               (bindanti lev bifrmList)
;               (setq lev (+ lev))]
;               (go next-leve)])
;        )
;      )))
; setc
; setc : set cParallel for newly generated rules
; (implicit: if cParallel == nil implies parallel rule
; not interesting [i.e. has delta value of either
; acceptable or impossible
; acceptible or 'impossible
; defun setc()
; (prog ([lev marker])
; (setq lev 1)
; (setq marker lmarker)
; next-level
; (cond [(and (< lev marker) (vref bifrmList lev))
;        (cond [[and (< lev marker) (vref bifrmList lev)]
;               ; below depth limit and
;               ; there are actually rules on this level
;               (bindanti lev bifrmList)
;               (setq lev (+ lev))]
;               (go next-leve)])
;        )
;      )))
; setc
; setc : set cAnti for all newly-generated rules
; (define setanti()
; (prog ([lev marker])
; ; set cAnti values for alpha rules
; (setq lev 1)
; (setq marker lmarker)
; next-level
; (cond [[and (< lev marker) (vref bifrmList lev)]
;       ; below depth limit and
;       ; there are actually rules on this level
;       (bindanti lev bifrmList)
;       (setq lev (+ lev))]
;       (go next-leve)])
; ; set cAnti for beta rules
; (setq lev 1)
; next-level
; (cond [[and (< lev marker) (vref bifrmList lev)]
;       ; below depth limit and
;       ; there are actually rules on this level
;       (bindanti lev bifrmList)
;       (setq lev (+ lev))]
;       (go next-leve)])
; ; bindanti : bind cAnti values for all new rules on given level
; ; Bindanti : Bind cAnti values for all new rules on given level
; (define bindanti(lev levlist)
; (defun bindanti(lev levlist)
; (purp (list afname blist bifrms antiList)
; (setq alfa (vref levlist lev))
; next-apply
; (cond [(null alista) (return v)] ; all rules on lev done
; (setq afname (eval (car alista)))
; (setq alista (cdr alista)))
; (setq alista (cdr alista)))
; (cond [(vref afname (car alista)) (go next-apply)]; canti known
; ; canti not known
; (setq anilis (makeanti (car (vref afname cl,b))))))
; ; match anilis against other c-rules on same level
; (setq blist (vref levlist lev))
; (cond [(vref afname (car alista)) (go next-apply)]; no antirule for anilis
; ; no try next rule
; (setq bifrms (eval (caar blist)))
; (setq blist (cdr blist)))
; (cond
;   (bmatch anilis (cdr (vref bifrms cl,b)); antirule found
;   (vref afname canti (vref bifrms cName))
;   (vref bifrms canti (vref afname cName)))
;   ))))
; setc
; setc : set cAnti for all newly-generated rules
; (define setanti()
; (prog ([lev marker])
; ; set cAnti values for alpha rules
; (setq lev 1)
; (setq marker lmarker)
; next-level
; (cond [[and (< lev marker) (vref bifrmList lev)]
;       ; below depth limit and
;       ; there are actually rules on this level
;       (bindanti lev bifrmList)
;       (setq lev (+ lev))]
;       (go next-leve)])
; ; set cAnti for beta rules
; (setq lev 1)
; next-level
; (cond [[and (< lev marker) (vref bifrmList lev)]
;       ; below depth limit and
;       ; there are actually rules on this level
;       (bindanti lev bifrmList)
;       (setq lev (+ lev))]
;       (go next-leve)])
; ; bindanti : bind cAnti values for all new rules on given level
; ; Bindanti : Bind cAnti values for all new rules on given level
; (define bindanti(lev levlist)
; (defun bindanti(lev levlist)
; (purp (list afname blist bifrms antiList)
; (setq alfa (vref levlist lev))
; next-apply
; (cond [(null alista) (return v)] ; all rules on lev done
; (setq afname (eval (car alista)))
; (setq alista (cdr alista)))
; (setq alista (cdr alista)))
; (cond [(vref afname (car alista)) (go next-apply)]; canti known
; ; canti not known
; (setq anilis (makeanti (car (vref afname cl,b))))))
; ; match anilis against other c-rules on same level
; (setq blist (vref levlist lev))
; (cond [(vref afname (car alista)) (go next-apply)]; no antirule for anilis
; ; no try next rule
; (setq bifrms (eval (caar blist)))
; (setq blist (cdr blist)))
; (cond
;   (bmatch anilis (cdr (vref bifrms cl,b)); antirule found
;   (vref afname canti (vref bifrms cName))
;   (vref bifrms canti (vref afname cName)))
;   ))))

```

```

[go next-avalue]
|> (cond
|>   ;cond
|>   ;prog
|>   ;defun
|>   ;match : match two sets of d-conditions
|>   ;(order of d-conditions may be scrambled)
|>   ;dfunc (match (value blist))
|>   ;(cond (xcond) (length alists) (length blists)) (return nil)
|>   ;next-second
|>   ;(cond (null alists) (return nil)); all d-conditions in alists are
|>   ;present in blist .. match
|>   ;(setq second (car alists))
|>   ;(setq alists (cdr alists))
|>   ;(cond
|>     ;member second blist (go next-second)); so far no good
|>   ;(return nil); one d-condition in alists missing from blist
|>   ;(cons possibly match
|>     ;cons
|>     ;prog
|>     ;defun
|>   ;makeanti : create antilhs corresponding to the
|>   ;warning : lhs must not contain 't' or SBR's,r
|>   ;(defun makeanti (lhs)
|>     ;(prog (xantilhs stcond cond-num)
|>       ;(setq xantilhs nil)
|>       ;next-second
|>       ;(cond ((null lrhs) (return xantilhs)); all second in the changed into
|>             ; antilhs in xantilhs
|>             ;(setq second (cadr lrhs))
|>             ;(setq lrhs (cdr lrhs))
|>             ; And bind in Conslist that matches second
|>             ;(setq blist (Antilhs))
|>             ;next-second
|>             ;(cond (null blist) ; no match .. oh-oh
|>               ;(princ "ERROR") (terpri)
|>               ;(print "lhs" "-") (princ lrhs) (terpri)
|>               ;(princ "Antilhs") (princ Antilhs) (terpri)
|>               ;(error))
|>             ;(setq lrhs (car blist))
|>             ;(cond
|>               ;(equal second (car lrhs)); match
|>               ; replace second by its anticond in antilhs
|>               ;(setq antilhs (append lrhs antilhs (cadr lrhs)))
|>               ;(go next-second)
|>             ;(setq blist (cdr blist)); blist doesn't match

```

```

[go next-bcond]; try next bcond
|> (cond
|>   ;prog
|>   ;defun
|>   ; functions to update the Anticon list
|>   ; needs to run after each update of Condlist
|>   ; du-and : update Anticon list
|>   ;(defun du-and ()
|>     ;(prog (second alist)
|>       ;(setq alist Condlist)
|>       ;next-second
|>       ;(cond ((null alist) (return nil)); all done ...
|>             ;(setq cond (car alist))
|>             ;(setq alist (cdr alist))
|>             ;(setq Antilhs (append Antilhs (list cond (mungeise cond))))
|>             ;(go next-second)
|>           ;(princ "next-bcond")
|>           ;(prog
|>             ;defun
|>             ;(defun mungeise (second)
|>               ;(defun mergeop (prop1 prop2)
|>                 ;(prog (op0 flag1 flag2 op1 var1 prop1 op2 var2 prop2)
|>                   ;(setq op0 (car second)) ; > or = or <
|>                   ;(setq flag1 (cadr second)) ; (setf t t)
|>                   ;(setq flag2 (noddle second)) ; (setf f t)
|>                   ;(setq op1 (car flag1))
|>                   ;(incf sub-particles of flag1)
|>                   ;(setq op1 (cadr flag1))
|>                   ;(extract sub-particles of flag2
|>                   ;(setq op2 (car flag2))
|>                   ;(setq var2 (cadr flag2))
|>                   ;(setq prop2 (addr flag2))
|>                   ;(cond
|>                     ;(and (equal var2 'x) (nequal var2 'y))
|>                       ;(tag2 not 'test x something) or 'test y something
|>                     ;(cond ((equal var1 'x) (setq du-and y))
|>                           ;(equal var1 'y) (setq du-and x))
|>                     ;(return (list op0 (list (get du-and prop1) flag2)))
|>                   ;(cond
|>                     ;(tag2 involves state variable (includes pseudo-state)
|>                     ;(cond
|>                       ;(equal prop1 prop2)
|>                         ;(cond ((equal var1 var2) (error)); same var of same node
|>                           ;(cond
|>                             ;(equal op0 '<) (return (list '> flag1 flag2))
|>                             ;(equal op0 '>) (return (list '< flag1 flag2))
|>                             ;(equal op0 '>) (return (list '< flag2))

```

```

;[equal op0 'equal] [return second]
|[error]
|second
|begin
|cond
|  ;prop1 := prop2

|cond [equal var1 var2]
|[equal var1 'x] [(setq var1 'y)]
|[equal var2 'y] [(setq var2 'x)]
|[error]
|  (return (list op0 (list 'get var-var prop1)
|                (list 'get var-var prop2)))
|;equal

|cond
|  ;prop1 := prop2 and var1 := var2
|  ;equal op0 '>' [return (list 'lessp var1 prop2)
|                  (list 'get var2 prop1)]]
|;cond

|cond
|  ;Equal op0 '<' [return (list 'lessp var1 prop2)
|                  (list 'get var2 prop1)]]
|;cond

|cond
|  ;Equal op0 'equal' [return (list 'equal (list 'get var1 prop2)
|                                 (list 'get var2 prop3))]

;functions that implement the asseyer

;assey : the asseyer

;defun assey (var1var2)
|(prog (list symb blist bysymb nsecond-list nvar-list sec-list)
|(princ "Asseyer") (terpri)
|(test nsecond-list nil)
|(test nvar-list (terpri var1var2))
|(test blist (append nvar-list bysymbol-list))
|(test sec-list (terpri Eymbol-list))
|(test Eymbol-list (terpri var1var2))
|(test nvar-list (terpri var1var2))

;nsecond-list : all symb's done
|(test symb (car alists))
|(test alists (cdr alists))
|(test sec-list (cdr sec-list)); do not compare a symbol with itself
|(test blist sec-list); compare symb with all bysymb's in blist

```

```

;next-baymb
|cond [(null blist) (go next-symb)]
|(setq baymb (car blist))
|(setq blist (cdr blist))
|cond

;Request Problem "value"
|cond [notg nsecond (valuefactcond symb baymb)]
|[setq nsecond-list (append nsecond-list nsecond)]]
|[equal Problem 'ide]
|cond [notg nsecond (listcond symb baymb)]
|[setq nsecond-list (append nsecond-list nsecond)]]
|[error]
|  (go next-baymb)
|broad
|;iprog
|;defun
;gen-cond : form the three conditions involving the two symbols
;defmacro gen-cond (symbol1 symbol2)
|(list 'list 'list 'lessp symbol1 symbol2)
|(list 'list 'list 'greaterp symbol1 symbol2)
|(list 'list 'list 'equal symbol1 symbol2)
|list
|;defmacro

;findcond : find conditions in a value description
;defun findcond (symbol1 symbol2)
|(prog (list symb blist bysymb)
|(setq var1 (cadr symb))
|(setq var2 (cadr symb)))
|(setq symb (cadr (caddr symb)))
|(setq bysym (cadr (caddr symb)))
|(setq bysymb (cadr (caddr symb)))
|(setq bysym (cadr (caddr symb)))
|(setq bysym (cadr (caddr symb)))
|(cond [and (equal var1 q1) (equal bysym 'q1)])
|[or (and (equal var1 symb) (equal bysym symb))
|[and (equal symb var1) (equal bysym symb)]
|[or (and (equal symb symb) (equal bysym symb))
|[both symb and bysymb are problem constants
|(return nil)]
|cond
;at least one of symb and bysymb is not a problem constant
|cond [(and (equal symb q1)
|(equal bysym 'q1))
|[one problem constant and one level constant
|(return nil)]
|cond
|cond [(and (equal bvar 'q1)
|(equal var1 'q1))
|[one problem constant and one level constant
|(return nil)]
|cond
|cond [(and (equal bvar 'q1)
|(member bvar '(x y)))
|[one problem constant and one level constant
|(return nil)]

```



```

(princ "new conditions: ") (princ au-const) [terpri]
;return au-const]
;hprog
;defun

; vvarh-fal : bad conditions between vvar and var
;-----[defun vvarh-fal (xysymb lyemb)
;-----(prog (xprop (prop ambi laubt suc-const))
;-----  (setq xprop (cdr (caddr xysymb)))
;-----  (setq iprop (car (caddr lyemb)))
;-----  (setq xsubt (set xprop 'SSubt))
;-----  (setq lsubt (get xprop 'LSubt))
;-----  (cond
;-----    ((and (equal xsubt 'shyb) 
;-----      (equal lsubt 'shyb))
;-----      (return lsubt 'enum))
;-----    (return (vvarh-fal xysymb lyemb prop iprop)))
;-----  ;cond
;-----  (cond ((and (equal xsubt 'true) (equal lsubt 'enum))
;-----        (return (vvarh-fal xysymb lyemb prop iprop)))
;-----      (return nil))
;-----    (cond ((and (equal xsubt 'true) (equal lsubt 'false))
;-----          (return (vvarh-fal xysymb lyemb prop iprop)))
;-----        (return nil))
;-----      (hprog
;-----        ;defun

;----- vvarh-fal : bad condition between state variable and summary variable
;-----[defun vvarh-fal (xysymb xysymb xprop sprop)
;-----  (prog (xprop iprop sumb laubt au-const zadr bdir)
;-----    (setq sprop (cddr (caddr xysymb)))
;-----    (setq iprop (cadr (caddr lyemb)))
;-----    (setq sumb (get xprop 'Subt))
;-----    (setq laubt (get iprop 'Subt))
;-----    (setq adir (set sprop 'Dir))
;-----    (setq bdir (set iprop 'Dir))
;-----    (cond
;-----      ((and
;-----        (equal sprop bprop)
;-----        (equal adir bdir)
;-----        (equal (get sprop 'Hypotype) 'at))
;-----      ) ;and
;-----      (cond
;-----        ((equal (caddr xysymb) 'x)
;-----          (setq au-const (gen-cond xysymb lyemb))
;-----          (setq au-const (gen-cond lyemb xysymb)))
;-----        ));cond
;-----        (princ "answer: ") (princ xysymb) (princ " and ")
;-----        (princ lyemb) (princ " have identical properties ") (terpri)
;-----        (princ " new conditions: ") (princ au-const) [terpri]
;-----        (return au-const)
;-----      );cond
;-----      (cond (and (equal sumb 'shyb) (equal laubt 'shyb))
;-----        (return nil))
;-----      (cond (and (equal sumb 'shyb) (equal laubt 'shyb))
;-----        (return nil))
;-----      );cond
;-----    );defun

```

```

(cond
  [and [equal aubt 'Any]
       [member bprop [set aprop 'DefVar]]
       [equal cdr bdt]
       [setq [cdr symb] (code symb)]]
    ;and
    [setq nu-const [gen-cond bdt symb symb])
    (princ "Answer: ") [princ aprop] [princ " is a hybrid of ")
    (princ bprop) [terpri]
    (princ "new conditions") [princ nu-const] [terpri]
    [princ nu-const])
  );cond

  [cond
    [and [equal bdtb 'Any]
         [member bprop [set aprop 'DefVar]]
         [equal cdr bdt]
         [setq [cdr symb] (code symb)]]
    ;and
    [setq nu-const [gen-cond symb symb)]
    (princ "Answer: ") [princ aprop] [princ " is a hybrid of ")
    (princ bprop) [terpri]
    (princ "new conditions") [princ nu-const] [terpri]
    [princ nu-const])
  );cond

  [return nil]
);function

;bad-cond : Bad condition between one state variable and our
;hybrid problem constraint

;bad-av-av : Bad condition between one state variable and our
;hybrid problem constraint

;defun
;bad-av-av (symbol symb)
;prog [sprop nu-const aprop]
;setq sprop (cad (caddr symb))
;setq aprop (cddr (caddr symb))
;cond [nequal (set aprop 'Sub1) 'true]
;      [return nil]
;  );cond
;cond [[member aprop [set aprop 'LimVal]]
;      [setq nu-const [gen-cond symb symb))]
;      (princ aprop) [terpri]
;      (princ aprop) [princ aprop] [princ " is a limit of ")
;      (princ "Answer: ") [princ aprop] [princ nu-const]
;  );cond
;cond
;      [return nil]
;    );cond
;  );defun

;bad-av-av : Bad condition between two state variables
;defun
;bad-av-av (symbol symb)
;prog [bprop bprop aubt bdt symb nu-const]
;setq aprop (cad (caddr symb))
;setq bprop (cad (caddr symb))
;cond
;  [and [equal aubt 'x] [member aubt '(x y)]]
;  [setq nu-const [gen-cond symb symb)]
;  (princ "Answer: identical property") [terpri]
;  (princ "new conditions") [princ nu-const] [terpri]
;  [princ nu-const])
;  );cond

;cond
;  [and [equal aprop bprop] [equal (code symb) 'x]]
;  [setq nu-const [gen-cond symb symb)]
;  (princ "Answer: identical property") [terpri]
;  (princ "new conditions") [princ nu-const] [terpri]
;  [princ nu-const])
;  );cond
;  [cond
;    [and [equal aprop bprop] [equal (code symb) 'x]]
;    [setq aubt (get aprop 'Sub1)]
;    [setq bdtb (get bprop 'Sub1)]
;    [setq nu-const [gen-cond aubt bdtb))]
;    (princ "Answer: ") [princ aubt] [princ bdtb]
;    [return [ind-av-av symb symb]]
;  );cond
);cond

```

```

(cond [(and (equal aubt 'shb) (equal boubt 'src))
       (return [iuvvv-fin0 hyemb symbbl])])
;cond
(cond [(and (equal aubt 'src) (equal boubt 'hyb))
       (return [iuvvv-fin0 asymbl hyemb])])
;cond
(cond [(and (equal aubt 'hyb) (equal boubt 'src))
       (return [iuvvv-fin0 hyemb symbbl])])
;cond
(cond [(and (equal aubt 'src) (equal boubt 'src))
       (return [iuvvv-fin1 asymbl hyemb])])
;cond
(cond [(and (equal aubt 'src) (equal boubt 'src))
       (return [iuvvv-fin0 hyemb symbbl])])
;cond
(cond [(and (equal aubt 'src) (equal boubt 'src))
       (return [iuvvv-fin2 asymbl hyemb])])
;cond
(cond [(and (equal aubt 'src) (equal boubt 'src))
       (return nil)])
;princ
;jdefun

;iuvvv-fin0 : find conditions between a src and a sum
(defun iuvvv-fin0 (ryemb symbbl)
  (prog [prop sprop rdefn sdefn nu-const]
    ; only if they share a common variable in their definitions,
    ; have opposite directions, and involve the same node
    (cond (nequal (caar rsymbol) (cadar symbbl))
          (return nil)))
;cond
(cond [(eq sprop [cadr (caddr symbbl)])
        (setq prop [cadf (caddr symbbl)]])
       (cond (lareval [test rprop 'Dir] (get sprop 'Dir))
             (return nil))]
;cond
(cons
  [setq sprop [cadf (caddr symbbl)]]
  [setq rdefn [get sprop 'xDelvar]]
  [setq sdefn [get sprop 'zDelvar]]]
  new-+entry)
;cond
(cond [null (member rdefn sdefn)]
      (do next-+entry)
;cond
  [setq entry (car rdefn)]
  [setq rdefn (cdr rdefn)]
  (cond [null (member entry sdefn)]
        (do next-+entry)
;cond
  [setq and sdefn share a common variable in their def
    (setq nu-const [len=cond ryemb symbbl])
    (princ "asserter: summary variable ") (princ sprop)
    (princ " and static variable ") (princ prop)
    (princ " invokes ") (princ entry) (terpri)
    (princ " new condition: ") (princ nu-const) (terpri)
    (return nu-const)
;princ
;jdefun

;iuvvv-fin0 : find conditions between a src and a elim
(defun iuvvv-fin2 (ryemb symbbl)
  (prog [prop sprop prop nu-const]
    (cond (nequal (caar rsymbol) (cadar symbbl))
          (return nil)))
;cond
(equal (caar hyemb) 'i)
(equal (caar hyemb) 'j)
(seq nu-const [len=cond hyemb symbbl])
(princ "asserter: ") [princ bprop]
(princ " is a hybrid static variable involving ")
(princ bprop) [terpri]
(princ " new condition: ") {princ nu-const} [terpri]
(princ nu-const)
(princ "asserter: ") [princ bprop]
(princ "asserter: ") [princ bprop]

```

*(princ " is a hybrid static variable involving ")  
*(princ bprop) [terpri]*  
*(princ " new condition: ") [princ nu-const] [terpri]*  
*(return nu-const)**

**b-|error|**

**jdefun**

*iuvvv-fin1 : find conditions between a src and a sum*  
*(defun iuvvv-fin1 (ryemb symbbl)*  
*(prog [prop sprop rdefn sdefn nu-const]*  
*; only if they share a common variable in their definitions,*  
*; have opposite directions, and involve the same node*  
*(cond (nequal (caar rsymbol) (cadar symbbl))*  
*(return nil))*

*(cond (eq sprop [cadr (caddr symbbl)])*  
*[setq prop [cadf (caddr symbbl)]]*  
*(cond (lareval [test rprop 'Dir] (get sprop 'Dir))*  
*(return nil))*

*(cond*  
*[setq sprop [cadf (caddr symbbl)]]*  
*[setq rdefn [get sprop 'xDelvar]]*  
*[setq sdefn [get sprop 'zDelvar]]*  
*new-+entry)*

*(cond*

*[null (member entry sdefn)]*

*(do next-+entry))*

*(cond*

*[ryemb and sdefn share a common variable in their def*

*(setq nu-const [len=cond ryemb symbbl])*

*(princ "asserter: summary variable ") (princ sprop)*

*(princ " and static variable ") (princ prop)*

*(princ " invokes ") (princ entry) (terpri)*

*(princ " new condition: ") (princ nu-const) (terpri)*

*(return nu-const)*

**bprop**

**jdefun**

*iuvvv-fin0 : find conditions between a src and a elim*  
*(defun iuvvv-fin0 (ryemb symbbl)*  
*(prog [prop sprop prop nu-const]*  
*(cond (nequal (caar rsymbol) (cadar symbbl))*  
*(return nil))*

*[eq eq prop [cadr (caddr symbbl)]]*  
*[eqq jprop [cadf (caddr symbbl)]]*  
*(cond [nequal (caar ryemb) (cadar symbbl))*  
*(return nil)]*

```

[return nil]
;cond
; :symbol is a limit for rymb but not a feasibility limit
; (which is satisfied by all feasible nodes)
(setq au-cond [gen-cond symb lymb])
(princ "answer: ") [print prop] [princ " is a limit of ")
(princ "prop") [terpri]
(princ " new conditions: ") [princ au-cond] [terpri]
(return nil)
)prog
)defun

*****+
; functions that implement the derivator
*****+
; derive : the derivator
; derive [var-list]
; (defun derive [var-list]
;   (princ "deriv blist aux-list prop slot evals [port sum-frame num-slots]")
;   (princ "Derivator") [terpri]
;   (cond [(or (Def-list Dbl-blst) [or And-cond])])
;   ; initial case -> read in basic variables
;   (setq iport [find var-blst])
;   (princ " derivator: reading set of initial variables") [terpri]
;   (setq num-frames [read iport])
;   num-frame
;   (cond [(notrop num-frames) (close iport)]
;         (t
;          (letrec ([num-frames (l- num-frames)])
;            (setq iprop (read iport))
;            (putprop iprop iprop 'Name)
;            (setq num-slots 0)
;            (setq slot
;                  (cond [(notrop num-slots) (dumpprop iprop)
;                         (go aux-frame)])
;                      (else (notrop [read iport]))
;                      (setq avalue [read iport])
;                      (putprop iprop avalue slot)
;                      (setq num-slots (l- num-slots))
;                      (go next-slot)
;                  )
;            )
;            (putprop
;              num-frame
;              (list ' (list 'get 'sp [list 'quote au-name])
;                    (list 'get QVar [list 'quote iprop])
;                    'Dcl) ; defining code
;              (putprop au-name (list au-name iprop) 'Defact) ; variables in defa
;              (putprop
;                au-name
;                (list ' (list 'get 'sp [list 'quote au-name])
;                      'Init) ; initialization code
;                (putprop au-name (list iprop 'xDefvar) 'Rel) ; related variables
;                (princ Dbl-blst [cons au-name Dbl-blst]) ; add to list of properties
;                (princ " derivator: new summary variable from ")
;                (princ iprop) [terpri]
;            )
;          )
;        )
;      )
;    )
;  )
;)
;
```

```

;setq blist [car (append Dbl-blst Dbl-blst)]
;next-aymb
; (cond [(null alist) (exit-auxprop Dbl-blst)
;        (return [close Dbl-blst])])
; (if (defn (car alist)) (return [close Dbl-blst])
; (cond* (defn (car alist) blist)
;       (cons alist (cdr alist)))
;       (setq alist (cdr blist))
;       (go aux-aymb)
;     )
; )
;cond
; self-def: derive variable from one existing variable
; (i.e. summary variable / view variable)
; (defun self-def (iprop)
;   (prog (typo auxb)
;     (cond
;       ((member (get iprop 'Sub) '(rel etc))
;        (make-sum iprop)
;        (cond [(and (equal Pbinds 'value) Private)
;               (make-view iprop)])
;        )
;       )
;     )
;   )
;   )
; )
;defun

; makeum : create a new summary variable
; (defun make-sum (iprop)
;   (prog (auxname)
;     (setq aux-name (consit 'R iprop))
;     ; all in its properties
;     (putprop aux-name 'aux 'Type) ; rpb evr [entity variable]
;     (putprop aux-name 'eum 'Sub) ; subtyp eum [summary variable]
;     (putprop aux-name 'loc 'xDir) ; direction loc [future decisions]
;     (putprop
;       iprop
;       iprop
;       auxname
;       (list ' (list 'get 'sp [list 'quote au-name])
;             (list 'get QVar [list 'quote iprop])
;             'Dcl) ; defining code
;       (putprop au-name (list au-name iprop) 'Defact) ; variables in defa
;       (putprop
;         au-name
;         (list ' (list 'get 'sp [list 'quote au-name])
;               'Init) ; initialization code
;         (putprop au-name (list iprop 'xDefvar) 'Rel) ; related variables
;         (princ Dbl-blst [cons au-name Dbl-blst]) ; add to list of properties
;         (princ " derivator: new summary variable from ")
;         (princ iprop) [terpri]
;       )
;     )
;   )
; )
;
```

```

; (dump-prop nu-name)
; (return nu-name)
; lprop
; )defun

; makeview : create a new view variable

(defun make-view (lprop)
  (prog (nu-name)
    (setq nu-name (cons 'V lprop))
    ; fill in its properties
    (putprop nu-name nu-name 'xName)
    {putprop nu-name 'var xType} ; type var (entity variable)
    {putprop nu-name 'view xSubs} ; subtype view (view variable)
    {putprop nu-name 'Var xDef} ; direction for [future decision]
    {putprop

nu-name
      ; no implied variables
      (list 'init (list 'V + (list 'get 'ql "xVar"))
            (list 'get 'qp (list 'quote sprop)))
      'xDef) ; defining code

(putprop nu-name (list lprop) 'xDef) ; only variables in formula
      ; calc-nurop : compute values of new properties for the search nodes
      ; calc-nurop : (defun calc-nurop (plists)
      ;   (let* ((prop initlist inst node-list strans tchr qt ql qt formula
      ;          normals lev-index root node qp)
      ;         (princ "Computing values of new symbols for problem instances")
      ;         (tchr)
      ;         (next-prop
      ;           (cond ((null plists) (return t)))
      ;                 (setq prop (car plists))
      ;                 (setq plists (cdr plists))
      ;                 (setq formula (get prop 'xDef))
      ;                 (setq formula (get prop 'xInit))
      ;                 (setq initlist init)
      ;                 (initlist)
      ;                 (cond ((null initlist) (go next-prop))
      ;                       (setq inst (car initlist))
      ;                       (setq initlist (cdr initlist))
      ;                       (putprop inst (cons prop (get inst 'propinst)) 'propinst)
      ;                       (setq tchr (get inst 'levlist))
      ;                       (cond
      ;                         ((or (nequal (get inst 'type) 'value)
      ;                               (nequal (get prop xType) 'var))
      ;                           (go process-var))
      ;                         (cond
      ;                           ((eq q (car tchr))
      ;                             (putprop q (eval formula) prop)
      ;                             next-q)
      ;                           (cond ((null tchr) (go next-prop))
      ;                                 (setq qp q)
      ;                                 (setq tchr (cdr tchr))
      ;                                 (putprop q (eval formula) prop)
      ;                                 (go next-q))
      ;                         (cond
      ;                           ((and (equal (get symb 'xType) 'var) (equal Phdbs 'value))
      ;                            (setq nu-list (append nu-list
      ;                               (list 'get 'qp (list 'quote symb))))))
      ;                           (setq get ' (list 'quote symb)))))

      ; symb is a variable
      ; assign variables set of each node with new variable
      ; process-var
      ; (setq tchr (eval (get inst 'lev)))
      ; (setq lev-index (cdr (rest tchr)))
```

```

; spatially process the root of the search tree
(defun sp-prop-root (eval-formula) prop)
  (cond ((and (equal (sprop 'eval) (equal bprop 'present))
              (not (make-sprop bprop)))
         (go next-bprop))
        ;cond
        ((and (equal (sprop 'eval) (equal bprop 'eval))
              (not (make-sprop bprop)))
         (go next-bprop))
        ;cond
        ((and (equal (sprop 'eval) (equal bprop 'eval))
              (not (make-sprop bprop)))
         (go next-bprop))
        ;process other levels in search tree
        next-level
        (cond ((null lev-index) (go next-level)))
        (setq node-list (cdr lev-index))
        (setq lev-index (cdr lev-index))
        (setq q (car lev-index))
        (setq thv (cdr lev-index))
        ;process nodes on this level
        next-node
        (cond ((null node-list) (go next-level)))
        (setq node (car node-list))
        (setq qd node)
        (setq qp (get node 'P))
        (putprop node (eval formula) prop)
        (go next-node)
        )prop
      )defun
  )

; cross-def : define new property by combining two properties
(defun cross-def (sprop bprop)
  (prog (bprop-type bprop)
    (setq at-type (get sprop 'xType))
    next-bprop
    (cond ((null bprop) (return t)))
    (setq bprop (car bprop))
    (setq bname (cdr bprop))
    (setq btype (set bprop 'xType))
    (cond ((and (equal at-type 'present) (equal btype 'eval))
           (ip-make-bprop sprop))
          (cond ((and (equal at-type 'eval) (equal btype 'eval))
                 (t-merge-prop sprop)
                 (putprop nu-name nil 'DifVar)
                 (putprop nu-name nil 'Init) ; a hack that works
                 ; no &DifFlag - don't know how to do this also
                 ; no &ConfFlag - don't know how to do this also
                 (putprop nu-name (list sprop 'xDifVar sprop 'xDifVar) 'Rel)
                 (dumpprop nu-name)
                 (go next-bprop)))
        
```

```

        )
      )Def
    )
  )
  (cond ((and (equal at-type 'eval) (equal btype 'eval))
         (t-merge-prop sprop)
         (putprop nu-name (list sprop 'xDifVar sprop) 'xDifVar)
         (putprop nu-name nil 'Init) ; a hack that works
         ; no &DifFlag - don't know how to do this also
         ; no &ConfFlag - don't know how to do this also
         (putprop nu-name (list sprop 'xDifVar sprop 'xDifVar) 'Rel)
         (dumpprop nu-name)
         (go next-bprop)))
  )
)
```

```

[setq Dx-list (cons nu-name Dx-list)]
  [return nu-name]
]prog
)defun

; find-defvar : returns the list of variables involved in the definition
; of the new variable

(defun find-defvar (sprop bprop)
  [prog
    (prog (list dprop plist)
      [setq dist (list sprop bprop)]
      ; now add those variable present in *Defvar of sprop and bprop
      ; if not already present in dist
      [setq plist (get sprop '*Defvar)]
      nextp0
      [cond [(null plist) (do done)]
            [setq dprop (car plist)]
            [setq plist (cdr plist)]
            [cond [(member dprop dist) t]
                  [setq dist (cons dprop dist)]]
            [cond
              [do nextp]
              done
              [setq plist (get bprop '*Defvar)]
              nextp1
              [cond [(null plist) (return dist)]
                    [setq dprop (car plist)]
                    [setq plist (cdr plist)]
                    [cond [(member dprop dist) t]
                          [setq dist (cons dprop dist)]]
                    [cond
                      [do nextp]
                      done
                      [prog
                        ; setmake : make new variables, if any, between a var and a var
                        ; setmake : make new variables, if any, between a var and a var
                        (defun re-make (sprop sprop)
                          [prog
                            (prog (e-subt e-subt)
                              [setq e-subt (get sprop '*Subt)]
                              [setq e-subt (set sprop '*Subt)]
                              [cond (and (equal e-subt 'svar) (equal e-subt 'evar))
                                    [return (s+ Info sprop sprop)]]
                                becond
                                [cond (and (equal e-subt 'senc) (equal e-subt 'evar))
                                      [return (s+ Info sprop sprop '*ConsFrag)]]
                                    becond
                                    [return nil])
                                ]prog
                        )defun
                      ]
                    ]
                  ]
                ]
              )
            )
          ]
        )
      )
    )
  )
]defun

; set-no : create new variable between the svar and the evar

(defun set-no (svar-prop evar-prop)
  [prog (list elist svar-prop nu-name)
    [setq elist (get svar-prop '*Defvar)]
    [setq elist (get evar-prop '*Defvar)]
    next-savary
    [cond (null elist) (return nil)]
    [setq senvry (car elist)]
    [setq elist (cdr elist)]
    [cond [(null (member senvry elist)) (go next-savary)]
          [setq nu-name (cons svar-prop 'evar-prop)]
          [putprop nu-name nu-name 'Name]
          [putprop nu-name svar 'xType]
          [putprop nu-name 'shy #'Subs)
          [cond
            [member evar-prop rel-list] t
            [setq rel-list (cons evar-prop rel-list)]
            [mapc
              [lambda (cond-prop)
                (cond
                  [rel-list nil]
                  [t
                    [setq rel-list (cons evar-prop rel-list)]]
                  [setq rel-list (cons evar-prop rel-list)])
                ]
              )
            ]
          )
        )
      )
    )
  )
]defun

; find-rel : copy related variables in *Defvar, *OpFrag, and *ConsFrag
; into xRel and annotate with source

(defun find-rel (*nu-prop)
  [prog (rel-list)
    [setq rel-list nil]
    [mapc
      [lambda (cond-prop)
        (cond
          [rel-list]
          [[member cond-prop rel-list] t
            [setq rel-list (cons cond-prop rel-list)]
            [cond
              [rel-list nil]
              [t
                [setq rel-list (cons cond-prop rel-list)]]
                [setq rel-list (cons cond-prop rel-list)])
              [setq rel-list (cons cond-prop rel-list)]]
            ]
          )
        )
      )
    )
  )
]defun

; derived vars do not have OptFrag and ConsFrag right now
```

```

)defun

; re-make

(defun re-make (sprop sprop)
  [prog
    (prog (e-subt e-subt)
      [setq e-subt (get sprop '*Subt)]
      [setq e-subt (set sprop '*Subt)]
      [cond (and (equal e-subt 'svar) (equal e-subt 'evar))
            [return (s+ Info sprop sprop)]]
          becond
          [cond (and (equal e-subt 'senc) (equal e-subt 'evar))
                [return (s+ Info sprop sprop '*ConsFrag)]]
              becond
              [return nil])
            ]prog
    )defun

; set-no : create new variable between the svar and the evar

(defun set-no (svar-prop evar-prop)
  [prog (list elist svar-prop nu-name)
    [setq elist (get svar-prop '*Defvar)]
    [setq elist (get evar-prop '*Defvar)]
    next-savary
    [cond (null elist) (return nil)]
    [setq senvry (car elist)]
    [setq elist (cdr elist)]
    [cond [(null (member senvry elist)) (go next-savary)]
          [setq nu-name (cons svar-prop 'evar-prop)]
          [putprop nu-name nu-name 'Name]
          [putprop nu-name svar 'xType]
          [putprop nu-name 'shy #'Subs)
          [cond
            [member evar-prop rel-list] t
            [setq rel-list (cons evar-prop rel-list)]
            [mapc
              [lambda (cond-prop)
                (cond
                  [rel-list nil]
                  [t
                    [setq rel-list (cons evar-prop rel-list)]]
                    [setq rel-list (cons evar-prop rel-list)])
                  [setq rel-list (cons evar-prop rel-list)]]
                )
              )
            ]
          )
        )
      )
    )
  )
]defun

; find-rel : copy related variables in *Defvar, *OpFrag, and *ConsFrag
; into xRel and annotate with source

(defun find-rel (*nu-prop)
  [prog (rel-list)
    [setq rel-list nil]
    [mapc
      [lambda (cond-prop)
        (cond
          [rel-list]
          [[member cond-prop rel-list] t
            [setq rel-list (cons cond-prop rel-list)]
            [cond
              [rel-list nil]
              [t
                [setq rel-list (cons cond-prop rel-list)]]
                  [setq rel-list (cons cond-prop rel-list)])
              [setq rel-list (cons cond-prop rel-list)]]
            ]
          )
        )
      )
    )
  )
]defun

; derived vars do not have OptFrag and ConsFrag right now
```

```

;dump-prop au-name]
{return au-name}
jprog
jdefn
; create : make new variable, if any, between two vars
{
  defun ex-make (pprop bprop)
  {
    pprop (symbol symbol au-prop au-var)
    (setq auvar (get approp 'Subst))
    (setq bumb (get bprop 'Symbol))
    (cond
      {and (equal auvar 'nil) (equal bumb 'src)}
      {return (ee-fn0 approp bprop 'OptFlag)})
    )
  }
  cond
  {and (equal auvar 'src) (equal bumb 'src)}
  {return (ee-fn0 approp bprop 'OptFlag)})
  )
}

(cond
  {and (equal auvar 'src) (equal bumb 'src)})
  {return (ee-fn1 approp bprop)})
)
;cond
)
cond
(cond
  {and (equal auvar 'src) (equal bumb 'src)})
  {return (ee-fn1 approp bprop)})
)
;cond
)
cond
(cond
  {and (equal auvar 'src) (equal bumb 'src)})
  {return (ee-fn0 approp bprop 'ConfFlag)})
)
;cond
)
cond
(cond
  {and (equal auvar 'src) (equal bumb 'src)})
  {return (ee-fn0 approp bprop 'ConfFlag)})
)
;cond
)
cond
(cond
  {and (equal auvar 'src) (equal bumb 'src)})
  {return (ee-fn0 approp bprop 'ConfFlag)})
)
;cond
)
cond
(cond
  {and (equal auvar 'src) (equal bumb 'src)})
  {return (ee-fn0 approp bprop 'ConfFlag)})
)
;cond
)
cond
(cond
  {and (equal auvar 'src) (equal bumb 'src)})
  {return (ee-fn0 approp bprop 'ConfFlag)})
)
;cond
)
cond
(cond
  {and (equal auvar 'src) (equal bumb 'src)})
  {return (ee-fn0 approp bprop 'ConfFlag)})
)
;cond
)
cond
(cond
  {and (equal auvar 'src) (equal bumb 'src)})
  {return (ee-fn0 approp bprop 'ConfFlag)})
)
;cond
)
cond
}
; ee-fn0

{defun ee-fn0 (pprop bprop abut)
  {
    pprop (list alist comb)
    (setq alist (get approp abut))
    next-abut
    (cond [(null alist) (return nil)])
    {setq alist (cons alist))
    (setq alist (cdr alist))
    (cond [(null (member bprop alist)) (go next-abut)]
      {setq comb (car alist))
      (setq comb (car alist))
      {cond [(null member bprop alist)] (go next-abut))
      {setq comb (car alist))
      (setq comb (car alist))
      (setq comb (car alist))
      (return (cons alist bprop)))
    )
  }
}
```

```

[putprop nu-name nu-name 'xName]
[putprop nu-name 'var 'xType]
[cond [equal (get 'aprop 'xSub) (get 'bprop 'xSub)]
      (putprop nu-name [get 'aprop 'xSub] 'xSub]]
      [putprop nu-name 'chb 'xSub]]
;cond

[putprop nu-name
  [list 'coml [list 'get QVar [list 'quote bprop]]]
  'Def]
[putprop nu-name [list 'get QVar [list 'quote bprop]]]
; don't know how to derive ConfFlag and OptFlag for these
[putprop nu-name [list 'aprop 'xDefvar bprop 'xDefvar] 'xRef]; 2 basic vars
[setq Da-list [cons nu-name Da-list]]
[dump-prop nu-name]
;return nu-name]

jprog
:jdefun
:jdefun
:jdefun
:jdefun

;defun ex-fn [ret-prop rec-prop]
[prog (nu-name)
; hard-coded to produce the ratio
[setq nu-name (coerce rec-prop / ret-prop)]
[putprop nu-name 'xName]
[putprop nu-name 'var 'xType]
[putprop nu-name 'chb 'xSub]
[putprop nu-name
  [list 'quotient [list 'get QVar [list 'quote rec-prop]]
  'Def]
[putprop nu-name [list 'ret-prop rec-prop] 'xDefvar]; 2 basic vars
[putprop nu-name nil 'xRef]
; no initialisation
; no OptFlag nor ConfFlag
[putprop nu-name [list 'ret-prop 'xDefvar rec-prop 'xDefvar] 'xRef]; 2 basic vars
[setq Da-list [cons nu-name Da-list]]
[dump-prop nu-name]
;return nu-name

jprog
:jdefun
;jdefun

;dump-prop : dump out a property frame
(defun dump-prop [prop]
[princ "Type "][princ [get 'aprop 'xType]])(lab 40)
(princ "Sib ")(princ [get 'aprop 'xSub])(setq
[lab 40] (princ "Def ")(princ [get 'aprop 'xDef])(setq
[lab 45] (princ "Init ")(princ [get 'aprop 'xInit])(setq

[lab 40] (princ "QVar ")(princ [get 'aprop 'xQVar])(setq
[lab 45] (princ "R ")(princ [get 'aprop 'xR])(setq
[lab 50] (princ "xDefvar ")(princ [get 'aprop 'xDefvar])(lab 50)
(princ "xRef ")(princ [get 'aprop 'xRef])(lab 55)
[lab 50] (princ "xDef ")(princ [get 'aprop 'xDef])(lab 55)

;initialization routines
;***** initialization routines *****

; init : carry out initialisation
[defun init [desc-fld]
[princ "Initialization"])(terpri)
; set size and elements of <frame [x-crule]
[setq c-size 15]
[setq cName 0]
[setq cType 1]
[setq cLhs 2]
[setq cRhs 3]
[setq cParent 4]
[setq cChild 5]
[setq cParent 6]
[setq cChildren 7]
[setq cDelta 8]
[setq cDepth 9]
[setq cEdeg 10]
[setq cHilite 11]
[setq cNameps 12]
[setq cNumang 13]
[setq cNumatt 14]
[setq cDv 15]
; set elements of a tframe [vec tframe]
[setq ctName 0]
[setq cNodeset 1]
;init-form desc-fld
; kludge to generate root of rule graphs and to ensure that the
; c-rules on the first level only invoke the return variable
; IMPORTANT : return state variable must be named R1
;defun

[setq CondList nil]
[setq oneCondList '(> (get x 'R) (get y 'R))
 (< (get x 'R) (get y 'R))
 (equal (get x 'R) (get y 'R))]

[setq cDv (new-vector 10)]
[setq cD0 ctName 'cD0]
[setq cD0 ctType 'cD0]
[setq cD0 cType 'cD0]

```

```

[set cdb <Cdb> {{Sibling * y}}]
[set cdb <Grand nil>]
[set cdb <Parent 'a0>]
[set cdb <Ancest 'a0>]
[set cdb <Children nil>]
[set cdb <Delta 'a0>]
[set cdb <Offset '(+)>]
[set cdb <Negat '(-)>]
[set cdb <Nullst '0>]
[set cdb <Chumpst 0>]
[set cdb <Numnst 0>]
[set cdb <Numau 0>]
[set cdb <New-vector 10>]
[set cdb <Name 'obj>]
[set cdb <Type 'beta>]
[set cdb <Class <Equal (set x 'y) (set y 'L)>>]
[set cdb <Cond nil>]
[set cdb <Parallel 'ao>]
[set cdb <Anst 'bo>]
[set cdb <Parent nil>]
[set cdb <Children nil>]
[set cdb <Delta 'ao>]
[set cdb <Offset '(+)>]
[set cdb <Negat '(-)>]
[set cdb <Nullst '0>]
[set cdb <Numau 0>]
[set cdb <Numang 0>]
[set cdb <Numau 0>]
[set cdb <Numang 0>]

; initialize asserted global variables

; set AxiomList nil]
; set AxiomList (new-vector 1 23)]
; (set AxiomList 0 [(c0)])
; (intervm [(c0 0)])
; (set CdbAxiomList nil)
; (make 0 'as_axiomList on!CdbAxiomList 2)
; (set blvAxiomList (new-vector 1 23))
; (set blvAxiomList 0 [(c0)])
; (intervm [(c0 0)])
; (make 0 'cb blvAxiomList on!ConditList 1)

; (defun int-form (descfile)
;   (prog (port numinst instname numinsts prop value numitems
;           hist insttree instframe tree)
;     (setq hist nil)
;     (setq insttree nil)
;     (setq instframe nil)
;     (setq instinst nil)
;     (cond ((eqlp numinst) (coi ingest-in-st))
;           ((eq numinst (1- numinst)))
;           (setq instances (read port))
;           (setq instances nil))

; (defun QVar 'q)
;   (set SymTabList nil)
; )
; ; edump0 : print summary of a crule
; ; (defun edump0 (crname)
;   (prog (crname)
;     (setq crname (eval crname))
;     (princ (ref crname cName))
;     (princ (ref crname clName))
;     (princ "LHS : ") (princ (ref crname clLhs)) (terpri)
;     (princ "RHS : ") (princ (ref crname cRule))
;     (princ "Parallel : ") (princ (ref crname cParallel))
;     (terpri)
;     (princ "Anst : ") (princ (ref crname cAnst))
;     (princ "Parents : ") (princ (ref crname cParents))
;     (princ "Children : ") (princ (ref crname cChildren))
;     (terpri)
;     (princ "Numang : ") (princ (ref crname cNumang))
;     (princ "Numau : ") (princ (ref crname cNumau))
;     (terpri)
;   )
; )
; ; edump1 : print a different summary of a crule
; ; (defun edump1 (crname)
;   (prog (crname)
;     (princ "LHS : ") (princ (ref crname clLhs)) (terpri)
;     (princ "RHS : ") (princ (ref crname clRhs)) (terpri)
;     (princ "Parallel : ") (princ (ref crname cParallel))
;     (terpri)
;     (princ "Anst : ") (princ (ref crname cAnst))
;     (princ "Parents : ") (princ (ref crname cParents))
;     (princ "Children : ") (princ (ref crname cChildren))
;     (terpri)
;   )
; )
; ; init-form : creates the frames for the problem instances
; ; assumes data file of form (illustrated by qf.lisp script):
; ;   1 : # of problem instances in data file
; ;   n : name of problem instance
; ;   t : # of properties of problem instance
; ;   i : (1 4 3 2 5) : property #1 and its value
; ;   w : (3 4 1 4) : #2
; ;   K : 8 : #3
; ;   numinsts : 5 : #4 - #initials in problem
; ;   numinsts : 5 : #5 - #initials in problem
; ;   proplist : (#1 #2 #3) : property list
; ;   type-value : #6 - description type
; ;   tree : 0 : #7 - name of the holding search tree
; ; (defun int-form (descfile)
;   (prog (port numinst instname numinsts prop value numitems
;           hist insttree instframe tree)
;     (setq hist nil)
;     (setq insttree nil)
;     (setq instframe nil)
;     (setq instinst nil)
;     (cond ((eqlp numinst) (coi ingest-in-st))
;           ((eq numinst (1- numinst)))
;           (setq instances (read port))
;           (setq instances nil))
;
```

```

(princ "reading instance ") (princ instname) (terpri)
(setq inst-lst (cons instname inst-lst))
(setq num-items (read (open)))
(puprop instname instname 'name)
;arity-ids
(cond ((terop num-items) (go cleanup-inst)))
(terop iprop (read iprop))
(setq ivalue (read iprop))
(puprop instname ivalue iprop)
(setq num-items (1- num-items))
(go next-inst)
cleanup-inst
(setq num-items (get instance 'num-items))
(setq last nil)
next-item
;cond #((lambda (num-items) (puprop instname list 'ivlist)
  ;(go next-inst))
  ;(cond
  ;(eq last (cons (concat instname ' num-items) list)))
  ;(setq num-items (1- num-items))
  ;(go next-item)
  ;)
  ; transform these frames into working data structures
  ; ingest-inst
  ;(setq instlist (cdr inst-lst))
  ;next-inst
  ;(cond ((null instlist) (return v)))
  ;(setq inst (car instlist))
  ;(setq instlist (cdr instlist))
  ; first, create tree frame and read in the search tree
  ;(setq tree (get inst 'tree))
  ;(setq tree (new-vector 4))
  ;(setq frame (eval tree))
  ;(setq frame (frame frame tree))
  (princ "reading search tree ") (princ tree) (terpri)
  ;3 inst (frame)
  ; now process the entity variables
  ;cond
  ;(equal (get inst 'type) 'value)
  ;(value-ingest inst)
  ;(equal (get inst 'type) 'iden)
  ;(iden-ingest inst)
  ;(terop)
  ;cond
  ;(go next-inst)
  ;puprop
  ;defun
)
```

```

; value-ingest : ingest the information in the instance frame into
; the working data structures
;defun value-ingest (inst)
;prog (terlist lvarname proplist prop)
  (terop (terlist lvarname proplist prop))
  (terop lvarlist (get inst 'terlist))
  (terop lvarname (car lvarlist))
  (terop lvarlist (cdr lvarlist))
  (puprop lvarname lvarname 'name)
  (puprop lvarname 0 'level)
  (mapc
    'lambda (x)
    (puprop lvarname nil x)
    lvarlist
  )
  (let (inst proplist)
    (lambda
      (princ
        (terop lvarnum 0)
        inst-lvarname
        (terop lvarname)
        (cond ((null lvarlist) (return v))
          (terop lvarlist (car lvarlist))
          (terop lvarname (car lvarlist))
          (terop lvarlist (cdr lvarlist))
          (terop proplist (get inst 'proplist))
          (puprop lvarname lvarname name)
          (puprop lvarname (1+ lvarnum) 'level)
        )
        next-prop
      )
      (cond
        ((null proplist)
          (terop lvarnum (1+ lvarnum))
          (go next-lvarname))
        (cond
          ((terop (terprop (car proplist)))
            (terop proplist (cdr proplist))
            (puprop lvarname (1+ lvarnum) (get inst prop)) prop)
          (go next-prop)
        )
        (princ
          (terop
            (defun
              ; iden-ingest : digest the entity variables in an identity description
              ; ingest the entity variables in each node
              (defun iden-ingest (inst)
                (terop (terframe)
                  (terop frame (eval (get inst 'tree)))
                  (mapc
                    'lambda (property)
                    (princ (list)
                      (terop (terframe)
                        (terop frame (eval (get inst 'tree)))
                        (terop next-set)
                        (cond ((null inst) (return v))
                          (terop
                            (lambda (node)

```

```

[putprop node [set [get node 'Op] property] property]
;lambda
;  (cdr llist)
;lambda
;mapc
;  (seq llist (cdr llist))
;  (go next-set)

;lprog
;lambda
;  (set init 'propinit)
;lambda
;lprog
;lprog
;lambda
;  (cond
;    ; num : utility function to find sum of numbers in a list.
;    (defun num (llist)
;      (prog (sum entry)
;        (setq sum 0)
;        next-entry
;        (cond ((null llst) (return sum)))
;          (setq entry (car llst))
;          (setq alias (cdr llst))
;          (setq sum (+ sum entry))
;          (so next-entry)
;        )lprog
;      )ldefun
;      ;(2 : read in search tree
;
;      (defun r3 (first lframe)
;        (prog (iport num name nodestk level llst q! pfirst)
;          (setq lev 0)
;          (setq name (ref lframe 'Name))
;          (setq llst (ref lframe 'Var llst))
;          (setq iport (info name))
;          (setq pfirst (read iport)) ; set of variables for each node
;          (setq nodestk nil)
;          (prog ()
;            (setq num (read iport)) ; # of nodes in current level
;            (cond
;              (null num) ; all
;              (setq level (list lev)) ; all col
;              (setq llst (list lev))
;              (setq q! (car llst))
;              (setq llst (cdr llst))
;            )lprog
;            (prog () ; read node
;              (setq nnode
;                (cond
;                  (equal num 0) ; all nodes on level k read
;                  (list q! node) ; read iport)
;                  (putprop node node 'N)
;                )lmapc
;
```

```

;lambda (x)
;  (putprop node [read iport] x)
;lambda
;plus
;lprog
;  (putprop node init 'q)
;  (putprop node q! 'q)
;  (setq num (1+ num))
;  (setq lvalst (append lvalst node))
;  (go extnode) ; go to next node

;lprog
;lambda
;  (cond
;    ; all nodes on level lev read in
;    (setq nodestk (append nodestk (lev)))
;    (setq lev (1+ lev))
;    (go nextlev) ; go to next level
;  )lprog
;lambda
;  (cond
;    ; (cons iport) ; all done
;    (setq lframe (NodeList (append (list lev) nodestk)))
;  )ldefun
;lambda
;  (cond
;    ;(2 : read in search tree
;
;    (defun Sibling (x y)
;      (prog ()
;        (cond [(equal x (set Y 'Sj)) (return t)]
;          (cond [(equal x (set Y 'Sj)) (return nil)])
;        )
;      )ldefun
;    )lprog
;
```

### Appendix B : A Trace for the 0/1-Knapsack Problem

This appendix describes one data set supplied to the searcher and the resulting trace produced by the searcher. The data supplied to the searcher is distributed over four files: knap, des, t0, and t1. The file knap [Figure A1] describes the symbols in the value-description of the 0/1-knapsack problem. The first entry in this file is the number of symbols. Each symbol has ten entries. The first entry for each symbol is its name. Each of the remaining nine entries describe a slot name and its corresponding slot filter.

The file des (Figure A2) describes the 0/1-knapsack problem instances. The first entry in this file is the number of problem instances. Each problem instance is described by a variable number of entries. Two entries must be present for each problem instance: the name of the problem instance and the number of properties that it possesses. Each remaining entry describes a slot name and its slot filter.

The file t0 (Figure A3) describes the search tree for problem instance i0. The first entry in t0 is a list which serves as a template for the property values of each node. The nodes are described by levels. Each set of entries for the set of equivalent nodes consists of a count of the number of nodes on that level and the appropriate entries describing the individual nodes. Similarly, Figure A4 shows t1, the search tree for problem instance i1.

The searcher is then invoked with arguments giving the names of the files containing the symbol descriptions and the file. The actual trace produced by the searcher is as follows:

```

->[searcher 'knap 'des]
Initialization
reading instance i0
reading search tree i0
reading search tree i1
Answer
Derivator
derivator: reading set of initial variables
w
  xType ev*           xSub err*
  xDef (get ql 'w)    xInit nil
  xOpFrag nil          xConfFrag {{K - read}}

```

Figure A1 Listing of file knap

	L	R	W	B	SI	Pk	Ch	Op
1	0	0	0	10	0d	0d	(0d1 0d2)	0d1
2	0	0	0	10	0d2	0d2	(0d3 0d4)	0d2
3	1	0	0	10	0d3	0d3	(0d5 0d6)	0d3
4	0	1	0	10	0d4	0d4	(0d7 0d8)	0d4
5	2	0	0	10	0d5	0d5	(0d9 0d10)	0d5
6	0	2	0	10	0d6	0d6	(0d11 0d12)	0d6
7	3	0	0	10	0d7	0d7	(0d13 0d14)	0d7
8	0	3	0	10	0d8	0d8	(0d15 0d16)	0d8
9	4	0	0	10	0d9	0d9	(0d17 0d18)	0d9
10	0	4	0	10	0d10	0d10	(0d19 0d20)	0d10
11	5	0	0	10	0d11	0d11	(0d21 0d22)	0d11
12	0	5	0	10	0d12	0d12	(0d23 0d24)	0d12
13	6	0	0	10	0d13	0d13	(0d25 0d26)	0d13
14	0	6	0	10	0d14	0d14	(0d27 0d28)	0d14
15	7	0	0	10	0d15	0d15	(0d29 0d30)	0d15
16	0	7	0	10	0d16	0d16	(0d31 0d32)	0d16
17	8	0	0	10	0d17	0d17	(0d33 0d34)	0d17
18	0	8	0	10	0d18	0d18	(0d35 0d36)	0d18
19	9	0	0	10	0d19	0d19	(0d37 0d38)	0d19
20	0	9	0	10	0d20	0d20	(0d39 0d40)	0d20
21	10	0	0	10	0d21	0d21	(0d41 0d42)	0d21
22	0	10	0	10	0d22	0d22	(0d43 0d44)	0d22
23	11	0	0	10	0d23	0d23	(0d45 0d46)	0d23
24	0	11	0	10	0d24	0d24	(0d47 0d48)	0d24
25	12	0	0	10	0d25	0d25	(0d49 0d50)	0d25
26	0	12	0	10	0d26	0d26	(0d51 0d52)	0d26
27	13	0	0	10	0d27	0d27	(0d53 0d54)	0d27
28	0	13	0	10	0d28	0d28	(0d55 0d56)	0d28
29	14	0	0	10	0d29	0d29	(0d57 0d58)	0d29
30	0	14	0	10	0d30	0d30	(0d59 0d60)	0d30
31	15	0	0	10	0d31	0d31	(0d61 0d62)	0d31
32	0	15	0	10	0d32	0d32	(0d63 0d64)	0d32
33	16	0	0	10	0d33	0d33	(0d65 0d66)	0d33
34	0	16	0	10	0d34	0d34	(0d67 0d68)	0d34
35	17	0	0	10	0d35	0d35	(0d69 0d70)	0d35
36	0	17	0	10	0d36	0d36	(0d71 0d72)	0d36
37	18	0	0	10	0d37	0d37	(0d73 0d74)	0d37
38	0	18	0	10	0d38	0d38	(0d75 0d76)	0d38
39	19	0	0	10	0d39	0d39	(0d77 0d78)	0d39
40	0	19	0	10	0d40	0d40	(0d79 0d80)	0d40

**Figure A2** Listing of file des.

```

    xDefVar ('w')          xRef ('w' xDefVar K xConstFrag)
    xSubst errt           xInit nil
    xType errt            xConfFrag nil
    xDef (set q1 't')      xRel (xDefVar w xDefVar K xConstFrag)
    xOpFrag nil           xRel (xDefVar w xDefVar K xConstFrag)
    xDefVar ('w')          xSubst errt
    xType errt            xInit nil
    xDef (set q0 'W')      xInit 0
    xOpFrag nil           xConfFrag ((K , feal))
    xDefVar (W w)          xRel (W xDefVar w xDefVar K xConstFrag)
    xType errt            xSubst errt
    xDef (set q1 'W')      xInit 0
    xOpFrag nil           xConfFrag 0
    xDefVar (R t)          xRel (R xDefVar w xDefVar)
    xType errt            xInit 0
    xDef (set q1 K)        xSubst nil
    xOpFrag nil           xConfFrag ((W - (set x 'W)) (w - (set x 'W)))
    xDefVar (K)            xRel (W xConstFrag w xConstFrag)
    ansver: (set x 'R) and (set y 'R) have identical properties
    new conditions: ((> (set x 'R) (set y 'R)) < (set x 'R) (set y 'R)))
    ansver: (set x 'W) and (set y 'W) have identical properties
    new conditions: ((> (set x 'W) (set y 'W)) < (set x 'W) (set y 'W)))
Category:
c1 LHS : ((Slicing x y) (> (set x 'R) (set y 'R)))
    Parent: (c40) Children: (c40 c41 c42)
    Parallel: ch1 Andi: ch2 Parent: (c40) Children: (c41 c42 c43)
c2 LHS : ((Slicing x y) < (set x 'R) (set y 'R))
    Parallel: ch2 Andi: ch2 Parent: (c40) Children: (c42 c43 c44)
c3 LHS : ((equal (get x 'I1) (equal (get x 'I1) (get y 'I1))) < (set x 'W) (set y 'W)))
    Parallel: ch3 Andi: ch3 Parent: (c40) Children: (c40 c41 c42)
c4 LHS : ((equal (get x 'I1) (get y 'I1)) > (set x 'R) (set y 'R))
    Parallel: ch4 Andi: ch4 Parent: (c40) Children: (c44 c45 c46)
c5 LHS : ((Slicing x y) (> (set x 'R) (set y 'R)) < (set x 'W) (set y 'W)))
    Parallel: ch5 Andi: ch5 Parent: (c40) Children: (c47 c48 c49)
c6 LHS : ((Slicing x y) (parallel (set x 'R) (set y 'R)) < (set x 'W) (set y 'W)))
    Parallel: ch6 Andi: ch6 Parent: (c40) Children: (c49 c50 c51)
c7 LHS : ((Slicing x y) (< (set x 'R) (set y 'R)) (> (set x 'W) (set y 'W)))
    Parallel: ch7 Andi: ch7 Parent: (c40) Children: (c51 c52 c53)
c8 LHS : ((Slicing x y) (< (set x 'R) (set y 'R)) < (set x 'W) (set y 'W)))
    Parallel: ch8 Andi: ch8 Parent: (c40) Children: (c52 c53 c54)
c9 LHS : ((Slicing x y) (< (set x 'R) (set y 'R)) (equal (set x 'W) (set y 'W)))
    Parallel: ch9 Andi: ch9 Parent: (c40) Children: (c53 c54 c55)
c10 LHS : ((Slicing x y) (equal (set x 'R) (set y 'R)) (> (set x 'W) (set y 'W)))
    Parallel: ch10 Andi: ch10 Parent: (c40) Children: (c54 c55 c56)
c11 LHS : ((Slicing x y) (equal (set x 'R) (set y 'R)) (< (set x 'W) (set y 'W)))
    Parallel: ch11 Andi: ch10 Parent: (c40) Children: (c55 c56 c57)
c12 LHS : ((Slicing x y) (equal (set x 'R) (set y 'R)) (equal (set x 'W) (set y 'W)))
    Parallel: ch12 Andi: ch12 Parent: (c40) Children: (c56 c57)

```

Figure A4 Listing of file t1.



cb13 LHS: [(Sibling x y) (> (set x 'R) (set y 'R)) (> (set x 'K-W) (set q 'R-w))]  
 Parallel : cb13 Anti: cb22 Parent: (set x 'R) Children: nil  
 cb14 LHS: [(Sibling x y) (> (set x 'R) (set y 'R)) (< (set x 'K-W) (set q 'R-w))]  
 Parallel : cb14 Anti: cb23 Parent: (set x 'R) Children: nil  
 cb15 LHS: [(Sibling x y) (> (set x 'R) (set y 'R)) (equal (set x 'K-W) (set q 'R-w))]  
 Parallel : cb15 Anti: cb24 Parent: (set x 'R) Children: nil  
 cb16 LHS: [(Sibling x y) (> (set x 'R) (set y 'R)) (> (set y 'K-W) (set q 'R-w))]  
 Parallel : cb16 Anti: cb20 Parent: (set x 'R) Children: nil  
 cb17 LHS: [(Sibling x y) (> (set x 'R) (set y 'R)) (< (set y 'K-W) (set q 'R-w))]  
 Parallel : cb17 Anti: cb19 Parent: (set x 'R) Children: nil  
 cb18 LHS: [(Sibling x y) (> (set x 'R) (set y 'R)) (equal (set y 'K-W) (set q 'R-w))]  
 Parallel : cb18 Anti: cb21 Parent: (set x 'R) Children: nil  
 cb19 LHS: [(Sibling x y) (< (set x 'R) (set y 'R)) (< (set y 'K-W) (set q 'R-w))]  
 Parallel : cb19 Anti: cb22 Parent: (set x 'R) Children: nil  
 cb20 LHS: [(Sibling x y) (< (set x 'R) (set y 'R)) (< (set x 'K-W) (set q 'R-w))]  
 Parallel : cb20 Anti: cb16 Parent: (set x 'R) Children: nil  
 cb21 LHS: [(Sibling x y) (< (set x 'R) (set y 'R)) (< (set y 'K-W) (set q 'R-w))]  
 Parallel : cb21 Anti: cb13 Parent: (set x 'R) Children: nil  
 cb22 LHS: [(Sibling x y) (< (set x 'R) (set y 'R)) (< (set x 'K-W) (set q 'R-w))]  
 Parallel : cb22 Anti: cb13 Parent: (set x 'R) Children: nil  
 cb23 LHS: [(equal (set x 'L) (set y 'L)) (> (set x 'R) (set y 'R)) (> (set x 'K-W) (set q 'R-w))]  
 Parallel : cb23 Anti: cb18 Parent: (set x 'R) Children: nil  
 cb24 LHS: [(equal (set x 'L) (set y 'L)) (< (set x 'R) (set y 'R)) (equal (set x 'K-W) (set q 'R-w))]  
 Parallel : cb24 Anti: cb18 Parent: (set x 'R) Children: nil  
 cb25 LHS: [(equal (set x 'L) (set y 'L)) (< (set x 'R) (set y 'R)) (> (set x 'K-W) (set q 'R-w))]  
 Parallel : cb25 Anti: cb28 Parent: (set x 'R) Children: nil  
 cb26 LHS: [(equal (set x 'L) (set y 'L)) (equal (set x 'R) (set y 'R)) (< (set x 'K-W) (set q 'R-w))]  
 Parallel : nil Anti: cb29 Parent: (set x 'R) Children: nil

---

cb27 LHS: [(equal (set x 'L) (set y 'L)) (equal (set x 'R) (set y 'R)) (equal (set x 'K-W) (set q 'R-w))]  
 Parallel : nil Anti: cb30 Parent: (set x 'R) Children: nil  
 cb28 LHS: [(equal (set x 'L) (set y 'L)) (equal (set x 'R) (set y 'R)) (> (set y 'K-W) (set q 'R-w))]  
 Parallel : nil Anti: cb25 Parent: (set x 'R) Children: nil  
 cb29 LHS: [(equal (set x 'L) (set y 'L)) (equal (set x 'R) (set y 'R)) (< (set y 'K-W) (set q 'R-w))]  
 Parallel : nil Anti: cb26 Parent: (set x 'R) Children: nil  
 cb30 LHS: [(equal (set x 'L) (set y 'L)) (equal (set x 'R) (set y 'R)) (equal (set y 'K-W) (set q 'R-w))]  
 Parallel : nil Anti: cb27 Parent: (set x 'R) Children: nil  
 Generator chooses rule cb13 of rank 2  
 Testing rule cb13 with data: nil  
 tester: rule cb13 is acceptable  
 LHS: [(Sibling x y) (> (set x 'R) (set y 'R)) (> (set x 'K-W) (set q 'R-w))]  
 #!+ve ex: \$#!-ve ex: 0 #!full ex: 0  
 tester: antiule cb22 is unacceptable  
 Scionger called with rule cb13  
 stronger: cb13 is parallel to acceptable c-rule cb13 -> increase rank of cb13  
 Generator chooses rule cb13 of rank 4  
 Testing rule cb13 with data: nil  
 tester: rule cb13 is acceptable  
 LHS: [(equal (set x 'L) (set y 'L)) (> (set x 'R) (set y 'R)) (> (set x 'K-W) (set q 'R-w))]  
 #!+ve ex: \$#!-ve ex: 0 #!full ex: 0  
 tester: antiule cb22 is unacceptable  
 Scionger called with rule cb13  
 stronger: cb13 is parallel to unacceptable c-rule cb22 -> cb22 is unacceptable  
 Generator chooses rule cb13 of rank 4  
 Testing rule cb13 with data: nil  
 tester: rule cb13 is acceptable  
 LHS: [(equal (set x 'L) (set y 'L)) (< (set x 'R) (set y 'R)) (< (set x 'K-W) (set q 'R-w))]  
 #!+ve ex: \$#!-ve ex: 0 #!full ex: 0  
 tester: antiule cb22 is unacceptable  
 Scionger called with rule cb13  
 stronger: cb13 is parallel to acceptable c-rule cb13 -> cb22 is unacceptable  
 Generator chooses rule cb13 of rank 2  
 Testing rule cb13 with data: nil  
 tester: rule cb13 is unacceptable  
 Scionger called with rule cb13  
 stronger: cb13 is parallel to unacceptable c-rule cb13 -> cb13 is unacceptable  
 Generator chooses rule cb13 of rank 2  
 Testing rule cb13 with data: nil  
 tester: rule cb13 is unacceptable  
 LHS: [(equal (set x 'L) (set y 'L)) (> (set x 'R) (set y 'R)) (equal (set x 'K-W) (set q 'R-w))]  
 #!+ve ex: \$#!-ve ex: 0 #!full ex: 0  
 tester: antiule cb22 is unacceptable  
 Scionger called with rule cb13  
 stronger: cb13 is parallel to unacceptable c-rule cb22 -> cb22 is unacceptable  
 Generator chooses rule cb13 of rank 4  
 Testing rule cb13 with data: nil  
 tester: rule cb13 is unacceptable  
 Scionger called with rule cb13

scrounger: cb18 is parallel to unacceptable rule cb18 -> cb18 is unacceptable  
 accounter: cb18 is parallel to unacceptable rule cb18 -> cb18 is unacceptable  
 Testing rule cb17 with delta: nil  
 letter: rule cb17 is unacceptable  
 tester: antituke cb19 is unacceptable  
 Scrounger called with rule cb17  
 scrounger: cb17 is parallel to unacceptable rule cb17 -> cb17 is unacceptable  
 accounter: cb20 is parallel to unacceptable rule cb20 -> cb20 is unacceptable  
 Generator chooses rule cb20 of rank 2  
 Testing rule cb19 with delta: nil  
 letter: rule cb19 is unacceptable  
 tester: antituke cb21 is unacceptable  
 Scrounger called with rule cb21  
 scrounger: cb21 is parallel to unacceptable rule cb21 -> cb21 is unacceptable  
 accounter: cb21 is parallel to unacceptable rule cb21 -> cb21 is unacceptable  
 Generator chooses rule cb22 of rank 1  
 Testing rule cb22 with delta: nil  
 letter: rule cb22 is acceptable  
 scrounger: cb22 is parallel to unacceptable rule cb22 -> cb22 is unacceptable  
 accounter: cb22 is parallel to unacceptable rule cb22 -> cb22 is unacceptable  
 Generator chooses rule cb23 of rank 1  
 Testing rule cb23 with delta: nil  
 letter: antituke cb23 is unacceptable  
 scrounger called with rule cb23  
 Generator chooses rule cb23 of rank 1  
 Testing rule cb23 with delta: nil  
 letter: antituke cb23 is unacceptable  
 tester: antituke cb24 is unacceptable  
 scrounger called with rule cb24  
 Stronger called with rule cb24  
 Generator chooses rule cb27 of rank 1  
 Testing rule cb27 with delta: nil  
 letter: rule cb27 is acceptable  
 scrounger: cb27 is unacceptable  
 letter: antituke cb28 is unacceptable  
 Scrounger called with rule cb27  
 Antituke  
 Derivation  
 $W + R \Rightarrow$   
 a. Type error  
 a.Def + [get x 'W'] [get q 'R+R\*']  
 a.OptProg nil  
 a.Defvar [w 'W+R']  
 a.Defvar [w 'W+R\*']  
 R + R  
 a.Type var  
 a.Def + [get x 'R'] [get q 'R+R\*']  
 a.Def + [get x 'R'] [get q 'R+R\*']  
 a.DefProg nil  
 a.Def [r 'R+R\*']  
 a.DefProg nil  
 a.Def [r 'R+R\*']  
 a.Def [r 'R+R\*']  
 Computing values of new symbols for problem instances  
 antituke: R + R is a hybrid of R  
 new condition: [> [get y 'R'] [get x 'R+R\*']] < [get x 'W'] [get y 'W+R\*']] [equal [get x 'W'] [get y 'W+R\*']]  
 antituke: R + R is a hybrid of R  
 new condition: [> [get y 'R'] [get x 'R+R\*']] < [get x 'R'] [get y 'R+R\*']] [equal [get x 'R'] [get y 'R+R\*']]  
 antituke: W + R\* is a hybrid of W  
 new condition: [> [get y 'W'] [get x 'W+R\*']] < [get x 'W'] [get y 'W+R\*']] [equal [get y 'W'] [get x 'W+R\*']]  
 antituke: W + R\* is a hybrid of W  
 new condition: [> [get y 'W'] [get x 'W+R\*']] < [get x 'W'] [get y 'W+R\*']] [equal [get x 'W+R\*'] [get y 'W+R\*']]  
 Parallels: cb24 Antituke: nil Parent: (c22) Children: nil  
 c24 LHS: [(Sibling x y) (< [get x 'R'] [get y 'R']) (get x 'W') (get y 'W+R\*')] Parent: (c22) Children: nil  
 Parallel: cb25 Antituke: nil Parent: (c22) Children: nil  
 c25 LHS: [(Sibling x y) (> [get x 'R'] [get y 'R']) (get x 'W+R\*')] Parent: (c22) Children: nil  
 Parallel: cb26 Antituke: nil Parent: (c22) Children: nil  
 c26 LHS: [(Sibling x y) (> [get x 'R'] [get y 'R']) (get x 'W+R\*')] Parent: (c22) Children: nil  
 Parallel: cb27 Antituke: nil Parent: (c22) Children: nil  
 c27 LHS: [(Sibling x y) (> [get x 'R'] [get y 'R']) (equal [get x 'R'] [get x 'R+R\*'])] Parent: (c22) Children: nil  
 Parallel: cb28 Antituke: nil Parent: (c22) Children: nil  
 c28 LHS: [(Sibling x y) (> [get x 'R'] [get y 'R']) (> [get x 'R'] [get y 'R+R\*'])] Parent: (c22) Children: nil  
 Parallel: cb29 Antituke: nil Parent: (c22) Children: nil  
 c29 LHS: [(Sibling x y) (> [get x 'R'] [get y 'R']) (< [get x 'R'] [get y 'R+R\*'])] Parent: (c22) Children: nil  
 Parallel: cb30 Antituke: nil Parent: (c22) Children: nil  
 c30 LHS: [(Sibling x y) (> [get x 'R'] [get y 'R']) (equal [get x 'R'] [get y 'R+R\*'])] Parent: (c22) Children: nil  
 Parallel: cb31 Antituke: nil Parent: (c22) Children: nil  
 c31 LHS: [(Sibling x y) (> [get x 'R'] [get y 'R']) (> [get x 'R'] [get y 'R']) (get x 'W+R\*')] Parent: (c22) Children: nil  
 Parallel: cb32 Antituke: nil Parent: (c22) Children: nil  
 c32 LHS: [(Sibling x y) (> [get x 'R'] [get y 'R']) (get x 'W+R\*')] Parent: (c22) Children: nil  
 Parallel: cb33 Antituke: nil Parent: (c22) Children: nil  
 c33 LHS: [(Sibling x y) (> [get x 'R'] [get y 'R']) (< [get x 'R'] [get y 'R']) (get x 'W+R\*')] Parent: (c22) Children: nil  
 Parallel: cb34 Antituke: nil Parent: (c22) Children: nil  
 c34 LHS: [(Sibling x y) (> [get x 'R'] [get y 'R']) (< [get x 'R'] [get y 'R']) (equal [get x 'R'] [get x 'R+R\*'])] Parent: (c22) Children: nil  
 Parallel: cb35 Antituke: nil Parent: (c22) Children: nil  
 c35 LHS: [(Sibling x y) (> [get x 'R'] [get y 'R']) (< [get x 'R'] [get y 'W']) (get x 'W+R\*')] Parent: (c22) Children: nil  
 Parallel: cb36 Antituke: nil Parent: (c22) Children: nil  
 c36 LHS: [(Sibling x y) (> [get x 'R'] [get y 'R']) (equal [get x 'W'] [get x 'W+R\*'])] Parent: (c22) Children: nil  
 Parallel: cb37 Antituke: nil Parent: (c22) Children: nil  
 c37 LHS: [(Sibling x y) (> [get x 'R'] [get y 'R']) (equal [get x 'W'] [get x 'W+R\*']) (get x 'R+R\*')] Parent: (c22) Children: nil  
 Parallel: cb38 Antituke: nil Parent: (c22) Children: nil  
 c38 LHS: [(Sibling x y) (> [get x 'R'] [get y 'R']) (> [get x 'R'] [get y 'R']) (get x 'W+R\*')] Parent: (c22) Children: nil  
 Parallel: cb39 Antituke: nil Parent: (c22) Children: nil  
 c39 LHS: [(Sibling x y) (> [get x 'R'] [get y 'R']) (< [get x 'R'] [get y 'R']) (get x 'W+R\*')] Parent: (c22) Children: nil  
 Parallel: cb40 Antituke: nil Parent: (c22) Children: nil  
 c40 LHS: [(Sibling x y) (< [get x 'R'] [get y 'R']) (> [get x 'R'] [get y 'R+R\*'])] Parent: (c22) Children: nil  
 Parallel: cb41 Antituke: nil Parent: (c22) Children: nil  
 c41 LHS: [(Sibling x y) (< [get x 'R'] [get y 'R']) (< [get x 'R'] [get y 'R']) (get x 'W') (get x 'W+R\*')] Parent: (c22) Children: nil  
 Parallel: cb42 Antituke: nil Parent: (c22) Children: nil  
 c42 LHS: [(Sibling x y) (< [get x 'R'] [get y 'R']) (> [get x 'R'] [get y 'W']) (get x 'W+R\*')] Parent: (c22) Children: nil  
 Parallel: cb43 Antituke: nil Parent: (c22) Children: nil  
 c43 LHS: [(Sibling x y) (< [get x 'R'] [get y 'R']) (> [get x 'W'] [get x 'W+R\*'])] Parent: (c22) Children: nil  
 Parallel: cb44 Antituke: nil Parent: (c22) Children: nil  
 c44 LHS: [(Sibling x y) (< [get x 'R'] [get y 'R']) (< [get x 'W'] [get x 'W+R\*'])] Parent: (c22) Children: nil  
 Parallel: cb45 Antituke: nil Parent: (c22) Children: nil  
 c45 LHS: [(Sibling x y) (< [get x 'R'] [get y 'R']) (equal [get x 'W'] [get x 'W+R\*'])] Parent: (c22) Children: nil  
 Parallel: cb46 Antituke: nil Parent: (c22) Children: nil  
 c46 LHS: [(Sibling x y) (< [get x 'R'] [get y 'R']) (equal [get x 'W'] [get x 'W+R\*']) (get x 'R+R\*')] Parent: (c22) Children: nil  
 Parallel: cb47 Antituke: nil Parent: (c22) Children: nil  
 c47 LHS: [(Sibling x y) (< [get x 'R'] [get y 'R']) (get x 'W') (get x 'W+R\*')] Parent: (c22) Children: nil  
 Parallel: cb48 Antituke: nil Parent: (c22) Children: nil  
 c48 LHS: [(Sibling x y) (< [get x 'R'] [get y 'R']) (equal [get x 'W'] [get x 'W+R\*']) (get x 'R+R\*')] Parent: (c22) Children: nil  
 Parallel: cb49 Antituke: nil Parent: (c22) Children: nil  
 c49 LHS: [(Sibling x y) (> [get x 'R'] [get y 'R']) (get x 'W') (get x 'W+R\*')] Parent: (c22) Children: nil  
 Parallel: cb50 Antituke: nil Parent: (c22) Children: nil  
 c50 LHS: [(Sibling x y) (> [get x 'R'] [get y 'R']) (equal [get x 'W'] [get x 'W+R\*']) (get x 'R+R\*')] Parent: (c22) Children: nil



```

Scrounger called with rule cb34
Generator chooses rule cb39 of rank 2
Testing rule cb39 with delta nil
tester: rule cb39 is unacceptable
Scrounger called with rule cb39
Scrounger: cb35 is parallel to unacceptable <rule cb39 -> cb35 is unacceptable
Generator chooses rule cb30 of rank 2
Testing rule cb30 with delta nil
tester: rule cb30 is acceptable
Scrounger called with rule cb30
Scrounger: cb35 is parallel to acceptable <rule cb30 -> increase rank of cb36
Generator chooses rule cb36 of rank 4
Testing rule cb36 with delta nil
tester: rule cb36 is acceptable

LHS : [(equal [set x 'R] [set y 'R]) (equal [set x 'R] [set y 'R + R1])]

#(t+ve ex) : cb36 #(t+ve ex) : 0 #(null ex) : 0

Scrounger called with rule cb36
Generator chooses rule cb37 of rank 2
Testing rule cb37 with delta nil
tester: rule cb37 is impossible
Scrounger called with rule cb37
Scrounger: cb37 is parallel to impossible <rule cb37 -> increase rank of cb37
Generator chooses rule cb37 of rank 4
Testing rule cb37 with delta nil
tester: rule cb37 is unacceptable
Scrounger called with rule cb37
Scrounger: rule cb37 is impossible
Scrounger called with rule cb37
Scrounger: cb38 is parallel to unacceptable <rule cb37 -> cb38 is unacceptable
Generator chooses rule cb32 of rank 2
Testing rule cb32 with delta nil
tester: rule cb32 is impossible
Scrounger called with rule cb32
Scrounger: cb38 is parallel to unacceptable <rule cb32 -> cb38 is unacceptable
Scrounger called with rule cb32
Scrounger: cb38 is parallel to unacceptable <rule cb32 -> increase rank of cb38
Generator chooses rule cb32 of rank 2
Testing rule cb32 with delta nil
tester: rule cb32 is unacceptable
Scrounger called with rule cb32
Scrounger: cb38 is parallel to unacceptable <rule cb32 -> cb38 is unacceptable
Scrounger called with rule cb32
Scrounger: cb38 is parallel to unacceptable <rule cb32 -> increase rank of cb38
Generator chooses rule cb33 of rank 2
Testing rule cb33 with delta nil
tester: rule cb33 is unacceptable
Scrounger called with rule cb33
Scrounger: cb38 is parallel to unacceptable <rule cb33 -> cb38 is unacceptable
Scrounger: cb38 is parallel to unacceptable <rule cb33 -> increase rank of cb38
Generator chooses rule cb33 of rank 2
Testing rule cb33 with delta nil
tester: rule cb33 is impossible
Scrounger called with rule cb33
Scrounger: cb38 is parallel to unacceptable <rule cb33 -> cb38 is unacceptable
Scrounger called with rule cb33
Scrounger: cb38 is parallel to unacceptable <rule cb33 -> increase rank of cb38
Generator chooses rule cb34 of rank 2
Testing rule cb34 with delta nil
tester: rule cb34 is unacceptable
Scrounger called with rule cb34
Scrounger: cb38 is parallel to unacceptable <rule cb34 -> cb38 is unacceptable
Generator chooses rule cb34 of rank 4
Testing rule cb34 with delta nil
tester: rule cb34 is unacceptable
Scrounger called with rule cb34
Scrounger: cb38 is parallel to unacceptable <rule cb34 -> cb38 is unacceptable
Scrounger called with rule cb34
Scrounger: cb38 is parallel to unacceptable <rule cb34 -> increase rank of cb38
Generator chooses rule cb34 of rank 4

```

Testing rule cb48 with delta nil  
 tester: rule cb48 is impossible  
 Scrounger called with rule cb48  
 Generator chooses rule cb48 of rank 2  
 Testing rule cb49 with delta nil  
 tester: rule cb49 is impossible  
 Scrounger called with rule cb49  
 Generator chooses rule cb49 of rank 2  
 Testing rule cb49 with delta nil  
 Testing rule cb44 with delta nil  
 tester: rule cb44 is unacceptable  
 Scrounger called with rule cb44  
 scrounger: cb44 is parallel to impossible c-rule cb44 -> increase rank of cb44  
 Generator chooses rule cb45 of rank 2  
 Testing rule cb45 with delta nil  
 tester: rule cb45 is impossible  
 Scrounger called with rule cb45  
 scrounger: cb45 is parallel to impossible c-rule cb45 -> cb45 is unacceptable  
 Generator chooses rule cb50 of rank 2  
 Testing rule cb50 with delta nil  
 tester: rule cb50 is unacceptable  
 Scrounger called with rule cb50  
 scrounger: cb50 is parallel to impossible c-rule cb44 -> cb50 is unacceptable  
 Generator chooses rule cb51 of rank 2  
 Testing rule cb51 with delta nil  
 tester: rule cb51 is unacceptable  
 Scrounger called with rule cb51  
 Generator chooses rule cb52 of rank 2  
 Testing rule cb52 with delta nil  
 tester: rule cb52 is impossible  
 Scrounger called with rule cb52  
 scrounger: cb52 is parallel to impossible c-rule cb48 -> increase rank of cb52  
 Generator chooses rule cb53 of rank 4  
 Testing rule cb52 with delta nil  
 tester: rule cb52 is impossible  
 Scrounger called with rule cb52  
 Generator chooses rule cb53 of rank 2  
 Testing rule cb53 with delta nil  
 tester: rule cb53 is unacceptable  
 Scrounger called with rule cb53  
 scrounger: cb53 is parallel to unacceptable c-rule cb48 -> cb53 is unacceptable  
 Generator chooses rule cb54 of rank 2  
 Testing rule cb54 with delta nil  
 tester: rule cb54 is impossible  
 Scrounger called with rule cb54  
 Generator chooses rule cb55 of rank 1  
 Testing rule cb55 with delta nil  
 tester: rule cb55 is impossible

Scrounger called with rule cb55  
 Generator chooses rule cb55 of rank 1  
 Testing rule cb56 with delta nil  
 tester: rule cb56 is unacceptable  
 Scrounger called with rule cb56  
 Generator chooses rule cb56 of rank 1  
 Testing rule cb57 with delta nil  
 tester: rule cb57 is impossible  
 Scrounger called with rule cb57  
 Generator chooses rule cb58 of rank 1  
 Testing rule cb58 with delta nil  
 tester: rule cb58 is unacceptable  
 Scrounger called with rule cb58  
 Generator chooses rule cb58 of rank 1  
 Testing rule cb59 with delta nil  
 tester: rule cb59 is impossible  
 Scrounger called with rule cb59  
 Generator chooses rule cb59 of rank 1  
 Testing rule cb59 with delta nil  
 tester: rule cb59 is impossible  
 Scrounger called with rule cb59  
 Generator chooses rule cb60 of rank 1  
 Testing rule cb60 with delta nil  
 tester: rule cb60 is impossible  
 Scrounger called with rule cb60  
 Generator chooses rule cb61 of rank 1  
 Testing rule cb61 with delta nil  
 tester: rule cb61 is impossible  
 Scrounger called with rule cb61  
 Generator chooses rule cb62 of rank 1  
 Testing rule cb62 with delta nil  
 tester: rule cb62 is unacceptable  
 Scrounger called with rule cb62  
 Generator chooses rule cb63 of rank 1  
 Testing rule cb63 with delta nil  
 tester: rule cb63 is acceptable  
 LHS: [equal (set x 'U) (set y 'L)] [equal (set x 'R) (set y 'R)] [equal (set y 'W) (set x 'W + R w)]  
 #[(+ ve ex) : 2    #(- ve ex) : 0    #([null ex] : 1)]

Scrounger called with rule cb63  
 Generator chooses rule cb63 of rank 1  
 Testing rule cb64 with delta nil  
 tester: rule cb64 is impossible  
 Scrounger called with rule cb64  
 Generator chooses rule cb65 of rank 1  
 Testing rule cb65 with delta nil  
 tester: rule cb65 is unacceptable  
 Scrounger called with rule cb65  
 Generator chooses rule cb65 of rank 1  
 Testing rule cb66 with delta nil  
 tester: rule cb66 is unacceptable  
 Scrounger called with rule cb66  
 Analyzer  
 Derivator  
 Trace ends

**VITA**

Chee Pen Yu was born in Johor Bahru, Malaysia, on November 3, 1957. He received the B.E. (Hons.) degree from University of Malaya, Malaysia, in 1980, and the M.S.E.E. degree from Purdue University, West Lafayette, in 1983.

**VITA**